

# **CodeBase 6.0™**

## **User's Guide**

**The Delphi Engine For Database Management**  
**Clipper Compatible**  
**dBASE Compatible**  
**FoxPro Compatible**

**Sequiter® Software Inc.**

© Copyright Sequiter Software Inc., 1988-1996. All rights reserved.

No part of this publication may be reproduced, or transmitted in any form or by any means without the written permission of Sequiter Software Inc. The software described by this publication is furnished under a license agreement, and may be used or copied only in accordance with the terms of that agreement.

The manual and associated software is sold with no warranties, either expressed or implied, regarding its merchantability or fitness for any particular purpose. The information in this manual is subject to change without notice and does not represent a commitment on the part of Sequiter Software Inc.

CodeBase<sup>®</sup> and CodeReporter<sup>™</sup> are trademarks of Sequiter Software Inc.

Borland C++<sup>®</sup> is a registered trademark of Borland International.

Borland Delphi<sup>®</sup> is registered trademark of Borland International.

Clipper<sup>®</sup> is a registered trademark of Computer Associates International Inc.

FoxPro<sup>®</sup> is a registered trademark of Microsoft Corporation.

Microsoft Visual C++<sup>®</sup> is a registered trademark of Microsoft Corporation.

Microsoft Windows<sup>®</sup> is a registered trademark of Microsoft Corporation.

OS/2<sup>®</sup> is a registered trademark of International Business Machines Corporation.

---

# Contents

---

Introduction .....	5
How To Use This Manual .....	6
1 Database Concepts.....	7
CodeBase File Format .....	7
2 DataBase Access .....	11
CodeBase Components .....	11
Learning CodeBase .....	14
The Delphi API .....	14
An Example CodeBase Program.....	14
Initializing The CODE4 Structure .....	17
Code fragment from ShowData .....	17
Opening Data Files .....	17
Closing Data Files.....	18
Moving Between Records .....	18
Accessing Fields .....	19
Retrieving A Field's Contents .....	21
Creating Data Files .....	22
Adding New Records.....	26
Assigning Field Values.....	27
Removing Records .....	28
Data File Information .....	31
Advanced Topics .....	33
3 Indexing.....	35
Indexes & Tags.....	35
Index Expressions.....	37
Creating An Index File .....	37
Maintaining Index Files .....	41
Opening Index Files .....	42
Referencing Tags.....	42
Selecting Tags.....	44
Tag Filters .....	44
Unique Keys .....	47
Seeking .....	48
Group Files .....	54
Bypassing Group Files .....	54
Reindexing.....	58
Advanced Topics .....	59
4 Queries And Relations .....	61
Relations.....	61
Complex Relations.....	65
Relation Types.....	67
Creating Relations .....	69
Setting The Relation Type.....	72
Moving Through The Composite Data File .....	73
Queries And Sorting.....	74
Accessing The Query Set.....	77
Queries On Multi-Layered Relation Sets .....	78
Lookups On Relations.....	79
Iterating Through The Relations .....	82

5 Query Optimization .....	83
What is Query Optimization .....	83
When is Query Optimization used.....	83
How To Use Query Optimization .....	86
6 Date Functions .....	89
Date Pictures .....	89
Date Formats .....	89
Code fragment from Date .....	92
Code fragment from Date .....	92
Code fragment from Date .....	92
Testing For Invalid Dates .....	92
Other Date Functions .....	92
Code fragment from Date .....	93
7 Memory Optimizations.....	95
Using Memory Optimizations .....	95
Memory Requirements.....	97
When To Use Memory Optimization .....	98
8 Multi-User Applications .....	101
Locking .....	101
Creating Multi-User Applications .....	102
Common Multi-User Tasks.....	104
Multi-User Optimizations.....	108
Avoiding Deadlock .....	111
Exclusive Access .....	112
Read Only Mode.....	113
Code fragment from Multi .....	113
Lock Attempts.....	113
Automatic Record Locking .....	114
Enforced Locking .....	114
9 Performance Tips.....	115
Memory Optimization.....	115
Memory Requirements.....	115
Observing the performance improvement .....	115
Functions that can reduce performance .....	115
Locking .....	115
Appending .....	116
Time Consuming Functions .....	116
Queries and Relations.....	116
General Tips .....	117
10 Transaction Processing .....	119
Transactions .....	119
Logging in the Stand-Alone Case.....	122
Logging in the Client-Server case .....	124
Locking .....	124
11 Index .....	127

---

# Introduction

---

CodeBase is a high performance database engine for C, C++, Delphi and Visual Basic programmers.

CodeBase is compatible with dBASE IV, FoxPro and Clipper file formats. CodeBase allows you to create, access or change their data, index or memo files. Applications written in CodeBase can seamlessly share files with concurrently running dBASE IV, FoxPro and Clipper applications.

CodeBase has four different interfaces depending on programming language. CodeBase supports the C++, C, Visual Basic languages and Pascal under Delphi. In addition, these interfaces can be used under a variety of operating systems, including DOS, Windows 3.1/95/NT or OS/2 and UNIX, depending on the language.

CodeBase is scalable to your needs, which means that it can be used in a stand-alone, multi-user or client-server configuration.

CodeBase has many useful features including a full set of query and relation functions. These functions use Query Optimization to greatly reduce the time it takes to perform queries and generate reports. CodeBase also has logging and transaction capabilities, which allow you to control the integrity of a database easily.

This *User's Guide* discusses the CodeBase Delphi API (Application Program Interface). The CodeBase Delphi API allows the programmer to write applications using CodeBase in Pascal under Borland Delphi. It is necessary to have Borland Delphi for Windows and some knowledge of how to program under Delphi.

Refer to the *Getting Started* booklet for information on installing CodeBase, running the example programs and contacting Sequiter Software.

---

## How To Use This Manual

This manual was designed for both novice and veteran database programmers. The *CodeBase 6.0 Getting Started* booklet details how to execute an example CodeBase application. This booklet is recommended reading for all users of CodeBase. The "Database Concepts" chapter of this guide discusses database terminology. Even if you are an expert database programmer, you should at least skim this chapter because some of the terms, such as tags and indexes, may be used differently than you expect. The remaining chapters deal with the basics of writing database applications with the CodeBase library. For your convenience, all example programs illustrated in this manual can be found on disk.

---

## Manual Conventions

This manual uses several conventions to distinguish between the various language constructs. These conventions are listed below:

CodeBase classes, functions, structures, constants, and defines are always shown highlighted as follows:

e.g. **d4create** , **CODE4**, **r4eof**, **E4DEBUG**

dBASE expressions are always contained by two double quotes ( " " ).

e.g. " 'Hello There ' + FIRST\_NAME"

You should note that the containing quotes are part of the Pascal syntax rather than part of the dBASE expression.

dBASE functions are displayed in upper case and are denoted by a trailing set of parenthesis.

e.g. UPPER() , STR(), DELETED()

Pascal language functions and constructs are shown in bold typeface.

e.g. **for**, **write**, **while**

# 1 Database Concepts

If you are not familiar with the concept of a *database*, simply think of it as a collection of information which has been organized in a logical manner. The most common example of a database is a telephone directory. It contains the names, phone numbers, and addresses of thousands of people. Each listing in the phone book corresponds to one *record*, and each piece of information in the record corresponds to one *field*. The phone book excerpt below illustrates this concept.

Name Field	Address Field	Phone # Field
Smith John	897 Elm Street	555-3456
Stevens Dave	456 Oak Lane	555-1234
Trumble Al	123 Maple Ave	555-4567
Tuna Peter	955 Pine Ridge	555-2345
Tunner Sam	634 Spruce Ave	555-6789
Victor Paul	344 Knottwood Blvd	555-6423

**Record # 5**

Figure 1.1 Phone Book Example

As shown in this example, each *data file* (also known as a table) is a collection of one or more fields (also known as a tuple). Each of these fields has a set of characteristics that determine the size and type of data to be stored. Collectively, these field descriptions make up the structure of your data file.

## CodeBase File Format

CodeBase uses the industry standard .DBF data files. This standard allows for several types of fixed width fields and one type of variable length field.

### Fields

The .DBF standard uses four attributes to describe each field. These attributes are Name, Type, Length, and Decimal. They are listed below. For detailed information on field attributes, please refer to the **FIELD4INFO** structure as described in the CodeBase *Reference Guide* under **d4create**.

- **Name:** This simply refers to the name that will be used to identify the field. Each field name can be a maximum of 10 characters, and each field name must be unique within a data file and must consist of alphanumeric or underscore characters. The first character of the name must be a letter.
- **Type:** The type of the field determines what kind of information should be stored in the field. There are eight different types that can be specified for any given field. They are Character, Numeric, Floating Point, Date, Logical, Memo, Binary and General.

- **Length:** This attribute refers to the number of characters or digits that can be stored in the field.
- **Decimal:** This attribute applies only to Numeric and Floating Point fields. It specifies the number of digits after the decimal point.

For a more in-depth discussion of the various field types and their attributes, refer to the "Field Functions" chapter of the *Reference Guide*.

---

## Records

A record consists of all information associated with one entry in a data file. Typically, every line in a phone book would correspond to a new record. Each data file record can be identified by its *record number*, which indicates its physical position in the data file.

In addition, each record includes an extra byte of storage that is used as a *deletion flag*. When a record is no longer needed in a data file, it can be 'flagged' for later removal by setting its deletion flag to true (non-zero). A record that is marked for deletion is not physically removed from the file until it undergoes a process known as packing. This method gives one the opportunity to recall any record scheduled for deletion by simply setting the deletion flag to false before packing the file.

---

## Indexes & Tags

An index is a way to sort data file information without actually reorganizing the data itself. The index at the back of a book is a good metaphor for data file indexes. The various topics of the book are sorted alphabetically for easy lookup, and each entry has a corresponding reference to the physical page number where the information can be found.

Indexes are also smaller than data files, and subsequently faster to sort and manipulate, because they are based on only a portion of the record, usually just a field value, rather than the entire record itself. This value is called the *index key*. The key can be a single field value or a combination of field values. For example, if a phone book data file has the fields LAST\_NAME and FIRST\_NAME, then an index key could be built based on a combination of the two fields. A valid expression for the key would be 'LAST\_NAME + FIRST\_NAME'.

In addition to including field values, key expressions can include a number of dBASE expression functions and operators. For example, the '+' operator concatenates two character fields together, or adds numeric values together. The dBASE function STR() , converts a non-character field into a string. The DELETED() function returns true if the record is marked for deletion. Here is a typical expression using these functions and operators:

```
'LAST_NAME + STR(NUM_FLD) + DTOS(DELETE_FLD)'
```

The previous expression combines a character field with the string representation of a Numeric field and a Date field.



In addition, dBASE IV, FoxPro and Clipper indexes can be built with 'filters' using expressions such as this,

`'LAST_NAME = "SMITH" .AND. .NOT. DELETED()'`

to cause only a subset of the data file records to have corresponding key entries in the index. Filters are discussed further in this chapter.

Besides containing a key for each valid record, the index also contains the physical record number that corresponds to each key. Thus, when you use an index to seek a certain value, the corresponding record can be located from the data file using the record number information. Here is how an index file might look, based on a key expression of the field NAME.

Figure 1.2 Data and Index Representation

Physical Order:

Rec. No.	NAME	COUNTRY
1	DAVIDSON	CANADA
2	SMITH	UNITED STATES
3	ALBERTSON	BRITAIN
4	FOREST	CHINA
5	BAKER	BRAZIL

Index File: key = NAME

Rec. No.	KEY VALUE
3	ALBERTSON
5	BAKER
1	DAVIDSON
4	FOREST
2	SMITH

The most common use of indexed is to implement quick lookups of one or more records. Another advantage to index files is that you can build as many as you need for any data file.

There are several formats of index files that CodeBase supports.

- **".MDX"** (dBASE IV)
- **".NTX"** (Clipper)
- **".CDX"** (FoxPro)

The **".CDX"** and **".MDX"** allow you to have multiple tags in each index file, and to have *production indexes*. A production index is an index file that is opened automatically when the associated data file is opened. The **".NTX"** format limits you to one tag per index file, and does not allow for production indexes. CodeBase does, however, provide this functionality through the use of group files. Refer to the "Indexing" chapter for more details on index file formats.

CodeBase uses the term *tag* to refer to any sort ordering, whether it be an index within a multiple index file, or a single index file. The term *selecting a tag* means putting an open index into use.

---

## Filters

Filters are used to obtain a subset of the available data file records. The subset is based on true/false conditions calculated from data file field information. Only records that pass through the filter have entries in the tag. This allows for fast access to a data subset. Refer to the "Indexing" chapter for details on using filters.

---

## Relations

A relation is a connection between two or more data files. A relation determines how records in related data files can be found from a record in the current data file. Refer to the "Queries And Relations" chapter for details on relations.

## 2 DataBase Access

Now that you have been introduced to the concepts of database management, you're ready to start programming some simple database programs using Delphi. This chapter take you through the details of manipulating the data files by explaining how to create and open data files, store and retrieve data, and how records are added and deleted.

To run any of the tutorial programs, refer to the *Getting Started* book. Tutorial source code is installed into subdirectory "EXAMPLES".

### CodeBase Components

In addition to the elements of standard Pascal programs, every CodeBase program will contain CodeBase functions, structures, and return codes and other constants.

### CodeBase Functions

All data manipulation is performed by CodeBase functions that can be identified by the format of their names. These names are mixed case lower case starting with one to six letters followed by the number four. The leading letters indicate the module of the function while the remaining portion of the name describes its purpose. For example, functions that manipulate data files start with **d4**, while all the date functions start with **date4**.

#### Function Success and Failure

CodeBase functions return -1 to indicate an error and 0 to indicate success. For this reason it is advisable to test the return values from CodeBase functions using the return code constants documented in Appendices A and B of the *Reference Guide* and defined in CODEBASE.PAS for Windows.

For example, the function **d4seek**, which does a look up in an index tag, returns **r4success** (zero) if the seek is successful and -1 if an error has occurred. The following are incorrect and correct examples of how to use this function.

Incorrect:

```
If d4seek( db, "JONES" ) Then
    writeln( 'We found Him' )
Else ...
```

Correct:

```
If d4seek( db, "JONES" ) = r4success Then
    writeln( 'We found Him' )
Else ...
```

The incorrect example assumes **d4seek** returns -1 or True upon success, thus when zero is returned by **d4seek** to indicate a CodeBase success the Else statement is executed instead of reporting that the seek value was found. In the correct version, the return code is compared to the constant **r4success** to determine whether the seek succeeded. Watch for this type of situation when writing applications.

The exceptions to this rule are the CodeBase functions that must return pointers. These functions returns zero when the function fails and usually set the **CODE4.errorCode** to a negative value.

Logical Parameters	Many CodeBase functions accept logical values as parameters. In these cases, a <i>true</i> value corresponds to the value 1, while a <i>false</i> value corresponds to 0. Thus, the value zero can represent both success and failure in functions depending on the return type and it can also represent false when used as a logical parameter.
Functions and Subs	A final note about our use of the word <i>function</i> to describe CodeBase routines. In Pascal, those routines that do not return a value are known as <i>Procedures</i> and those that do return a value are referred to as <i>Functions</i> . For the sake of consistency, the CodeBase <i>Reference Guide</i> and <i>User's Guide</i> refer to all routines as functions regardless of whether they return a value or not. Each function listed in the <i>Reference Guide</i> has a "usage" section that indicates whether the function returns a value.
<b>CodeBase Structures</b>	CodeBase structures are used for storing information and for referencing objects such as data files and fields. CodeBase structures follow the same naming conventions as CodeBase functions except that they are all upper case.
<b>CodeBase Constants</b>	In addition to CodeBase structures and functions, you can also use CodeBase constants. Most constants are integers which are used for return values and error codes. These constants have names which usually start with <b>r4</b> , or <b>e4</b> . Some examples include <b>r4success</b> , <b>r4locked</b> , <b>e4memory</b> , and <b>e4open</b> . These constants are listed in more detail in "Appendix A: Error Codes" and "Appendix B: Return Codes" of the CodeBase <i>Reference Guide</i> . All constants are defined in CODEBASE.PAS, and can be included in any of your projects.
<b>Overview</b>	The next section introduces some of the important CodeBase structures and examines how they are related to data files.

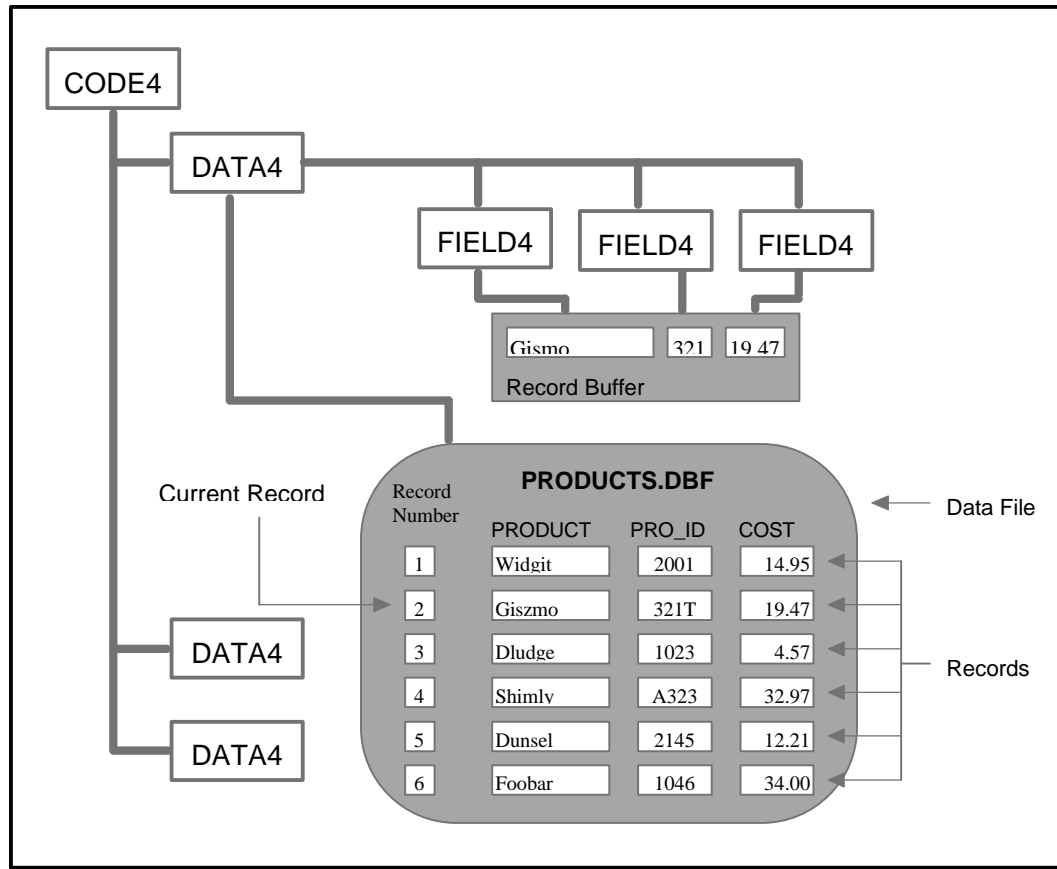


Figure 2.1 CodeBase Structures

**The CODE4 Structure** Foremost of the CodeBase structures is the **CODE4**. This structure contains settings and information used by most of the other CodeBase functions. All **CODE4** elements are accessed through **code4...** functions, such as **code4safety** and **code4optimize**. This structure is created internally and is initialized to its default values when **code4init** is called.



### Note

Most applications only require one **CODE4** structure. If you are using more than one **CODE4** structure, you will most likely be wasting memory resources. Exceptions to this rule are discussed later in this chapter.

**The DATA4 structure** The **DATA4** structure is used to reference a particular data file. When a data file is opened, the information for the **DATA4** structure is allocated and stored in the **CODE4** structure. When you call a function that operates on a data file, you specify which data file you want by passing the function a pointer to its **DATA4** structure. ( For convenience, the terms "pointer to its **DATA4** structure" and "**DATA4** pointer" will be used interchangeably. ) Note that more than one **DATA4** structure can reference the same open data file. However, since the information for the structure is stored within the **CODE4** structure, each new structure constructed with **code4data** contains the same information, including the same current record and optimization settings.

The record buffer Each data file has a memory area associated with it called the *record buffer*, which is used to store one record from the data file. The record that is stored in the record buffer is called the *current record*. For example, when the data file is positioned to the top record, record number one becomes the current record and its contents are read from disk and stored in the record buffer. Access to the record is then performed through the buffer, which is really just a copy of the disk data.

Changes made to the record buffer using the field functions are automatically written to the data file before any new current record is read from disk.

The **FIELD4** structure A **FIELD4** structure is used to reference a particular field in the record buffer. When the data file is opened, a **FIELD4** structure is automatically allocated for each of its fields. When you call any function that uses a specific field of the record buffer, you specify which field by passing a **FIELD4** pointer. Access to the **FIELD4** pointer is obtained with the **d4field** function.

---

## Learning CodeBase

As explained in the Introduction, the code samples for all the subsequent chapters in the *User's Guide* can be found in the `..\EXAMPLES` sub-directory.

Each example presented in this guide has a name, which corresponds to a program name in `..\EXAMPLES`.

The Delphi API Some of the example code presented here deals with the Delphi API itself, such as filing list boxes, displaying data and so forth. These details are not discussed because they are not directly related to the functionality of CodeBase. If you have any questions about these sections, refer to the *Delphi Component Writer's Guide*.

---

## An Example CodeBase Program



### PROGRAM ShowData

At this point, you are ready to write your first useful program that uses the CodeBase library. Refer to the *Getting Started* book for information on how to run the example programs.

The first example program is called ShowData. This example opens a data file and prints its contents to the screen. Following the code listing is a brief explanation of how the program works. Don't worry if some things aren't totally clear after reading the explanation. The details will be explained further in this chapter and the rest of the guide.

```
{ showdata.pas (c)Copyright Sequiter Software Inc., 1993-1994. All rights reserved. }
{ See User's Guide, page 34 }

{$N+}

{$I CODEBASE.INC}

program SHOWDATA ;

{$IFDEF WINDOWS}
    uses CodeBase, WinCrt, Sysutils ;
{$ELSE}
    uses CodeBase, Crt, Strings ;
{$ENDIF}

var
    code_base : CODE4 ;
    data_file : DATA4 ;
    field_ref : FIELD4 ;
    rc, j, num_fields : integer ;
    field_contents : PChar ;
```

```

    dataname : array[0..255] of char ;
begin
    if ParamCount <> 1 then
    begin
        writeln( ' USAGE: SHOWDATA <FILENAME.DBF>' ) ;
        Halt ;
    end ;

    code_base := code4init ;

    StrPCopy( dataname, ParamStr(1) ) ;
    data_file := d4open( code_base, dataname ) ;
    error4exitTest( code_base ) ;

    num_fields := d4numFields( data_file ) ;
    rc := d4top( data_file ) ;
    while rc = r4success do
    begin
        for j := 1 to num_fields do
        begin
            field_ref := d4fieldJ( data_file, j ) ;
            field_contents := f4memoStr( field_ref ) ;
            write(field_contents, ' ' ) ;
        end ;
        writeln( ' ' ) ;
        rc := d4skip( data_file, Longint(1) ) ;
    end ;

    d4close( data_file ) ;
    code4initUndo( code_base ) ;

    {$IFDEF WINDOWS}
        readKey ;
        DoneWinCrt ;
    {$ENDIF}

    Halt ;
end.

```

---

**Explanation :** The next section contains a brief explanation of the ShowData program, which is followed by a more detailed explanation.

The first step required in any CodeBase application is the initialization of the library. This step must occur before you attempt to call any other functions. Failing to do so will lead to unpredictable results. Initialization is accomplished through **code4init**.

**code4init** returns a pointer to an internal **CODE4** structure that CodeBase creates and initializes. In most cases, this return value should be stored in a variable that has a global application scope.

```

var
    code_base: CODE4;          'Declare a CODE4 pointer
...
begin
    code_base:= code4init;      'Initialize CodeBase

```

If **code\_base** equals nil after **code4init** has been called, an error has occurred. When 'ShowData' is executed, **d4open** is called to open the data file. The second parameter of **d4open** specifies which data file to open, in this case SHOWDATA.DBF. At this point, if **data\_file** or **code\_base** equal nil, an error has occurred and the subroutine is exited.

```

db := d4open( code_base, 'SHOWDATA' ) ;
error4exitTest( code_base ) ;

```

**Note**

If an error does occur, CodeBase sets **CODE4.errorCode** to a negative value corresponding to the error that occurred. Other CodeBase functions will not respond until this error code is reset to zero. This is accomplished through the **code4errorCode** function.

```
rc := code4errorCode( cb, 0 ) ;
```

If the file is opened without error, the program then moves to the top of the file by calling **d4top**, which loads the first record into the record buffer.

```
rc := d4top( data_file ) ;
```

Next, a **While** loop is used to load each subsequent record into the record buffer. The loop terminates when it has iterated through all the records in the file. The number of records in the file is returned **d4recCount**.

```
while rc = r4success do
begin
.
.
.
end ;
```

Inside this **While** loop, a **For** loop is used to iterate through the fields by incrementing a counter from one to the total number of fields in the data file. The number of fields in the data file is obtained by a call to **d4numFields**.

```
for j := 1 to num_fields do
begin
.
.
.
end ;
```

This counter is used in conjunction with the **d4fieldJ** function to obtain a pointer to a field's **FIELD4** structure. The **FIELD4** pointer is passed to **f4memoStr**, which in turn returns a pointer to the contents of that field so it can be printed out.

```
field_ref := d4fieldJ( data_file, j ) ;
field_contents := f4memoStr( field_ref ) ;
write(field_contents, ' ' ) ;
```

After the inner loop has completed, one record has been printed to the screen. The next record is then loaded in the record buffer with the **d4skip** function.

```
rc := d4skip( data_file, Longint(1) ) ;
```

If everything has executed properly, the next step is to close the data file with the **d4close** function.

```
d4close( data_file ) ;
```



Before the program is done, the last step is to place a call to **code4initUndo**. This function is important for freeing any resources that CodeBase has allocated.

```
code4initUndo( code_base ) ;
```

## Initializing The CODE4 Structure

Code fragment  
from ShowData

Since the **CODE4** structure contains random data when the program starts, the **CODE4** structure must be initialized before any CodeBase functions can be used.

```
code_base := code4init;
```

When the **CODE4** structure is initialized, all of the flags and member variables are set to their default values. The **code4init** function should normally only be called once in the application.

After **code4init** has been called, you can then change the value of the **CODE4** structure's flags. Refer to the "CodeBase Members and Functions" chapter in the *Reference Guide* for an explanation on all the members of the **CODE4** structure and how to change their default values.

## Opening Data Files

Code fragment  
from ShowData

Only after a **CODE4** structure is initialized, can data files be opened with **d4open**.

```
data_file := d4open( code_base, dataname ) ;
```

**d4open** accepts the **CODE4** pointer as its first parameter and a string containing the path and file name of a data file as its second. If the specified file name does not contain an extension, the default extension of **".DBF"** is used. If no path is specified, the current directory is searched.

Obtaining a **DATA4**  
pointer

If the file is located, it is opened and the address to an internal **DATA4** structure is returned. This address is used as a reference to the data file.



### Note

Almost every function that starts with **d4** takes a **DATA4** pointer as a parameter. This pointer specifies which data file the function should operate on.

Checking for errors

If the file does not exist, or if CodeBase detects any other error, an error message is displayed and a nil is returned by **d4open**. **error4exitTest** is called to check for any errors. It will terminate the program if one has occurred.

Code fragment  
from ShowData

```
error4exitTest( code_base ) ;
```

An alternative to calling **error4exitTest** and exiting the application, is to check **CODE4.errorCode**. If **d4open** has failed, **CODE4.errorCode** will be set to a negative value such as **e4open**. This error code must be reset to zero before other CodeBase function calls can be made.

```
rc := code4errorCode( code_base, 0 ) ;
```

After the data file has been successfully opened, the current record number is set to '-1' to indicate that a record has not yet been read into the record buffer.

---

## Closing Data Files

Before your application completes, it is important that any open data files are closed. This ensures that any changes to the data file are updated correctly.

There are two functions that you can use. They are **d4close** and **code4close**. The **d4close** function closes the data file whose **DATA4** pointer was provided as a parameter. The function also closes any index files (see the "Indexing" chapter) or memo files associated with the data file.

Code fragment  
from ShowData

```
d4close( db ) ;
```

The function **code4close** closes all open data, index, and memo files. This function takes a **CODE4** pointer as its only parameter. **code4initUndo** will also close any open data files.

---

## Moving Between Records

CodeBase provides ten functions for moving between records in the data file: **d4top**, **d4bottom**, **d4skip**, **d4go**, **d4seek**, **d4seekDouble**, **d4seekN**, **d4seekNext**, **d4seekNextDouble** and **d4seekNextN**. These functions change the current record by loading a new record into the record buffer.

- The **d4top** function sets the current record to the first record of the data file. The **d4bottom** performs a similar function for the last record.
- The **d4skip** function moves a specified number of records above or below the current record. This function must have a current record in order to function correctly.
- The **d4go** function loads the record buffer with the record whose record number was provided as a parameter.
- The **d4seek**, **d4seekDouble** and **d4seekN** functions locate the first record in a sorted order that matches a search key. These functions are described in detail in the "Indexing" chapter.
- The **d4seekNext**, **d4seekNextDouble** and **d4seekNextN** functions are similar to the **d4seek** functions except that they locate the next record in a sorted order that matches a search key. These functions are described in detail in the "Indexing" chapter.

---

## Scanning Through The Records

A common task in database programming is to display a list of records. The ShowData sample code illustrated this by using the number of records as the outer loop counter. Another way this same task can be accomplished is by using the return codes from the positioning functions. The following code fragment demonstrates this idea.

Scanning through  
records

```
...
rc := d4top( db ) ;
while rc = r4success do
begin
  For i = 1 To d4numFields( db ) do
  begin
    writeln( f4str( d4fieldJ( db, i ) ) ) ;
    i := i + 1 ;
  end;
  rc := d4skip( db, 1 ) ;
end;
...
```

**d4top** **d4top** is called, before the **While** loop, to position the data file to a valid record. This function causes the first record in the data file to be loaded into the record buffer. If the function is successful, it returns a value of **r4success**.

**d4skip** At the end of each iteration through the **While** loop, **d4skip** is called. **d4skip** allows you to skip either forwards or backwards through the data file. The numeric argument is used to specify how many records are skipped. If the number is positive, you move towards the end of the data file; if it is negative you move back towards the top. In this example the value is 1, which indicates that the next record should be loaded. The **d4skip** function also returns a value of **r4success** when it executes successfully.

The **While** loop is terminated when **rc** is no longer equal to **r4success**. This happens when either **d4skip** or **d4top** encounters an error or the end of the data file.

## Accessing Fields



### PROGRAM CustList

Before any information can be stored to or retrieved from a field, you must first obtain a pointer to its **FIELD4** structure. There are two ways to accomplish this: the first involves knowing the field's position in the data file, and the second requires that you already know the field's name.

This program lists the customer information contained in the CUSTLIST.DBF data file.

```
{ $I CODEBASE.INC }

program CUSTLIST;

uses CodeBase, WinCrt, Sysutils;

var
  code_base : CODE4;
  data_file : DATA4;
  f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4;
  rc, j, age_value : integer;
  amount_value : Double;
  name_str : array[0..24] of char;
  address_str : array[0..19] of char;
  date_str : array[0..8] of char;
  married_str : array[0..1] of char;
  comment_str : PChar;

begin
  code_base := code4init;

  data_file := d4open(code_base, 'DATA1.DBF');
  error4exitTest(code_base);

  f_name := d4field(data_file, 'F_NAME' );
  l_name := d4field(data_file, 'L_NAME' );
  address := d4field(data_file, 'ADDRESS' );
  age := d4field(data_file, 'AGE' );
  birth_date := d4field(data_file, 'BIRTH_DATE' );
  married := d4field(data_file, 'MARRIED' );
  amount := d4field(data_file, 'AMOUNT' );
```

```

comment      := d4field(data_file, 'COMMENT' );

rc := d4top(data_file);
while rc = r4success do
begin
  StrLCopy(name_str, f4str(f_name), SizeOf(name_str));
  StrLCopy(address_str, f4str(address), SizeOf(address_str));
  age_value := f4int(age);
  amount_value := f4double(amount);

  StrLCopy(date_str, f4str(birth_date), SizeOf(date_str));
  StrLCopy(married_str, f4str(married), SizeOf(married_str));
  comment_str := f4memoStr(comment);

  writeln('-----');
  writeln('Name      : ', name_str);
  writeln('Address   : ', address_str);
  writeln('Age       : ', age_value, ' Married : ', married_str);
  writeln('Comment    : ', comment_str);
  writeln('Purchased this year : ', amount_value:5:2);
  writeln('');

  rc := d4skip(data_file, Longint(1));
end;

d4close(data_file);
code4initUndo(code_base);

readKey;
DoneWinCrt;

Halt;
end.

```

## Referencing By Field Number

Each field in the data file has a unique *field number*. Field numbers range from one to the total number of fields in the data file. This value denotes the field's physical order in the record. The **d4fieldJ** function uses a field's field number to return a pointer to the field's **FIELD4** structure.

### Iterating through the fields

Iterating through the fields in a record is a common use of this method of field referencing. Before you can do this, **d4numFields** is called to determine the number of fields in the data file. A **For** loop is then used to iterate through each field to generically obtain a **FIELD4** pointer.

### Code fragment from ShowData

```

for j := 1 to num_fields do
begin
  field_ref := d4fieldJ( data_file, j );
  field_contents := f4memoStr( field_ref );
  write(field_contents, ' ');
end ;

```

## Referencing By Field Name

If you know ahead of time what the names of the fields are, you can obtain a field's **FIELD4** pointer using its name. This is the method used in the **CustList** program.

The **d4field** function returns a pointer to a **FIELD4** structure of the field whose name is specified. This is demonstrated in the following section of code.

Code fragment  
from CustList

```
var
f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4;

begin
... data file is opened ....

f_name      := d4field(data_file, 'F_NAME'   );
l_name      := d4field(data_file, 'L_NAME'   );
address     := d4field(data_file, 'ADDRESS'  );
age         := d4field(data_file, 'AGE'      );
birth_date  := d4field(data_file, 'BIRTH_DATE');
married     := d4field(data_file, 'MARRIED'  );
amount      := d4field(data_file, 'AMOUNT'   );
comment     := d4field(data_file, 'COMMENT'  );
```

In this code fragment, a **FIELD4** pointer is obtained and stored in separate variables for each field of the record.

When the data file is closed, all **FIELD4** pointers for the data file become invalid.

## Retrieving A Field's Contents

The Pascal language has many different data types such as **Integer**, **String**, **Char**, **Double** and **Longint**. As a result, the CodeBase provides a number of functions that retrieve the contents of fields and automatically perform any necessary conversions to the Pascal data type.

All the field functions accept a **FIELD4** pointer that specifies which field is to be operated upon.

## Retrieving field contents as Strings

If you need to access the field's contents in the form of a string, you can use **f4str** or **f4memoStr**. They both return a Pchar pointer to a copy of the field's contents.



### Note

Each time a **f4str** or **f4memoStr** function is called, the internal CodeBase buffer is overlaid with data from the new field. Consequently, if the field's value needs to be saved, it is necessary for the application to make a copy of the field's contents.

The **f4str** function can be used on any type of field except Memo fields. **f4memoStr** works on all types of fields including Memo fields. The reason for having two almost identical functions is that Memo fields require special handling. If you know a field is not a Memo field, it is appropriate to use the **f4str** function instead, since it is slightly faster.

To enable the program to work on any type of field, the ShowData program uses **f4memoStr** function.

Code fragment  
from ShowData

```
field_ref := d4fieldJ( data_file, j ) ;
field_contents := f4memoStr( field_ref ) ;
```

The CustList program uses **f4str** on non-Memo fields.

Code fragment  
from CustList

```
StrLCopy(date_str, f4str(birth_date), SizeOf(date_str));
StrLCopy(married_str, f4str(married), SizeOf(married_str));
```

Using **f4str** and  
**f4memoStr**

On Character fields

Character fields are returned by **f4str** and **f4memoStr** as left justified strings padded out with blanks to the full length of the field.

On Date fields	When these functions are used on Date fields, they return the date in the form of an eight character string. Please refer to the "Date Functions" chapter for more information.
On Numeric fields	The contents of Numeric fields are returned as right justified strings padded out to the full length of the field. If the field has any decimal places, the string will contain exactly that many decimal places.
On Logical fields	When used on Logical fields, these functions simply return a string containing "Y", "N", "y", "n", "T", "F", "t", or "f" (depending on the value stored on disk in the Logical field).
<b>Retrieving Numerical Data</b>	In addition to returning a field's contents as a string, CodeBase can also convert its contents to numerical data.
<b>f4double</b>	<b>f4double</b> returns the contents of the field as a <b>double</b> value. If <b>f4double</b> is unable to convert the field's contents, a value of <b>0</b> is returned.
<b>f4int and f4long</b>	The other two numeric operators, <b>f4int</b> and <b>f4long</b> convert the field's contents to <b>integer</b> or <b>longint</b> values. Any digits to the right of the decimal place (if any are stored in the field) are truncated.

## Creating Data Files

Not only does CodeBase allow you to access existing files, you can also create new files. This is accomplished by a single function call to **d4createData**.



### PROGRAM NewList

The NewList program creates a new data file if one of the same name does not already exist. It then appends several new records and assigns values to the fields in each record.

```
{ $I CODEBASE.INC }

program NEWLIST;

uses CodeBase, WinCrt, Sysutils;

const
  field_info : array[1..9] of FIELD4INFO = (
    (name:'F_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
    (name:'L_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
    (name:'ADDRESS'     ; atype:integer(r4str) ; len:15; dec:0),
    (name:'AGE'         ; atype:integer(r4num) ; len: 2; dec:0),
    (name:'BIRTH_DATE'  ; atype:integer(r4date) ; len: 8; dec:0),
    (name:'MARRIED'     ; atype:integer(r4log) ; len: 1; dec:0),
    (name:'AMOUNT'      ; atype:integer(r4num) ; len: 7; dec:2),
    (name:'COMMENT'     ; atype:integer(r4memo) ; len:10; dec:0),
    (name:nil           ; atype:0              ; len: 0; dec:0));

var
  code_base : CODE4;
  data_file : DATA4;
  f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4;

Procedure OpenDataFile;

begin
  data_file := d4open(code_base, 'DATA1.DBF');
  if data_file = nil then
    data_file := d4create(code_base, 'DATA1.DBF', @field_info, nil);

  f_name      := d4field(data_file, 'F_NAME'      );
  l_name      := d4field(data_file, 'L_NAME'      );
  address     := d4field(data_file, 'ADDRESS'     );
  age         := d4field(data_file, 'AGE'         );
  birth_date  := d4field(data_file, 'BIRTH_DATE' );
  married     := d4field(data_file, 'MARRIED'     );
  amount      := d4field(data_file, 'AMOUNT'      );
```

```

    comment    := d4field(data_file, 'COMMENT'    );
end;

Procedure PrintRecords;

var
    rc, j, age_value : integer;
    amount_value : Double;
    name_str      : array[0..24] of char;
    address_str   : array[0..19] of char;
    date_str      : array[0..8]  of char;
    married_str   : array[0..1]  of char;
    comment_str   : PChar;

begin
    rc := d4top(data_file);
    while rc = r4success do
        begin
            StrLCopy(name_str, f4str(f_name), SizeOf(name_str));
            StrLCopy(address_str, f4str(address), SizeOf(address_str));
            age_value := f4int(age);
            amount_value := f4double(amount);

            StrLCopy(date_str, f4str(birth_date), SizeOf(date_str));
            StrLCopy(married_str, f4str(married), SizeOf(married_str));
            comment_str := f4memoStr(comment);

            writeln('-----');
            writeln('Name      : ', name_str);
            writeln('Address   : ', address_str);
            writeln('Age       : ', age_value, ' Married : ', married_str);
            writeln('Comment  : ', comment_str);
            writeln('Purchased this year : ', amount_value:5:2);

            writeln('');

            rc := d4skip(data_file, Longint(1));
        end;
    end;

Procedure AddNewRecord(f_name_str, l_name_str, address_str : PChar;
    age_value, married_value : integer;
    amount_value : Double; comment_str : PChar);

begin
    d4appendStart(data_file, 0);

    f4assign(f_name, f_name_str );
    f4assign(l_name, l_name_str );
    f4assign(address, address_str);

    f4assignInt(age, age_value);

    if married_value = 1 then
        f4assign(married, 'T')
    else
        f4assign(married, 'F');

    f4assignDouble(amount, amount_value);
    f4memoAssign(comment, comment_str);

    d4append(data_file);
end;

begin
    code_base := code4init;

    code4errOpen(code_base, 0 );
    code4safety(code_base, 0 );

    OpenDataFile;

    PrintRecords;

    AddNewRecord('Sarah', 'Webber', '132-43 St.', 32, 1, 147.99,
        'New Customer');
    AddNewRecord('John', 'Albridge', '1232-76 Ave.', 12, 0, 98.99, nil);

    PrintRecords;

    code4close(code_base);

```

```
code4initUndo(code_base);

readKey;
DoneWinCrt;

Halt;
end.
```

## FIELD4INFO Structures

The **FIELD4INFO** structure is an integral component of the data file creation process. This structure contains the information which defines a field.

The **FIELD4INFO** structure contains a member for each of the field's four attributes. These members are described below:

- **name** This is a string containing the name of the field. This name must be unique to the data file and any characters after the first ten are ignored. Valid field names can only contain letters, numbers, and/or underscores. The first character of the name must be a letter.
- **type** This member contains an abbreviation for the field type. CodeBase supports Character, Date, Floating Point, Logical, Memo, Numeric, Binary, and General field types. Valid abbreviations are either the character constants 'C', 'D', 'F', 'L', 'M', 'N', 'B' and 'G'. In addition you can also use the equivalent CodeBase constants: **r4str**, **r4date**, **r4num**, **r4log**, **r4memo**, **r4num**, **r4bin** and **r4gen**, respectively.
- **len** This **integer** value determines the length of the field.
- **dec** This **integer** member determines the number of decimal places in Numeric or Floating Point fields.

Refer to **d4create** in the CodeBase *Reference Guide* for more detailed information on field attributes.

### FIELD4INFO arrays

Before a data file can be created, an array of **FIELD4INFO** structures must be defined. Each element of this array describes a single field.

An example **FIELD4INFO** array is defined below:

#### Code fragment from NewList

```
field_info : array[1..9] of FIELD4INFO = (
  (name:'F_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
  (name:'L_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
  (name:'ADDRESS'     ; atype:integer(r4str) ; len:15; dec:0),
  (name:'AGE'         ; atype:integer(r4num) ; len: 2; dec:0),
  (name:'BIRTH_DATE'  ; atype:integer(r4date); len: 8; dec:0),
  (name:'MARRIED'     ; atype:integer(r4log) ; len: 1; dec:0),
  (name:'AMOUNT'      ; atype:integer(r4num) ; len: 7; dec:2),
  (name:'COMMENT'     ; atype:integer(r4memo); len:10; dec:0),
  (name:nil           ; atype:0              ; len: 0; dec:0));
```



## Creating The Data File

The actual creation of the data file is performed by the **d4create** function. This function uses the information in a **FIELD4INFO** array as the field specifications for a new data file. **d4createData** only creates a data file and possibly a memo file. The function **d4create** builds a data file and production index file as well. The "Indexing" chapter discusses the use of **d4create** and building indexes.

The second parameter of the **d4createData** function is a string containing a file name. This file name can contain any extension, but if no extension is provided, the default ".DBF" extension is used. Like the **d4open** function, the current directory is assumed if no path is provided as part of the file name.

If the data file is successfully created, a pointer to its **DATA4** structure is returned. If for any reason the data file could not be created, zero is returned instead and **CODE4.errorCode** is set accordingly. The **DATA4** pointer that is returned from **d4createData** or **d4create** can be passed on to other data functions such as **d4top** and **d4skip**.

The **CODE4.safety** flag

The **CODE4.safety** flag determines whether any existing file is replaced when CodeBase attempts to create a file. If the file exists and the **CODE4.safety** flag is set to its default value of true (non-zero), the new file is not created. When this flag is set to false (zero), the old file and its index and memo files are replaced by the newly created file and its associated files. If the new file is not created because the file already exists and **CODE4.errCreate** is true, an error message is generated. The **CODE4.safety** flag can be changed by calling **code4safety**.



**WARNING**

As with all **CODE4** members, the **CODE4.safety** flag can only be changed after **code4init** has been called.

## Opening Or Creating

It is often the case that you want to open a data file if it exists or create it if it does not. The following section of code demonstrates how this can be accomplished:

Code fragment  
from NewList

```
save1 = code4errOpen( cb, 0 )
save2 = code4safety( cb, 0 )

. . .

db = d4open( cb, fPath + "DATA1" )

If db = 0 Then
  InitField4
  db = d4createData( cb, "DATA1", fieldInfo( ) )
  If cbError( ) Then Exit Function
End If
```

The  
**CODE4.errOpen**  
flag

If the data file does not already exist, **d4open** will normally generate an error message when it fails to open the file. This error message can be suppressed by changing the **CODE4.errOpen** flag. When this flag is set to true (its default value), an error message is generated when CodeBase is unable to open a file. If this flag is set to false (zero), CodeBase does not display this error message. It does however still return a zero **DATA4** pointer, which can be tested. This is the recommended method for opening a data file if it exists and creating it if it does not. The **CODE4.errOpen** can be changed by calling **code4errOpen**.

## Adding New Records

There are two methods that can be used to add new records to the data file. They both append a new record to the end of the data file.

### Appending A Blank Record

If all you need to do is add a blank record to the bottom of the data file, **d4appendBlank** should be used.

```
rc := d4appendBlank( db ) ;
```

This function adds a new record to the data file, blanks out all of its fields, and makes the new record the current record.

### The Append Sequence

The second method involves three steps. This method is more efficient and should be used when you intend to immediately assign values to the new record's fields. Using this method results in less disk writes and therefore improves performance.

starting the  
append sequence

The first step is performed by the **d4appendStart** function. This function sets the current record number to zero to let CodeBase know that a new record is about to be appended. Additionally, **d4appendStart** temporarily disables automatic flushing for the current record, so that if changes need to be aborted, the record is not flushed to disk.

Code fragment  
from NewList

```
d4appendStart( data_file, 0 ) ;
```

The second step involves assigning values to the fields. After the **d4appendStart** function is called, the record buffer contains a copy of the previous current record. If no changes are made to the record buffer, a copy of this record is appended. If you prefer an empty record buffer, a call to **d4blank** will clear it.

```
d4appendStart( data_file, 0 ) ;  
d4blank( db ) ;
```

Assigning values to the fields will be discussed in the next section.



### Note

The first byte of the record buffer is reserved for the 'deleted' flag. This flag is also brought forward when you call **d4appendStart**. If you are not calling **d4blank**, call **d4recall** to ensure that the current record is not also marked as 'deleted'.

Appending the  
record

The final step is accomplished by **d4append**. This function causes a new record to be physically added to the end of the data file and its key values to be added to any open index tags.

Code fragment  
from NewList

```
d4append( data_file );
```

## Assigning Field Values

### CODE4.lockEnforce

Just as there are several CodeBase functions for retrieving information from fields, there are corresponding functions for storing information in fields.

In multi-user configurations, only one application should edit a record at a time. One way to ensure that only one application can modify a record is to explicitly lock the record before it is changed. To ensure that a record is locked before it is changed, call **code4lockEnforce** to set **CODE4.lockEnforce** to true (non-zero). When **CODE4.lockEnforce** is true (non-zero) and an attempt is made to modify an unlocked record using a field function or **d4blank**, **d4changed**, **d4delete** or **d4recall**, an **e4lock** error is generated and the modification is aborted.

An alternative method of ensuring that only one application modifies a record is to deny all other applications write access to the data file. In this case explicit locking is not required, even when **CODE4.lockEnforce** is true, since only no outside applications can write to the data file. Write access can be denied to other applications by passing **OPEN4DENY\_WRITE** or **OPEN4DENY\_RW** to **code4accessMode** before the data file is opened.

Refer to the "Multi-User Applications" chapter of this guide for more information.

## Copying Strings To Fields

Copying strings to  
Character fields

**f4assign** and **f4memoAssign** are two functions that copy strings to fields. The **f4assign** function works on any type of field except Memo fields, while **f4memoAssign** works on all types.

If the length of the assigned string is less than the size of the field, the field's contents are left justified, and the remaining spaces are filled with blanks. If the string is larger than the field, its contents are truncated to fit. When an assignment is done with a memo field, the length is automatically adjusted to accommodate any value.

Code fragment  
NewList

```
f4assign(f_name, f_name_str );  
...  
f4memoAssign(comment, comment_str);
```

Copying  
strings to Numeric  
fields

**f4assign** and **f4memoAssign** can also store strings into Numeric fields, but generally this is not a good idea as it is very easy to format the information incorrectly. A correctly formatted number is right justified, zero-padded right to the number of decimals in the field, and space-padded left for any unused units.



### Note

The functions **f4assign** and **f4memoAssign** store the information exactly as it appears in the string. This can cause problems when they are used to store non-numeric data into Numeric type fields. It is recommended that the functions **f4assignInt**, **f4assignDouble** or **f4assignLong** be used to store data in Numeric fields.

Copying strings to Date fields Any strings that are copied to Date type fields should be exactly eight characters long and should contain a date in standard format (CCYYMMDD). For details please refer to the "Date Functions" chapter.

Copying strings to Logical fields The only strings that should be copied to Logical fields are "T", "F", "t", "f", "Y", "N", "y" or "n".

Code fragment  
from NewList

```
if married_value = 1 then
  f4assign(married, 'T')
else
  f4assign(married, 'F');
```

## Storing Numeric Data

Instead of worrying about the formatting of strings containing numerical data, you can let CodeBase perform the formatting for you. **f4assignDouble**, **f4assignInt** and **f4assignLong** automatically convert and format numerical data.

Assigning **double**  
values to a field

You can assign a **double** value to a field using **f4assignDouble**. This function can automatically format the numerical value according to the field's attributes. For example, if the field has a length of six and has two decimal places, and the **double** value of 34.12345 assigned to it, the field will contain " 34.12" after **f4assignDouble** is called. If the **f4assignDouble** function is used on a Character field, it assumes the field has no decimals and right justifies its contents.

Code fragment  
from NewList

```
f4assignDouble(amount, amount_value);
```

Assigning integer  
values to a field

**f4assignInt** and **f4assignLong** are used to copy and format **integer** and **longint** values to fields. Right justification is used, and if the field is of type Numeric, any decimal places are filled with zeros.

```
f4assignLong( age, age_value );
```

## Removing Records

CodeBase provides a two level method for removing records. You can delete a record by simply changing the status of its deletion flag to true. The records that are flagged for deletion can then be physically removed by "packing" the data file.



### PROGRAM Deletion

The Deletion program demonstrates the effects of deleting, recalling, and packing records.

```
{ $I CODEBASE.INC }

program DELETION ;

uses CodeBase, WinCrt, Sysutils ;

const
  field_info : array[1..3] of FIELD4INFO = (
    (name:'DUMMY' ; atype:integer('C') ; len:10 ; dec:0),
    (name:'MEMO' ; atype:integer('M') ; len:10 ; dec:0),
    (name:nil ; atype:0 ; len: 0 ; dec:0) ) ;

var
  code_base : CODE4 ;
  data_file : DATA4 ;
  count : integer ;

Procedure print_delete_status( status : integer ; rec_no : Longint ) ;

begin
```

```

if status <> 0 then
    writeln( 'Record ', rec_no, ' - DELETED' )
else
    writeln( 'Record ', rec_no, ' - NOT DELETED' ) ;
end ;

Procedure print_records( data_file : DATA4 ) ;

var
    rc, status : integer ;
    rec_no : Longint ;

begin
    writeln( ' ' ) ;

    rc := d4top( data_file ) ;
    while rc <> r4eof do
        begin
            rec_no := d4recno( data_file ) ;
            status := d4deleted( data_file ) ;
            print_delete_status( status, rec_no ) ;
            rc := d4skip( data_file, Longint(1) ) ;
        end ;
    end ;

begin
    code_base := code4init ;
    data_file := d4create( code_base, 'TUTOR5', @field_info, nil ) ;
    error4exitTest( code_base ) ;

    for count := 0 to 3 do
        d4appendBlank( data_file ) ;

    print_records( data_file ) ;

    d4go( data_file, Longint(3) ) ;
    d4delete( data_file ) ;
    d4go( data_file, Longint(1) ) ;
    d4delete( data_file ) ;
    print_records( data_file ) ;

    d4go( data_file, Longint(3) ) ;
    d4recall( data_file ) ;
    print_records( data_file ) ;

    d4pack( data_file ) ;
    d4memoCompress( data_file ) ;
    print_records( data_file ) ;

    code4close( code_base ) ;
    code4initUndo( code_base ) ;

    readKey ;
    DoneWinCrt ;

    Halt ;
end.

```

**Explanation:** This program creates the data file (or re-creates it, if it exists) and then appends five records. The program then displays the deletion status of all the records:

Deletion output  
part 1

```

Record #    1 - NOT DELETED
Record #    2 - NOT DELETED
Record #    3 - NOT DELETED
Record #    4 - NOT DELETED

```

Record numbers one and three are then marked for deletion and the deletion status of the records are again displayed.

Deletion output  
part 2

```
Record # 1 - DELETED
Record # 2 - NOT DELETED
Record # 3 - DELETED
Record # 4 - NOT DELETED
```

Record number three is recalled (ie. the deletion mark is removed) and the deletion status of the records are displayed.

Deletion output  
part 3

```
Record # 1 - DELETED
Record # 2 - NOT DELETED
Record # 3 - NOT DELETED
Record # 4 - NOT DELETED
```

Finally the data file is packed. Displaying the record status shows that record one was removed from the data file. Record two now occupies the space of record one.

Deletion output  
part 4

```
Record # 1 - NOT DELETED
Record # 2 - NOT DELETED
Record # 3 - NOT DELETED
```



### Note

Since the record numbers are contiguous and always start at one, the record number for a specific record may have changed after the data file has been packed

## Determining The Deletion Status

Each record has its own deletion flag. The status of the deletion flag of the current record can be checked using **d4deleted**.

Code fragment  
from Deletion

```
status := d4deleted( data_file ) ;
```

**d4deleted** returns a true (non zero) value when the current record has been marked for deletion.

## Marking A Record For Deletion

**d4delete** is used to set the current record's deletion status to true. The next time the data file is packed, the record is physically removed from the file and any reference to it in any open index file is also removed.

Code fragment  
from Deletion

```
d4delete( data_file ) ;
```



### Note

When a record is marked for deletion, it is not physically removed from the data file. In fact you can still access the record normally. The main importance of the deletion flag is that the marked record is physically removed when the data file is packed. The deletion flag may be used in tag filter expressions. Refer to the "Indexing" chapter for details.

## Recalling Records

If the data file has not yet been packed, any deleted records can be recalled back to non-deleted status. This is accomplished by a call to **d4recall**.

```
d4recall( data_file ) ;
```

## Packing The Data File

Physically removing records from the data file is called packing. If your data file is quite large, this process can be quite time consuming. That is why records are first marked for deletion and then physically removed. The cumulative time necessary for removing individual records would be substantially greater than removing the many "deleted" records at one time. Packing the data file is performed by **d4pack**.



### Note

Consider opening the file exclusively before packing to keep others from accessing the data file while packing is occurring. Otherwise, the data may appear as corrupted to other users for the duration of the pack.

## Compressing The Memo File

When **d4pack** removes a record from the data file, it does not remove the corresponding memo entry (assuming that there is at least one Memo field). This results in unreferenced memo entries. Unreferenced memo entries cause no difficulties in an application except that they simply waste disk space. To remove any wasted memo file space, the memo file can be compressed by calling **d4memoCompress**. Again, since compressing a memo file can consume a large amount of time, it may be appropriate to compress a data file sometime other than when the data file is packed



### Note

Wasted memo file space is mainly caused by packing records that have memo entries without doing a memo compress. If you are using Clipper compatibility, wasted memo space can also occur when memo entries are increased beyond 504 bytes. In either case, memo compressing reduces disk space usage.

## Data File Information



### PROGRAM DataInfo

CodeBase contains a set of functions that provide information about the data files, their records and fields.

This program displays information about a data file whose name is provided in an **Input Box** when the program is run. It displays the data file's alias, record count, record width, and the attributes of its fields.

```
program DATAINFO ;

uses CodeBase, Sysutils, WinCrt ;

var
  code_base : CODE4 ;
  data_file : DATA4 ;
  field_ref : FIELD4 ;

  j, num_fields, len, dec, rec_width : integer ;
  name, alias : PChar ;
  atype : integer ;
  rec_count : Longint ;
  dataname : array[0..255] of char ;

begin

  code_base := code4init ;
  code4autoOpen( code_base, 0 ) ;
  writeln( 'DataInfo - displays the structure of a database.' ) ;
  write( 'Enter database name: ' ) ;
  readln( dataname ) ;

  data_file := d4open( code_base, dataname ) ;
  error4exitTest( code_base ) ;
```

```

rec_count := d4reccount ( data_file ) ;
num_fields := d4numFields ( data_file ) ;
rec_width := d4recWidth( data_file ) ;
alias      := d4alias      ( data_file ) ;

writeln( '-----' ) ;
writeln( '| Data File: ', dataname:12, ' |' ) ;
writeln( '| Alias      : ', alias:12, ' |' ) ;
writeln( '|-----|' ) ;
writeln( '| Number of Records: ', rec_count:7, ' |' ) ;
writeln( '| Length of Record : ', rec_width:7, ' |' ) ;
writeln( '| Number of Fields : ', num_fields:7, ' |' ) ;
writeln( '|-----|' ) ;
writeln( '| Field Information : |' ) ;
writeln( '|-----|' ) ;
writeln( '| Name      | type | len | dec |' ) ;
writeln( '|-----|' ) ;

for j := 1 to d4numFields( data_file ) do
begin
    field_ref := d4fieldJ( data_file, j ) ;
    name := f4name( field_ref ) ;
    atype := f4type( field_ref ) ;
    len := f4len( field_ref ) ;
    dec := f4decimals( field_ref ) ;

    writeln( '|', name:10, '|', atype, '|', len:4, '|', dec:4, '|' ) ;
end ;
writeln( '-----' ) ;

d4close( data_file ) ;
code4initUndo( code_base ) ;

ReadKey ;
DoneWinCrt ;

Halt ;
end.

```

The data  
file alias

Each data file that you have open has a character string label called the *alias*. The alias is mainly used for qualifying field names in dBASE expressions or looking up a **DATA4** pointer using **code4data**.

**d4alias** returns a string containing the data file's alias. When the file is opened, the alias is set to the same name as the data file without the extension. This string is valid as long as the data file is open.

Code fragment for  
DataInfo

```

alias := d4alias( data_file ) ;
...
writeln( '| Alias      : ', alias:12, ' |' ) ;

```

The alias of a data file can be changed by using **d4aliasSet**. This would be required in a situation where the same file is being opened twice or if two data files in different directories have the same file name. In either case, open the first data file and change its alias to something other than the file name and then open the second data file. In this way CodeBase can distinguish the data files based on their differing aliases.

Finding the number  
of records

An important data file statistic is the record count, which can be obtained by calling **d4recCount**. This function does not take into account the deleted status of any records, nor any filter condition of an open tag.

```

rec_count := d4recCount( data_file ) ;

```



Determining the record width

Another useful function is **d4recWidth**. This function returns the length of the record buffer. This value is the total length of all the fields plus one byte for the deletion flag. This function may be used to save copies of the record buffer.

Code fragment for DataInfo

```
rec_width := d4recWidth( data_file );
Print "Record Length: "; d4recWidth( db )
```

## Field Information

If you have access to a field's **FIELD4** pointer, you can retrieve the field attributes.

Field name

**f4name** returns a pointer to the field's name. This pointer is valid as long as the data file is open.

Code fragments from DataInfo

```
name := f4name( field_ref );
```

Field type

**f4type** returns the ASCII value of the field type:

```
atype := f4type( field_ref );
```

Field length and decimals

The length of the field and the number of decimals are returned by **f4len** (**f4memoLen**) and **f4decimals** respectively. Both functions return **integers**.

```
len := f4len( field_ref );
dec := f4decimals( field_ref );
```

## Advanced Topics

This section documents how to copy the structure of an existing file and use the information to create a new file with the same structure.

### Copying Data File Structures

You may occasionally require a new data file that has an identical structure to an existing data file.

To satisfy this need CodeBase provides the function **d4fieldInfo**, which returns a copy of the **FIELD4INFO** array of an existing file. NOTE that this function returns a pointer to the field information that is already in the native 'C' format of the DLL.

You can also refer to the discussion of **i4tagInfo** in the "Indexing" chapter. This function returns an array of index information, which can be used to create index files.



### Note

The **FIELD4INFO** array pointer that is returned by **d4fieldInfo** is allocated dynamically and should therefore be deallocated after you are finished using it. The **u4free** function can be used for this purpose.



### PROGRAM CopyData

The CopyData program retrieves a file name from a Input Box and uses that file's structure to create a new empty file.

```
program COPYDATA ;

uses CodeBase, WinCrt, Sysutils;

var
  code_base : CODE4 ;
  data_file, data_copy : DATA4 ;
  field_info : PFIELD4INFO ;
  dataname1 : array[0..255] of char ;
  dataname2 : array[0..255] of char ;
```

## 34 CodeBase

```
begin
  if ParamCount <> 2 then
    begin
      writeln( ' USAGE: COPYDATA <FROM FILE> <TO FILE>' ) ;
      Halt ;
    end ;

    code_base := code4init ;
    code4safety( code_base, 0 ) ;
    code4autoOpen( code_base, 0 ) ;

    StrPCopy( dataname1, ParamStr(1) ) ;
    data_file := d4open( code_base, dataname1 ) ;
    error4exitTest( code_base ) ;

    field_info := d4fieldInfo( data_file ) ;

    StrPCopy( dataname2, ParamStr(2) ) ;
    data_copy := d4create( code_base, dataname2, field_info, nil ) ;

    u4free( PVoid(field_info) ) ;

    code4close( code_base ) ;
    code4initUndo( code_base ) ;

    readKey ;
    DoneWinCrt ;

    Halt ;
end.
```

## 3 Indexing

The purpose of a database is to organize information in a useful manner. So far you have seen how the information is divided into fields and records. Now it is time to order the records in purposeful ways. One approach is to physically sort records in your data files. Unfortunately, maintaining data file sorts involves continually shuffling large amounts of data. In addition, it is only possible to maintain a single sort order using this method.

A preferable method is to leave the records in the data file in their original order and store the sorted orderings in a separate file called an index file. When you create an index file, you are effectively sorting the data file. Index files are efficiently maintained and you can have an unlimited number of sorted orderings continually available.

### Indexes & Tags

Each index file can contain one or more sorted orderings. These sorted orderings are identified by a *tag*. That is, each index file tag corresponds to a single sorted ordering. CodeBase supports four types of index files. Their attributes and differences are described below:

File Format	Compatibility	Number of Tags	Production Indexes	Group Files	Index Filtering	Descending Ordering
.MDX	dBase IV	1 - 47	yes	no	yes	yes
.CDX	FoxPro 2.x FoxPro 3.0	No Limit	yes	no	yes	yes
.NTX	Clipper 87 Clipper 5	1	Only with group files	yes	no	Special Case

Figure 3.1 Index File Formats

The difference between indexes and tags are illustrated in Figure 3.2.

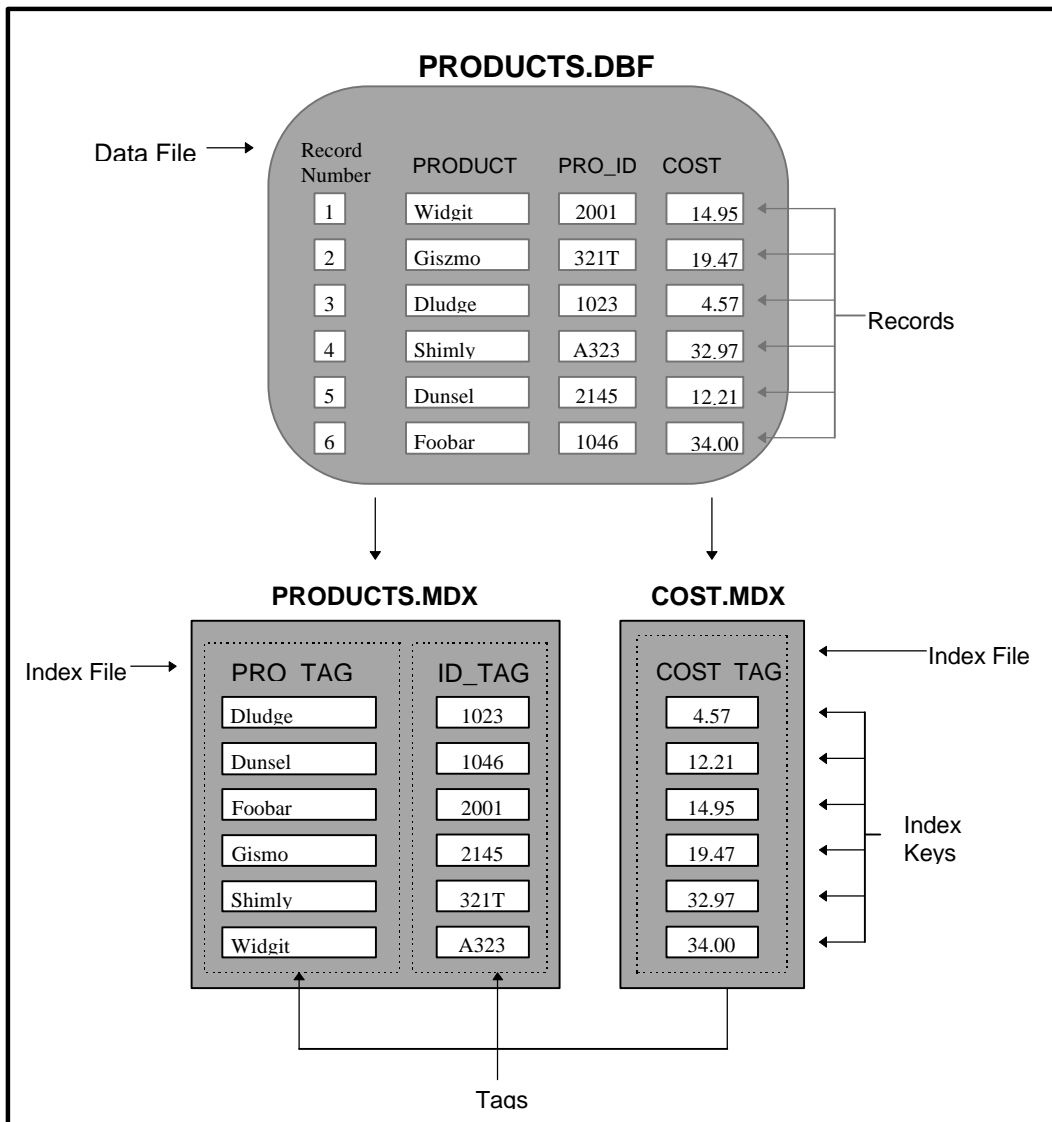


Figure 3.2 Index Files

## Index Expressions

An *index expression* is a *dBASE expression* that is used to determine the *index key* for each record. An index expression must evaluate to a value of one of the following types: Numeric, Character, or Date types. When you are using **.CDX** (FoxPro) indexes, you can also index on Logical expressions. Refer to "Appendix C: dBASE Expressions" in the CodeBase Reference Guide for details on dBASE expressions.

The most commonly used index expression is a field name. For example, the index expression for the PRO\_TAG tag in Figure 3.2 is "PRODUCT". When this expression is evaluated for record number 6, the value of the field PRODUCT is returned; for record number 6, this value is "Foobar ". To put it simply, tag PRO\_TAG is ordered by the contents in field PRODUCT.

Other common index expressions involve generating tags based on two or more fields. This is known as a *compound index key*. For example, if we assume that field PRO\_ID is also a character field, we can base a tag ordering with the following index expression: "PRODUCT + PRO\_ID". This produces an index key consisting of the PRODUCT field concatenated with the PRO\_ID field. For record number 6, the resulting index key is "Foobar 1046".

In addition, most *dBASE functions* are also permitted. See "Appendix C: dBASE Expressions" in the CodeBase Reference Guide for a list of dBASE functions supported. If we wanted to base a tag on the PRODUCT field, but remain case insensitive, we can use the dBASE function UPPER() to convert the index key to upper case. In this case the index expression would be "UPPER(PRODUCT)".

## Creating An Index File



### PROGRAM NewList2

Creating an index file is similar to creating a data file; in fact you can even create both at the same time. There are two types of index files that you can create: *production indexes* and *non-production indexes*.

The following is the program listing for NewList2. This modified version of the NewList program, creates an index file along with the data file.

```
program NEWLIST2;

uses CodeBase, WinCrt, Sysutils;

const
  field_info : array[1..9] of FIELD4INFO = (
    (name:'F_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
    (name:'L_NAME'      ; atype:integer(r4str) ; len:10; dec:0),
    (name:'ADDRESS'     ; atype:integer(r4str) ; len:15; dec:0),
    (name:'AGE'         ; atype:integer(r4num) ; len: 2; dec:0),
    (name:'BIRTH_DATE'  ; atype:integer(r4date); len: 8; dec:0),
    (name:'MARRIED'     ; atype:integer(r4log) ; len: 1; dec:0),
    (name:'AMOUNT'      ; atype:integer(r4num) ; len: 7; dec:2),
    (name:'COMMENT'     ; atype:integer(r4memo); len:10; dec:0),
    (name:nil           ; atype:0              ; len: 0; dec:0));

  tag_info : array[1..6] of TAG4INFO = (
    (name:'NAME_TAG' ; expression:'F_NAME+L_NAME' ; filter:'' ; unique:0;
     descending:0),
    (name:'ADDR_TAG' ; expression:'ADDRESS'       ; filter:'' ; unique:0;
     descending:0),
    (name:'AGE_TAG'  ; expression:'AGE'           ; filter:'' ; unique:0;
     descending:0),
    (name:'DATE_TAG' ; expression:'BIRTH_DATE'    ; filter:'' ; unique:0;
     descending:0),
    (name:'AMNT_TAG' ; expression:'AMOUNT'        ; filter:'' ; unique:0;
     descending:0),
```

```

        (name:nil          ; expression:nil          ; filter:nil; unique:0;
         descending:0));

var
  code_base : CODE4;
  data_file : DATA4;
  f_name, l_name, address,
    age, birth_date, married, amount, comment : FIELD4;
  name_tag, age_tag, amount_tag : TAG4 ;

Procedure CreateDataFile;

begin
  data_file := d4create(code_base, 'DATA1.DBF', @field_info, @tag_info);

  f_name      := d4field(data_file, 'F_NAME'      );
  l_name      := d4field(data_file, 'L_NAME'      );
  address     := d4field(data_file, 'ADDRESS'     );
  age         := d4field(data_file, 'AGE'         );
  birth_date  := d4field(data_file, 'BIRTH_DATE'  );
  married     := d4field(data_file, 'MARRIED'     );
  amount      := d4field(data_file, 'AMOUNT'      );
  comment     := d4field(data_file, 'COMMENT'     );

  name_tag    := d4tag(data_file, 'NAME_TAG');
  age_tag     := d4tag(data_file, 'AGE_TAG' );
  amount_tag  := d4tag(data_file, 'AMNT_TAG');
end;

Procedure PrintRecords;

var
  rc, j, age_value : integer;
  amount_value : Double;
  f_name_str : array[0..14] of char;
  l_name_str : array[0..14] of char;
  address_str : array[0..19] of char;
  date_str : array[0..8] of char;
  married_str : array[0..1] of char;
  comment_str : PChar;

begin
  rc := d4top(data_file);
  while rc = r4success do
    begin
      StrLCopy( f_name_str, f4str( f_name ), SizeOf( f_name_str ) );
      StrLCopy( l_name_str, f4str( l_name ), SizeOf( l_name_str ) );
      StrLCopy( address_str, f4str( address ), SizeOf( address_str ) );
      age_value := f4int(age);
      amount_value := f4double(amount);
      StrLCopy( date_str, f4str( birth_date ), SizeOf( date_str ) );
      StrLCopy( married_str, f4str( married ), SizeOf( married_str ) );
      comment_str := f4memoStr(comment);

      writeln('-----');

      writeln('Name      : ', f_name_str:10, ' ', l_name_str:10);
      writeln('Address   : ', address_str:15);
      writeln('Age       : ', age_value:3, ' Married : ', married_str:1);
      writeln('Comment  : ', comment_str);
      writeln('Amount purchased this year : ', amount_value:5:2);

      rc := d4skip(data_file, Longint(1));
    end;
  end;

Procedure AddNewRecord(f_name_str, l_name_str, address_str : PChar;
  age_value, married_value : integer;
  amount_value : Double; comment_str : PChar);

begin
  d4appendStart(data_file, 0);

  f4assign(f_name, f_name_str);
  f4assign(l_name, l_name_str);
  f4assign(address, address_str);
  f4assignInt(age, age_value);
  if married_value <> 0 then
    f4assign(married, 'T')
  else
    f4assign(married, 'F');

```

```

f4assignDouble(amount, amount_value);
f4memoAssign(comment, comment_str);

d4append(data_file);
end;

begin
code_base := code4init;
code4safety(code_base, 0);
CreateDataFile;

AddNewRecord('Sarah', 'Webber', '132-43 St.', 32, 1, 147.99,
'New Customer');
AddNewRecord('John', 'Albridge', '1232-76 Ave.', 12, 0, 98.99, nil);

PrintRecords;

d4tagSelect(data_file, name_tag);
PrintRecords;

d4tagSelect(data_file, age_tag);
PrintRecords;

d4tagSelect(data_file, amount_tag);
PrintRecords;

code4close(code_base);
code4initUndo(code_base);

readKey;
DoneWinCrt;

Halt;
end.

```

<b>Production Indexes</b>	A production index is an index that is opened automatically when its associated data file is opened. A data file can have only one production index and this index has the same name as the data file, but with a different extension.
<b>Non-Production Indexes</b>	All other index files are non-production indexes. A data file can have an unlimited number of non-production indexes. To use these indexes, they must explicitly be opened after the data file has been opened.
<b>The TAG4INFO Structure</b>	Just as the attributes of a field in the data file are defined by a <b>FIELD4INFO</b> structure, the attributes of a tag in an index file are specified by a <b>TAG4INFO</b> structure. This structure contains the following five members:
The <b>TAG4INFO</b> members	<p>The first two members must be defined for every <b>TAG4INFO</b> structure. The last three members are used to specify optional properties of the tag which will be discussed later in this chapter.</p> <ul style="list-style-type: none"> <li>• <b>name</b> This is a <b>string</b> containing the name of the tag. The name may be composed of letters, numbers and underscores. Only letters and underscores are permitted as the first character of the name. This member cannot be nil.</li> <li>• When using FoxPro or <b>.MDX</b> formats, the name must be unique to the data file and have a length of ten characters or less excluding the extension. If you are using the <b>.NTX</b> index format, then this name can include a tag name with a path. In this case, the tag name within the path is limited to eight characters or less, excluding the extension.</li> </ul>

- **expression** This is a **string** containing the tag's index expression. This expression determines the sorting order of the tag. (see the Index Expression section of this chapter for details) This member cannot be nil.
- **filter** This is a **string** containing the tag's filter expression. If this member is set to a blank **string** or nil, there is no filter for the tag.
- **unique** This **integer** determines how duplicate keys are treated. See the Unique Keys section of this chapter for more details. If this member is set to 0, the tag is non-unique.
- **descending** This **integer** determines if the index should be generated in descending order. Set this member to **0** to specify ascending order, or **r4descending** if descending order is desired.

Before the index file can be created, an array of **TAG4INFO** structures must be defined. Each element of this array denotes a single tag.



### Note

Each member of the last element in a **TAG4INFO** array must be set to null. This indicates that there are no more elements in the array. Failure to provide an element filled with null's will result in errors when the index file is created.

The **TAG4INFO**  
array

Before the index file can be created, an array of **TAG4INFO** structures must be defined. Each element of this array denotes a single tag.

An example **TAG4INFO** array is defined below.

Code fragment  
from NewList2

```
const
  tag_info : array[1..6] of TAG4INFO = (
    (name:'NAME_TAG'; expression:'F_NAME+L_NAME';
     filter:'' ; unique:0; descending:0),
    (name:'ADDR_TAG'; expression:'ADDRESS';
     filter:'' ; unique:0; descending:0),
    (name:'AGE_TAG' ; expression:'AGE';
     filter:'' ; unique:0; descending:0),
    (name:'DATE_TAG'; expression:'BIRTH_DATE';
     filter:'' ; unique:0; descending:0),
    (name:'AMNT_TAG'; expression:'AMOUNT';
     filter:'' ; unique:0; descending:0),
    (name:nil ; expression:nil;
     filter:nil; unique:0; descending:0));
```

### Using d4create To Create Index Files

The most common method for creating an index file is to create it when the data file is created. In this case the index will be a production index. If you recall, the **d4create** accepts four parameters, only three of which were used in the last chapter. The fourth parameter is a pointer to a **TAG4INFO** array which is used for creating a production index file. This is illustrated in the following code segment:

Code fragment  
from NewList2

```
data_file := d4create(code_base, 'DATA1.DBF',@field_info, @tag_info);
```



## Using i4create To Create Index Files

The other method for creating index files involves the use of the **i4create** function. In addition to a **TAG4INFO** array, this function also accepts a string containing a file name. An index file is created with this file name. If no extension is specified, an appropriate extension for the file compatibility used ( **.MDX**, **.CDX**, or **.NTX** ) is provided.

For example, the program NewList2 could be modified to use **i4create** instead.

```
var
  ind: INDEX4;
...
db := d4createData( cb, 'DATA1', @field_info );
ind := i4create( db, 'DATAINDEX', @tag_info );
```

**i4create** returns a pointer to the index file's **INDEX4** structure. The pointer is used by many low level index functions.

### Creating Production Indexes with i4create

Normally, **i4create** does not produce production indexes, even if the file name that you provide is the same as the data file's. **i4create** can be used to create production indexes in the following manner. The data file must be opened exclusively before **i4create** is called. This can be achieved by setting the **CODE4.accessMode** to **OPEN4DENY\_RW** before the data file is opened. After the data file is opened or created exclusively, then **i4create** is used to create a production index by passing nil as the file name parameter. This is illustrated as follows:

```
var
  ind: INDEX4;
...
code4accessMode( cb, OPEN4DENY_RW );
db := d4create( cb, 'DATA1', @fldInfo, nil );
...
ind := i4create( db, nil, @tagInfo );
```

## Maintaining Index Files

CodeBase automatically updates all open index files for the data file. When a record is added, modified, or deleted, the appropriate tag entries for all of the open index tags are modified automatically. In general application programming, there is no need to manually add, delete, or modify the tag entries. CodeBase handles all of the key manipulation in the background, letting you concentrate on application programming.

### Code fragment

```
var
  cb: CODE4;
  db: DATA4;
  nameFld: FIELD4;
begin
  cb := code4init ;
  db := d4open( cb, 'DATAFILE' ) ;
  nameFld := d4field( db, 'NAME' ) ;

  d4top( db ) ;
  f4assign( nameFld, 'ROBERT WILKHAM' ) ;

  d4skip( db, 1 ) ;

{ the changes are flushed to disk. If "NAME" is a
key field, the index tag entry is automatically
updated. A seek for "ROBERT WILKHAM" would succeed }
```

## Opening Index Files

Opening index files can be accomplished by one of two functions. If the data file has a production index, **d4open** automatically opens it. Other index files may be opened using function **i4open**.

```
var
  cb: CODE4;
  db: DATA4;
  ind: INDEX4;
begin
  cb := code4init ;
  db := d4open( cb, 'DATAFILE' ) ;

  ind := i4open( db, 'INDEXFILE' );
...
```

Disabling the automatic opening of production indexes

There may be occasions when you prefer to leave the production index file closed. You can disable the automatic opening of production index files by setting flag **CODE4.autoOpen** to false (zero). The production index can then be opened like any other index file:

```
...
cb := code4init ;
db := d4open( cb, 'DATAFILE' ) ;

ind := i4open( db, 'INDEXFILE' );
ind2 := i4open( db, 'DATAFILE' ) ;
...
```

## Referencing Tags

Before a tag is used, you must first obtain a pointer to its **TAG4** structure. When calling tag functions, this pointer specifies which tag the function operates on. There are six functions that can be used to retrieve a tag's **TAG4** pointer: **d4tag**, **d4tagDefault**, **d4tagNext**, **d4tagPrev**, **d4tagSelected** and **t4open**.



### PROGRAM ShowData2

The program ShowData2 is a modified version of ShowData. This version displays the records of the data file in natural order and according to any tags in its production index.

```
program SHOWDAT2 ;

uses CodeBase, WinCrt, Sysutils ;

var
  code_base : CODE4 ;
  data_file : DATA4 ;
  field_ref : FIELD4 ;
  tag       : TAG4 ;
  rc, j : integer ;
  field_contents : PChar ;
  dataname : array[0..255] of char ;

Procedure print_records ;

begin
  rc := d4top( data_file ) ;
  while rc = r4success do
    begin
      for j := 1 to d4numFields( data_file ) do
        begin
          field_ref := d4fieldJ( data_file, j ) ;
          field_contents := f4memoStr( field_ref ) ;
          write( field_contents, ' ' ) ;
        end ;
        writeln( ' ' ) ;
        rc := d4skip( data_file, Longint(1) ) ;
      end ;
    end ;
  end ;

begin
  if ParamCount <> 1 then
```

```

begin
    writeln( ' USAGE: SHOWDAT2 <FILENAME.DBF>' ) ;
    Halt ;
end ;

code_base := code4init ;

{$IFDEF N4OTHER}
    code4autoOpen( code_base, False ) ;
{$ENDIF}

StrPCopy( dataname, ParamStr(1) ) ;
data_file := d4open( code_base, dataname ) ;
error4exitTest( code_base ) ;

writeln( ' ' ) ;
writeln( 'Data File ', dataname, ' in Natural Order' ) ;
writeln( ' ' ) ;
print_records ;

tag := d4tagNext( data_file, nil ) ;
while tag <> nil do
begin
    writeln( ' ' ) ;
    writeln( 'Press ENTER to continue:' ) ;
    readln ;

    writeln( 'Data File ', dataname, ' sorted by Tag ', t4alias( tag ) ) ;
    writeln( ' ' ) ;
    d4tagSelect( data_file, tag ) ;
    print_records ;
    tag := d4tagNext( data_file, tag ) ;
end ;

d4close( data_file ) ;
code4initUndo( code_base ) ;

readKey ;
DoneWinCrt ;

Halt ;
end.

```

## Referencing By Name

If you know the tag's name, you can retrieve its **TAG4** pointer by using **d4tag**. This function looks through all the open index files of the specified data file for a tag matching the name. This method is used by program NewList2.

Code Fragment  
from NewList2

```

name_tag := d4tag(data_file, 'NAME_TAG');
age_tag  := d4tag(data_file, 'AGE_TAG' );
amount_tag := d4tag(data_file, 'AMNT_TAG');

```

## Obtaining the Default Tag

The default tag is the currently selected tag. If there are no tags selected, the default tag is the first tag in the first opened index file. You can obtain a pointer to the default tag's **TAG4** structure by calling function **d4tagDefault**.

```

var
    tag1: TAG4 ;
    db: DATA4;
...
db := d4open( cb, 'DATAFILE' ) ;
tag1 := d4tagDefault( db ) ;

```

## Iterating Through The Tags

The functions **d4tagNext** and **d4tagPrev** are used for iterating through the tags.

**d4tagNext** The function **d4tagNext** returns the next tag after the specified tag. The first tag of the first index is returned by passing a nil pointer as the tag parameter. If the last tag in the index file is passed into the tag parameter, the first tag in the next index file is returned.

Code fragment  
from ShowData2

```
tag := d4tagNext( data_file, nil ) ;
while tag <> nil do
begin
    ...
    tag := d4tagNext( data_file, tag ) ;
end ;
tagPtr = d4tagNext( db, 0& )
```

**d4tagPrev** The function **d4tagPrev** works like **d4tagNext** except that it starts at the last tag of the last index file and works its way toward the first tag of the first index file.

---

## Selecting Tags

Initially a data file does not have a selected tag and is therefore in *natural order*. Natural order is the order in which the records were added to the data file. When you want to use a particular sort ordering, select a tag by calling **d4tagSelect**.

Code fragment  
from NewList2

```
d4tagSelect(data_file, name_tag);
PrintRecords;
```

Selecting Natural  
Order

If you want to return to natural order simply pass nil to **d4tagSelect** instead of a **TAG4** pointer.

Obtaining the  
currently selected  
tag

The function **d4tagSelected** obtains the currently selected tag. It behaves like **d4tagDefault** except that it returns zero when no tag is selected.

---

## The Effects Of Selecting a Tag

Once a tag is selected, the behaviour of **d4top**, **d4bottom**, and **d4skip** changes. Functions **d4top** and **d4bottom** then set the current record to the first and last data file entries respectively, using the tag's sort ordering. Similarly, **d4skip** skips using the tag ordering instead of the natural ordering.

---

## Tag Filters

A tag filter is used to obtain a subset of the available data file records. The subset is based on true/false conditions calculated from a dBASE expression. Only records that pass through the filter have entries in the tag. A tag filter is created when the tag is created and cannot later be changed without destroying and recreating the tag.



### PROGRAM ShowList2

In program ShowList2, when the index file is created, several filters are present.

```
program SHOWLIST2 ;

uses CodeBase, WinCrt, Sysutils ;

const
    tag_info : array[1..6] of TAG4INFO = (
        (name:'NAME'      ; expression:'L_NAME+F_NAME' ; filter:'.NOT. DELETED()' ;
          unique:0 ; descending:0),
        (name:'ADDRESS'   ; expression:'ADDRESS'      ; filter:''      ; unique:0 ;
          descending:0),
        (name:'AGE_TAG'   ; expression:'AGE'          ; filter:'AGE >= 18' ;
          unique:0 ; descending:0),
        (name:'DATE_TAG'  ; expression:'BIRTH_DATE'   ; filter:''      ; unique:0 ;
          descending:0),
        (name:'AMNT_TAG'  ; expression:'AMOUNT'       ; filter:''      ; unique:0 ;
          descending:0),
        (name:nil         ; expression:nil            ; filter:nil     ; unique:0 ;
```

```

                                descending:0) ) ;

var
  code_base : CODE4 ;
  data_file : DATA4 ;
  f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4 ;
  name_tag, age_tag, amount_tag, address_tag, birthdate_tag : TAG4 ;

Procedure OpenDataFile ;

begin
  code4autoOpen( code_base, 0 ) ;
  code4safety( code_base, 0 ) ;
  code4accessMode( code_base, OPEN4DENY_RW ) ;
  data_file := d4open( code_base, 'DATA1.DBF' ) ;

  i4create( data_file, nil, @tag_info ) ;

  f_name      := d4field( data_file, 'F_NAME' ) ;
  l_name      := d4field( data_file, 'L_NAME' ) ;
  address     := d4field( data_file, 'ADDRESS' ) ;
  age        := d4field( data_file, 'AGE' ) ;
  birth_date  := d4field( data_file, 'BIRTH_DATE' ) ;
  married     := d4field( data_file, 'MARRIED' ) ;
  amount      := d4field( data_file, 'AMOUNT' ) ;
  comment     := d4field( data_file, 'COMMENT' ) ;

  name_tag    := d4tag( data_file, 'NAME' ) ;
  address_tag := d4tag( data_file, 'ADDRESS' ) ;
  age_tag     := d4tag( data_file, 'AGE_TAG' ) ;
  birthdate_tag := d4tag( data_file, 'DATE_TAG' ) ;
  amount_tag  := d4tag( data_file, 'AMNT_TAG' ) ;
end ;

Procedure PrintRecords ;

var
  rc,j,age_value : integer ;
  amount_value : Double ;
  f_name_str : array[0..14] of char ;
  l_name_str : array[0..14] of char ;
  address_str : array[0..19] of char ;
  date_str : array[0..8] of char ;
  married_str : array[0..1] of char ;
  comment_str : PChar ;

begin
  rc := d4top( data_file ) ;
  while rc = r4success do
    begin
      StrLCopy( f_name_str, f4str( f_name ), SizeOf( f_name_str ) ) ;
      StrLCopy( l_name_str, f4str( l_name ), SizeOf( l_name_str ) ) ;
      StrLCopy( address_str, f4str( address ), SizeOf( address_str ) ) ;
      age_value := f4int( age ) ;
      amount_value := f4double( amount ) ;
      StrLCopy( date_str, f4str( birth_date ), SizeOf( date_str ) ) ;
      StrLCopy( married_str, f4str( married ), SizeOf( married_str ) ) ;
      comment_str := f4memoStr( comment ) ;

      writeln( '-----' ) ;
      writeln( 'Name      : ', f_name_str:10, ' ', l_name_str:10 ) ;
      writeln( 'Address   : ', address_str:15 ) ;
      writeln( 'Age       : ', age_value:3, ' Married : ', married_str:1 ) ;
      writeln( 'Comment  : ', comment_str ) ;
      writeln( 'Amount purchased : ', amount_value:5:2 ) ;
      writeln( ' ' ) ;

      rc := d4skip( data_file, Longint(1) ) ;
    end ;
  end ;

begin
  code_base := code4init ;

  OpenDataFile ;

  d4tagSelect( data_file, name_tag ) ;
  PrintRecords ;

  d4tagSelect( data_file, age_tag ) ;
  PrintRecords ;

```

```

d4tagSelect( data_file, amount_tag ) ;
PrintRecords ;

code4close( code_base);
code4initUndo( code_base ) ;

readKey ;
DoneWinCrt ;

Halt ;
end.

```

## Filter Expressions

A *filter expression* is a dBASE expression that returns a Logical result and is used as a tag filter. The expression is evaluated for each record as its tag entry is updated. If the filter expression evaluates to true, an entry for that record is included in the tag. If it evaluates to false, that record's tag entry is omitted from the tag.

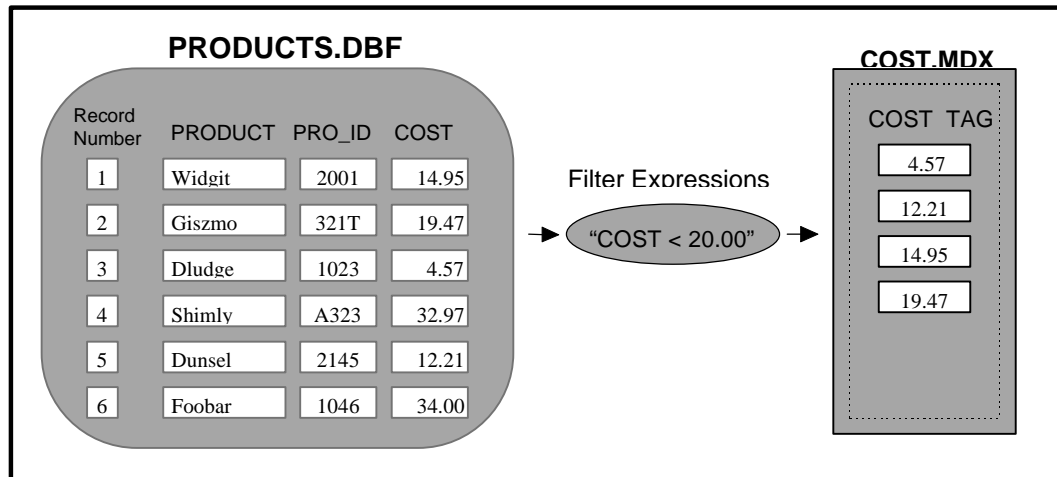


Figure 3.3 Filter Expressions

## Creating A Tag Filter

The tag filter is specified, when the index is initially created, through the **TAG4INFO.filter** structure member. If this member is a zero length string or nil, then no records are filtered. The following code segment shows several tags with filters:

Code Fragment  
from ShowList2

```

tag_info : array[1..6] of TAG4INFO = (
  (name:'NAME' ; expression:'L_NAME+F_NAME' ;
   filter:'.NOT. DELETED()' ; unique:0 ; descending:0),
  (name:'ADDRESS' ; expression:'ADDRESS' ;
   filter:'' ; unique:0 ; descending:0),
  (name:'AGE_TAG' ; expression:'AGE' ;
   filter:'AGE >= 18' ; unique:0 ; descending:0),
  (name:'DATE_TAG' ; expression:'BIRTH_DATE' ;
   filter:'' ; unique:0 ; descending:0),
  (name:'AMNT_TAG' ; expression:'AMOUNT' ;
   filter:'' ; unique:0 ; descending:0),
  (name:nil ; expression:nil ;
   filter:nil ; unique:0 ; descending:0) ) ;

```

The first tag has a filter expression of ".NOT. DELETED()". This filters out any records that have been marked for deletion. The third tag's filter expression is "AGE > 18". This excludes any record which has an AGE field value of less than or equal to 18.

## Unique Keys

The fourth member of the **TAG4INFO** structure, **unique**, determines how duplicate keys are handled. This member can have four possible values: **0**, **r4uniqueContinue**, **e4unique** and **r4unique**.

It is often desirable to ensure that no two index keys are the same. To meet this requirement, you can specify that a tag must only contain unique keys. As a result, when a new key is to be added to the tag, a search is first made to see if that key already exists in the tag. If it does exist, the new entry is not added. CodeBase provides three methods of handling the addition of non-unique keys: **r4uniqueContinue**, **e4unique** and **r4unique**.

The descriptions of the values for the **TAG4INFO.unique** member are as follows:

- **0** Duplicate keys are allowed.
- **r4uniqueContinue** Any duplicate keys are discarded. However, the record continues to be added to the data file. In this case, there may not be a tag entry for a particular record.
- **e4unique** An error message is generated if a duplicate key is encountered.

**r4unique** The data file record is not added or changed. Regardless, no error message is generated. Instead a value of **r4unique** is returned.

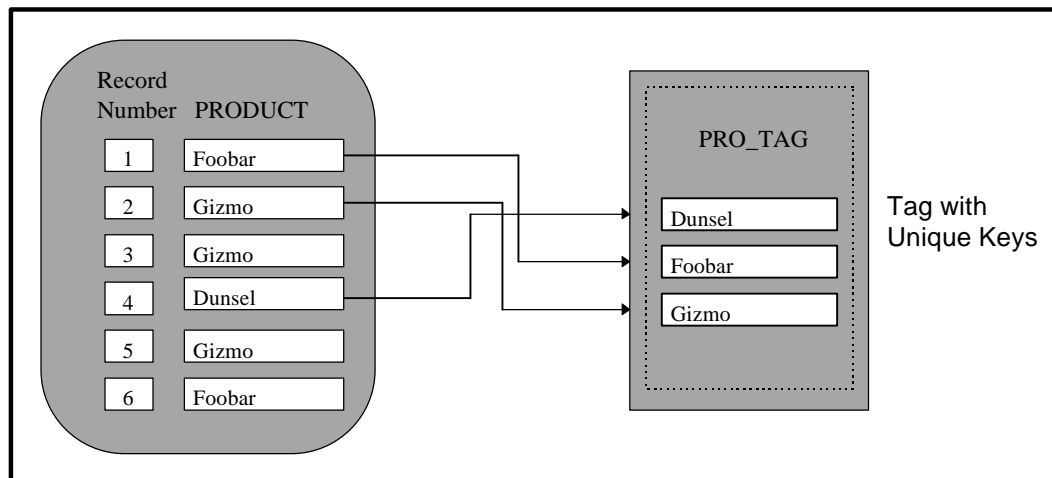


Figure 3.4 Unique Keys

Here is an example **TAG4INFO** array that uses unique keys:

```
tag_info : array[1..2] of TAG4INFO = (
  (name:'NAME'      ; expression:'L_NAME+F_NAME' ;
   filter:''       ; unique: 1 ; descending:0),
  (name:nil        ; expression:nil             ;
   filter:nil      ; unique: 0 ; descending:0) ) ;
```



### WARNING

Although a specific unique code for a tag is specified when the tag is created, only a true/false flag, which indicates whether a tag is unique or non-unique, is actually stored on disk. No information on how to respond to a duplicate key for unique key tags is saved.

As a result, when a tag is re-opened, the value contained in **CODE4.errDefaultUnique** is used to set the tag's unique code. If the tag was designed to be **r4unique** or **e4unique**, then the default value of **r4uniqueContinue** must be changed before any records are modified or appended.

This change can be accomplished in two ways. The first involves setting the **CODE4.errDefaultUnique** to another value, in which case all indexes opened subsequently will have this value for their unique code. The function that sets this unique code is **code4errDefaultUnique**. This setting has no effect on non-unique tags.

The second method is to set the unique code for one or more tags individually after the index file has been opened. Pass the appropriate unique code to the function **t4uniqueSet**. This function only has an effect on unique tags.



## Note

**t4uniqueSet** can only be used to change the setting of a unique tag. A CodeBase error is generated when an attempt is made to set a unique tag to a non-unique tag or a non-unique tag to a unique tag.

## Seeking

One of the most useful features of a database is the ability to find a record by providing a *search key*. This process is referred to as *seeking*. When a seek is performed, the search key, which is usually a string, is compared against the index keys in the selected tag. When a match occurs, the corresponding record is loaded in the record buffer.



### PROGRAM Seeker

The Seeker program demonstrates seeks on the various types of tags. When it performs the seeks, it displays the record that was found and the return value that was returned from the seek.

```
program SEEKER ;

uses CodeBase, WinCrt, Sysutils ;

var
  rc : integer ;
  code_base : CODE4 ;
  data_file : DATA4 ;
  f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4 ;
  name_tag, age_tag, amount_tag, address_tag, birthdate_tag : TAG4 ;

Procedure OpenDataFile ;

begin
  data_file := d4open( code_base, 'DATA1.DBF' ) ;

  f_name      := d4field( data_file, 'F_NAME' ) ;
  l_name      := d4field( data_file, 'L_NAME' ) ;
  address     := d4field( data_file, 'ADDRESS' ) ;
  age         := d4field( data_file, 'AGE' ) ;
  birth_date  := d4field( data_file, 'BIRTH_DATE' ) ;
  married     := d4field( data_file, 'MARRIED' ) ;
  amount      := d4field( data_file, 'AMOUNT' ) ;
  comment     := d4field( data_file, 'COMMENT' ) ;

  name_tag    := d4tag( data_file, 'NAME_TAG' ) ;
  address_tag := d4tag( data_file, 'ADDR_TAG' ) ;
  age_tag     := d4tag( data_file, 'AGE_TAG' ) ;
  birthdate_tag := d4tag( data_file, 'DATE_TAG' ) ;
```



```

    amount_tag := d4tag( data_file, 'AMNT_TAG' ) ;
end ;

Procedure seek_status( rc : integer ; status : PChar ) ;

begin
    case rc of
        r4success : StrLCopy(status, 'r4success', SizeOf( status )) ;
        r4eof      : StrLCopy(status, 'r4eof',      SizeOf( status )) ;
        r4after    : StrLCopy(status, 'r4after',    SizeOf( status )) ;
    else
        StrLCopy( status, 'other', SizeOf( status )) ;
    end ;
end ;

Procedure print_record( rc : integer ) ;

var
    age_value      : integer ;
    amount_value   : Double ;
    f_name_str     : array[0..14] of char ;
    l_name_str     : array[0..14] of char ;
    address_str    : array[0..19] of char ;
    date_str       : array[0..8] of char ;
    married_str    : array[0..1] of char ;
    comment_str    : PChar ;
    status         : array[0..14] of char ;

begin
    StrLCopy( f_name_str, f4str( f_name ), SizeOf( f_name_str )) ;
    StrLCopy( l_name_str, f4str( l_name ), SizeOf( l_name_str )) ;
    StrLCopy( address_str, f4str( address ), SizeOf( address_str )) ;
    age_value := f4int( age ) ;

    amount_value := f4double( amount ) ;

    StrLCopy( date_str, f4str( birth_date ), SizeOf( date_str )) ;
    StrLCopy( married_str, f4str( married ), SizeOf( married_str )) ;

    comment_str := f4memoStr( comment ) ;

    seek_status( rc, status ) ;

    writeln( 'Seek status ', status ) ;
    writeln( '-----' ) ;
    writeln( 'Name      : ', f_name_str, ' ', l_name_str ) ;
    writeln( 'Address   : ', address_str ) ;
    writeln( 'Age: ', age_value, ' BirthDate: ', date_str, ' Married : ', married_str ) ;
    writeln( 'Comment: ', comment_str ) ;
    writeln( 'Amount purchased : ', amount_value ) ;
    writeln( ' ' ) ;
end ;

begin
    code_base := code4init ;

    OpenDataFile ;

    d4tagSelect( data_file, address_tag ) ;

    rc := d4seek( data_file, '123 - 45 Ave ' ) ;
    print_record( rc ) ;

    rc := d4seek( data_file, '12' ) ;
    print_record( rc ) ;

    rc := d4seek( data_file, '12 ' ) ;
    print_record( rc ) ;

    d4tagSelect( data_file, name_tag ) ;

    rc := d4seek( data_file, 'Krammer David ' ) ;
    print_record( rc ) ;

    d4tagSelect( data_file, birthdate_tag ) ;

    rc := d4seek( data_file, '19500101' ) ;
    print_record( rc ) ;

    d4tagSelect( data_file, amount_tag ) ;

```

```

rc := d4seekDouble( data_file, 47.23 ) ;
print_record( rc ) ;

rc := d4seek( data_file, ' 47.23' ) ;
print_record( rc ) ;

code4close( code_base ) ;
code4initUndo( code_base ) ;

readKey ;
DoneWinCrt ;

Halt ;
end.

```

## Performing Seeks

There are six functions in the CodeBase library that perform seeking: **d4seek** and **d4seekNext** may seek for character data, **d4seekDouble** and **d4seekNextDouble** may seek for numeric data and **d4seekN** and **d4seekNextN** may seek for binary data.

**d4seek** with character data

**d4seek** accepts a string as the search key and performs a seek using the default tag, which is the selected tag if one is selected. If there is no selected tag, the default tag is the first tag of the first opened index. **d4seek** can perform seeks using Character, Numeric or Date tags.

Code fragment from Seeker

```

d4tagSelect( data_file, address_tag ) ;

rc := d4seek( data_file, '123 - 45 Ave ' ) ;
print_record( rc ) ;

```

**d4seek** on Date tags

When performing seeks using Date tags, **d4seek** must be passed an eight character string containing a date in the standard "CCYYMMDD" format. See chapter 7 for details on using the date functions.

Code fragment from SEEKER

```

d4tagSelect( data_file, birthdate_tag ) ;

rc := d4seek( data_file, '19500101' ) ;
print_record( rc ) ;

```

**d4seek** on Numeric tags

When **d4seek** uses a Numeric tag and the search key is a string, the string is converted into **double** value before searching. As a result, you do not have to worry about formatting the character parameter to match the key value.

Code fragment from SEEKER

```

d4tagSelect( data_file, amount_tag ) ;

rc := d4seek( data_file, ' 47.23' ) ;
print_record( rc ) ;

```

**d4seekDouble** on Numeric Tags

**d4seekDouble** accepts a **double** value as the search key and this function should only be used on numeric keys. All numeric tags, regardless of the number of decimals, store their key entries as **double** values. When seeking on numeric tags, it is more efficient to seek with a **double** value than a character representation.

Code fragment from SEEKER

```

d4tagSelect( data_file, amount_tag ) ;

rc := d4seekDouble( data_file, 47.23 ) ;
print_record( rc ) ;

```

**d4seekN** on Binary Data Use the **d4seekN** function to seek on character tags that contain binary data. This seek function can be used to seek without regard for nulls, since *len* specifies the length of the search key.

Code fragment

```
var
  key: array[0..20] of Char;
...
  StrCopy( key, '0000111010 ');
  d4tagSelect( db, binaryTag );
  rc := d4seekN( db, key, StrLen( key ) );
```

**d4seekNext**,  
**d4seekNextDouble**  
and **d4seekNextN**

The **d4seekNext**, **d4seekNextDouble** and **d4seekNextN** functions only differ from their **d4seek** counterparts in that the searches begin at the current position in the tag. This provides the capability of performing an incremental search through the data base. Consequently, all index keys matching the search key can be found by repeated calls to **d4seekNext** functions. On the other hand, **d4seek** functions always begin their searches from the first logical record as determined by the selected tag and therefore, they only locate the first index key in the tag that matches the search key.

All three **d4seekNext** functions work in the following manner. If the index key at the current position in the tag does not match the search key, **d4seekNext** calls **d4seek** to find the first occurrence of the search key. Conversely, if the index key at the current position does match the search key, **d4seekNext** tries to find the next occurrence of the search key. Similarly, **d4seekNextDouble** and **d4seekNextN** call **d4seekDouble** and **d4seekNextN** respectively to find the first occurrence of the search key.

```
d4tagSelect( db,name_tag )
rc = d4seekNext( db, "John      Albridge  " )

while rc = r4success do
begin
  PrintRecordSeek( rc )
  rc := d4seekNext( db, "John      Albridge  " )
Loop
```

**d4seek**, **d4seekDouble**, **d4seekN**, **d4seekNext**, **d4seekNextDouble** and **d4seekNextN** return an integer value indicating the success or failure in locating the key value in the selected tag.

## Exact Matches

CodeBase gives you the ability to perform both exact and partial matches on tags. An exact match occurs when the index key in the tag is the same as the search key. When searching using a character tag, the criteria that all **d4seek** and **d4seekNext** functions use for determining an exact match is as follows:

- The characters in the search string must be the same as corresponding characters in the index key.
- The comparison is case sensitive. Therefore a search key of "aaaa" does not match an index key of "AAAA".

To do a case insensitive search, one must change the indexing strategy, not the search method. Create a tag with the index expression "UPPER(fieldName)" and then a search for "AAAA" will be case insensitive. Therefore, both "Aaaa" and "aAaA" will be found.

- The strings are only compared up to end of the shortest value. Therefore a search key of "abcd" exactly matches an index key of "abcdefg" because only four characters are compared.



### Note

You can force CodeBase to compare additional characters by padding the search key out to the full length of the index key. Therefore a search key of "abcd " does not exactly match "abcdefg"

When the **d4seek** or **d4seekNext** functions find an exact match, a value of **r4success** is returned.

## Partial Matches

When a seek is performed, the tag is positioned to the place in the tag where the search key should be located. If this position lies between two index keys (ie. the tag was not exactly found), then a partial match has occurred. The record whose index key is directly after the "correct" position is loaded into the record buffer. When this happens, a value of **r4after** is returned by the **d4seek** functions and **r4entry** or **r4after** is returned by the **d4seekNext** functions. Consider the following three seeks.

Code fragment  
from SEEKER

```
d4tagSelect( data_file, address_tag ) ;

rc := d4seek( data_file, '123 - 45 Ave ' ) ;
print_record( rc ) ;

rc := d4seek( data_file, '12' ) ;
print_record( rc ) ;

rc := d4seek( data_file, '12          ' ) ;
print_record( rc ) ;
```

Assuming that there is a record in the data file containing "'123 - 45 Ave " in the compound tag NAME\_TAG: The first time **d4seek** is called, it returns **r4success** because its search key exactly matched and was the same length as the index key. The second call will also return **r4success** because it only compares the first three characters. The third call however, returns **r4after** because it could not find an exact match. All three, however, result in the "'123 - 45 Ave " record being loaded into the record buffer.

possible returns  
from **d4seekNext**  
functions

When **d4seekNext** is called and the index key at the current position in the tag does not match that of the search key, **d4seek** is called to find the first instance of the search key. Should this seek fail, **r4after** is returned and tag is positioned as discussed above.

On the other hand, if **d4seekNext** is called and index key at the current position does match the search key, **d4seekNext** searches for the next occurrence of the search key. If this seek fails, **r4entry** is returned and the tag is positioned in the same manner as **r4after**.

**d4seekNextDouble** and **d4seekNextN** function in the same manner as **d4seekNext**.

### No Match

If there are no index keys in the tag or if the search key's position is after the last index key, then the **d4seek** and **d4seekNext** functions return **r4eof**.

### Seeking On Compound Keys

Performing seeks on compound index keys is slightly more difficult. Before a seek can be performed on a tag with a compound index expression, a search key must be created in a similar manner. For example, if the index expression for the tag is "F\_NAME + L\_NAME", an example is as follows:

```
d4tagSelect( db, nameTag ) ;  
  
rc := d4seek( db, `Sarah      Webber      ` ) ;  
PrintRecordSeek( rc ) ;
```

The first part of the search key, "Sarah ", matches the format of field F\_NAME and is padded out with blanks to the length of that field ( 10 ). The second part, "Webber ", matches field L\_NAME, which also has a length of 10. The total size of the string is 20 characters.



### WARNING

The various parts of the search key must be padded out to the full length of their respective fields. Failure to do so may cause an incorrect seek.

A similar strategy is used when seeking on compound index keys built with fields of different types. If the index expression is "PRODUCT + STR(COST, 7, 2)" where 'product' is a Character field of length eight and Cost is a Numeric field, then a valid search key is "Dunsel 1234.21". Since the STR() function converts Numeric values to Character values, the search string is constructed with the first eight characters matching the PRODUCT field and the "1234.21" matching the return value of STR().



### WARNING

You should note the difference between the dBASE "+" and "-" operators when they are applied to Character values. The "+" includes any trailing blanks when the Character values are concatenated. For example if the index expression is "L\_NAME + F\_NAME", the search key should be "Krammer David ". The "-" removes any trailing blanks from the first string. The second string is concatenated and the previously removed blanks are added to the end of the concatenated string. As a result, if the index expression is "L\_NAME - F\_NAME", the search key should be "KrammerDavid ".

---

## Group Files

In Clipper, the **.NTX** indexes must be manually opened each and every time the data file is opened. This causes programming time to be increased, as well as having a good chance of forgetting to open an index file. dBASE IV and FoxPro use compound index files (more than one tag in a file) and production index files (automatically opened when data file opens) to avoid these problems.

CodeBase has introduced group files in order to compensate for this limitation of the **.NTX** file format. A group file allows you to use the same function calls when using **.NTX** index files as you would when using **.CDX** and **.MDX** index files.

This is accomplished by emulating production indexes and multiple tags per index file. The same database that was shown in Figure 3.2 is depicted in Figure 3.5 using **.NTX** index files and group files.

---

## Creating Group Files

A group file is automatically created by **i4create** or **d4create** when using **.NTX** index compatibility. In addition, a group file can be created for existing index files using a text editor. Simply create a file with a file name extension of **".CGP"**. Then you enter the names of the index files to be grouped together. Enter one index file name per line.



GROUP FILE  
PRODUCTS.CGP

For example, here is the actual contents of group file 'PRODUCTS.CGP' from Figure 3.5:

```
PRO_TAG  
P_NUM_TAG
```

---

## Bypassing Group Files

Because group files are unique to CodeBase, index files generated with Clipper do not have group files. Under these conditions, it may be more convenient for you to bypass the group files and access the **.NTX** group files directly instead of creating your own group files.

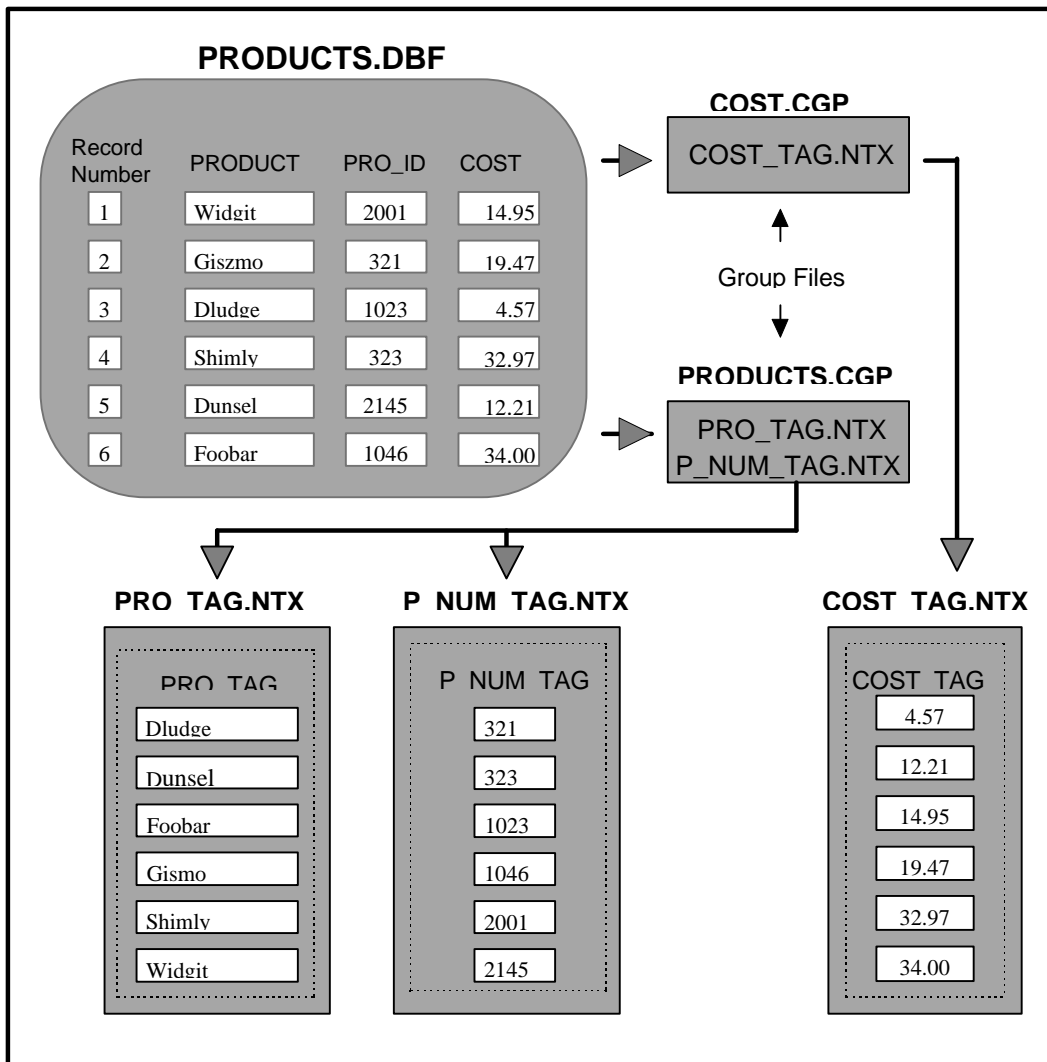


Figure 3.5 Group Files

### Disabling The Auto Open

When using **.NTX** index compatibility, CodeBase automatically tries to open a group file when a data file is opened. If you have tried to open a data file that does not have a group file, you will encounter the following error:

Error #: -60  
Opening File  
File Name:  
DATA1.CGP

This results from CodeBase being unable to locate a group file. This behavior can be disabled by setting the **CODE4.autoOpen** flag to false (zero).

## Creating Index Files

You can create **.NTX** files without a corresponding group file by passing a nil as the file name to **i4create**. As documented earlier, the tag names also become file names.



### PROGRAM NoGroup

This program demonstrates how **.NTX** index files can be created without group files.

```
program NOGROUP1 ;

uses CodeBase, WinCrt, Sysutils ;

const
  tag_info : array[1..5] of TAG4INFO = (
    (name:'NAME' ; expression:'L_NAME+F_NAME' ; filter:'' ; unique:0 ;
     descending:0),
    (name:'ADDRESS' ; expression:'ADDRESS' ; filter:'' ; unique:0 ;
     descending:0),
    (name:'AGE_TAG' ; expression:'AGE' ; filter:'' ; unique:0 ;
     descending:0),
    (name:'DATE' ; expression:'BIRTH_DATE' ; filter:'' ; unique:0 ;
     descending:0),
    (name:nil ; expression:nil ; filter:nil ; unique:0 ;
     descending:0) ) ;

var
  code_base : CODE4 ;
  data_file : DATA4 ;
  name_tag, address_tag, age_tag, date_tag : TAG4 ;

begin
  {$IFDEF N4OTHER}
    writeln( 'CLIPPER FORMAT ONLY' ) ;
  {$ELSE}
    code_base := code4init ;

    code4autoOpen( code_base, False ) ;
    code4safety( code_base, False ) ;

    data_file := d4open( code_base, 'DATA1.DBF' ) ;

    i4create( data_file, nil, @tag_info ) ;
    name_tag := d4tag( data_file, 'NAME' ) ;
    address_tag := d4tag( data_file, 'ADDRESS' ) ;
    age_tag := d4tag( data_file, 'AGE_TAG' ) ;
    date_tag := d4tag( data_file, 'DATE' ) ;

    d4closeAll( code_base ) ;
    code4initUndo( code_base ) ;
  {$ENDIF}

  readKey ;
  DoneWinCrt ;
end.
```

## Opening Index Files

A single **.NTX** index file can also be opened without using group files.



### WARNING

The following example program contains code segments that are specific to CodeBase libraries built with **.NTX** index file compatibility.



### PROGRAM NoGroup2

This program demonstrates how **.NTX** index files can be opened without using a group file.

```
program NOGROUP2;

uses CodeBase, WinCrt, Sysutils;

var
  code_base : CODE4;
```



```

data_file : DATA4;
f_name, l_name, address, age, birth_date, married, amount, comment : FIELD4;
name_tag, address_tag, age_tag, date_tag : TAG4 ;

Procedure print_records;

var
  rc, j, age_value : integer;
  amount_value : Double;
  f_name_str : array[0..14] of char;
  l_name_str : array[0..14] of char;
  address_str : array[0..19] of char;
  date_str : array[0..8] of char;
  married_str : array[0..1] of char;
  comment_str : PChar;

begin
  rc := d4top(data_file);
  while rc = r4success do
    begin
      StrLCopy( f_name_str, f4str( f_name ), SizeOf( f_name_str ) );
      StrLCopy( l_name_str, f4str( l_name ), SizeOf( l_name_str ) );
      StrLCopy( address_str, f4str( address ), SizeOf( address_str ) );
      age_value := f4int(age);
      amount_value := f4double(amount);
      StrLCopy( date_str, f4str( birth_date ), SizeOf( date_str ) );
      StrLCopy( married_str, f4str( married ), SizeOf( married_str ) );
      comment_str := f4memoStr(comment);
      writeln('-----');
      writeln('Name      : ', f_name_str:10, ' ', l_name_str:10);
      writeln('Address   : ', address_str:15);
      writeln('Age       : ', age_value:3, ' ' Married : ', married_str:1);
      writeln('Comment  : ', comment_str);
      writeln('Amount purchased : ', amount_value:5:2);
      writeln('');
      rc := d4skip(data_file, Longint(1));
    end;
  end;

begin
  {$IFDEF N4OTHER}
    writeln('This program can only be used for Clipper file format.');
```

```

    writeln('Check the "Index file compatibility options" in CODEBASE.INC.');
```

```

  {$ELSE}
    code_base := code4init;

    code4autoOpen(code_base, False);
    code4safety(code_base, False);

    data_file := d4open(code_base, 'DATA1.DBF');
```

```

    name_tag := t4open(data_file, nil, 'NAME_TAG');
    address_tag := t4open(data_file, nil, 'ADDR_TAG');
    age_tag := t4open(data_file, nil, 'AGE_TAG');
    date_tag := t4open(data_file, nil, 'DATE_TAG');
```

```

    d4tagSelect(data_file, name_tag);

    f_name := d4field(data_file, 'F_NAME');
    l_name := d4field(data_file, 'L_NAME');
    address := d4field(data_file, 'ADDRESS');
    age := d4field(data_file, 'AGE');
    birth_date := d4field(data_file, 'BIRTH_DATE');
    married := d4field(data_file, 'MARRIED');
    amount := d4field(data_file, 'AMOUNT');
    comment := d4field(data_file, 'COMMENT');
```

```

    print_records;

    code4close(code_base);
    code4initUndo( code_base );
  {$ENDIF}

  readKey;
  DoneWinCrt;
end.
```

The function that provides this functionality is **t4open**. It opens the specified file and returns the **TAG4** pointer of its single tag.

Code  
fragment from  
NoGroup2

```
name_tag := t4open(data_file, nil, 'NAME_TAG');
address_tag := t4open(data_file, nil, 'ADDR_TAG');
age_tag := t4open(data_file, nil, 'AGE_TAG');
date_tag := t4open(data_file, nil, 'DATE_TAG');
```

Alternatively, **i4open** may be used to open an **.NTX** index file. When doing so, the **.NTX** file name extension must be provided.

---

## Reindexing

Index files become out of date if they remain closed when their data file is modified. Alternatively, a computer could be turned off in the middle of an update. When an index is out of date, the data file may contain records that do not have corresponding tag entries, or the index file may have tag entries that specify the wrong record.

To rectify this problem, it is sometimes necessary to reindex the index files. You could accomplish this in two ways. First you could delete the index file from your system and then recreate the index, or you could open an existing index file and call one of the reindex functions.

---

### Reindexing All Index Files

If you suspect one or all of your index files may be out of date, you can use the **d4reindex** function. This function reindexes any index files that are associated with that data file and are currently open. This includes both production and non-production index files.

```
rc := d4reindex( data_file );
```

---

### Reindexing A Single Index

If you have several index files but only one is out of date, you can reindex just that index file by using the **i4reindex** function. This function reindexes the index file corresponding to the **INDEX4** pointer you pass to it. All tags associated with the **INDEX4** structure are automatically updated.

```
var
  db: DATA4;
  ind: INDEX4;
...
db := d4open( cb, 'TEST' );
ind := d4index( db, 'INDEX' );

i4reindex( ind );
```

You can obtain the index's **INDEX4** pointer either from the **i4open** or **i4create** functions, or by using **d4index** with the index file's name.

## Advanced Topics

The next section contains details on copying the index file structures and determining the current index key.

### Copying Index File Structures

You may occasionally require a new index file that has an identical structure as an existing index file. This situation usually arises when you are making a copy of a data file.

As a parallel function to **d4fieldInfo**, **CodeBase** provides the function **i4tagInfo**, which returns the **TAG4INFO** array of an existing index file in its native 'C' format. You can pass this array pointer **ByVal** to **d4createCB** or **i4createCB** to create a new index file.



#### WARNING

The **TAG4INFO** returned by **i4tagInfo** is allocated dynamically and should therefore be deallocated after you are finished using it. The **u4free** function may be called to free up the memory for this purpose.



#### PROGRAM CopyData2

Program CopyData2 uses the structure from the data and index files to create a new data and index file. An **Input Box** is used to retrieve the name of the file to be copied.

```
program COPYDAT2 ;

uses CodeBase, WinCrt, Sysutils;

var
  code_base : CODE4 ;
  data_file, data_copy : DATA4 ;
  index : INDEX4 ;
  field_info : PFIELD4INFO ;
  tag_info : PTAG4INFO ;
  dataname1 : array[0..255] of char ;
  dataname2 : array[0..255] of char ;

begin
  if ParamCount <> 2 then
    begin
      writeln( ' USAGE: COPYDAT2 <FROM FILE> <TO FILE>' ) ;
      Halt ;
    end ;

    code_base := code4init ;
    code4safety( code_base, 0 ) ;

    StrPCopy( dataname1, ParamStr(1) ) ;
    data_file := d4open( code_base, dataname1 ) ;
    error4exitTest( code_base ) ;

    index := d4index( data_file, dataname1 ) ;
    if index <> nil then
      tag_info := i4tagInfo( index ) ;

    field_info := d4fieldInfo( data_file ) ;
    StrPCopy( dataname2, ParamStr(2) ) ;
    d4create( code_base, dataname2, field_info, tag_info ) ;
    u4free( PVoid(field_info) ) ;
    u4free( PVoid(tag_info) ) ;
    code4close( code_base ) ;

    readKey ;
    DoneWinCrt ;
  end.
```



---

## 4 Queries And Relations

---

One of the most powerful features of CodeBase is its relation and query capabilities. By using relations, you can turn a collection of separate data files into a fully relational database. Relations are used for two tasks: lookups and queries.

Lookups are used to automatically locate records between related data files. Essentially, this makes several related data files act like a single large data file.

Queries retrieve groupings of records from the database that meet conditions that you can specify at run-time. Queries use CodeBase's Query Optimization which greatly reduces the time it takes to perform queries.

---

### Relations

In its simplest form, a relation is a connection between two data files that specifies how the records from the first data file are used to locate one or more records from the second.

#### Master & Slave Data File

The data file that controls the others is called the master. The controlled data files are called slaves. A master can have any number of slaves which in turn can be masters of other slaves.

#### Master Expression & Slave Tag

To create a relation, you usually provide two pieces of information. The first is the *master expression*, and the second is a tag based on the slave data file called the *slave tag*.

The purpose of the master expression is to generate an index key called the *lookup key*. The master expression is evaluated for each record in the master data file and the resulting value is that record's lookup key. A record in the slave is then found by performing a seek using the lookup key on the slave tag.

The most common type of master expressions are ones which only contain the name of a field from the master data file. The slave tag ordering is usually based on an identical field in the slave data file. Figure 4.1 shows a relation of this type.




Figure 4.1 Relation on two data files

In this example, the master expression is "ID" and the slave tag is the tag STU\_ID\_TAG. The index expression for the slave tag is "STU\_ID". If, for example, the current record for the master data file is set to record number 4, the master expression would then evaluate to "423232". The lookup on the slave data is then performed, locating record number 3.

### Composite Records

Related records from the master and slaves data files are collectively referred to as a *composite record*. The composite record consists of the fields from the master along with the fields from the slave data files. The following is a composite record from the previous relation example:




Figure 4.2 A Composite Record

When using fields from the composite record in a dBASE expression, you must precede any field names from slave data files with the field name qualifier. The field name qualifier is the name of the data file followed by the "->" characters. This avoids any potential naming conflicts caused by data files with field names in common.

For example, " ENROLL->STU\_ID = 876987" is a dBASE expression containing fields from the slave data file. This expression could be used as part of a query (discussed later in this chapter).

### Composite Data Files

The composite data file is the set of composite records that are produced by the relation. Figure 4.3 lists the composite data file for the example relation.

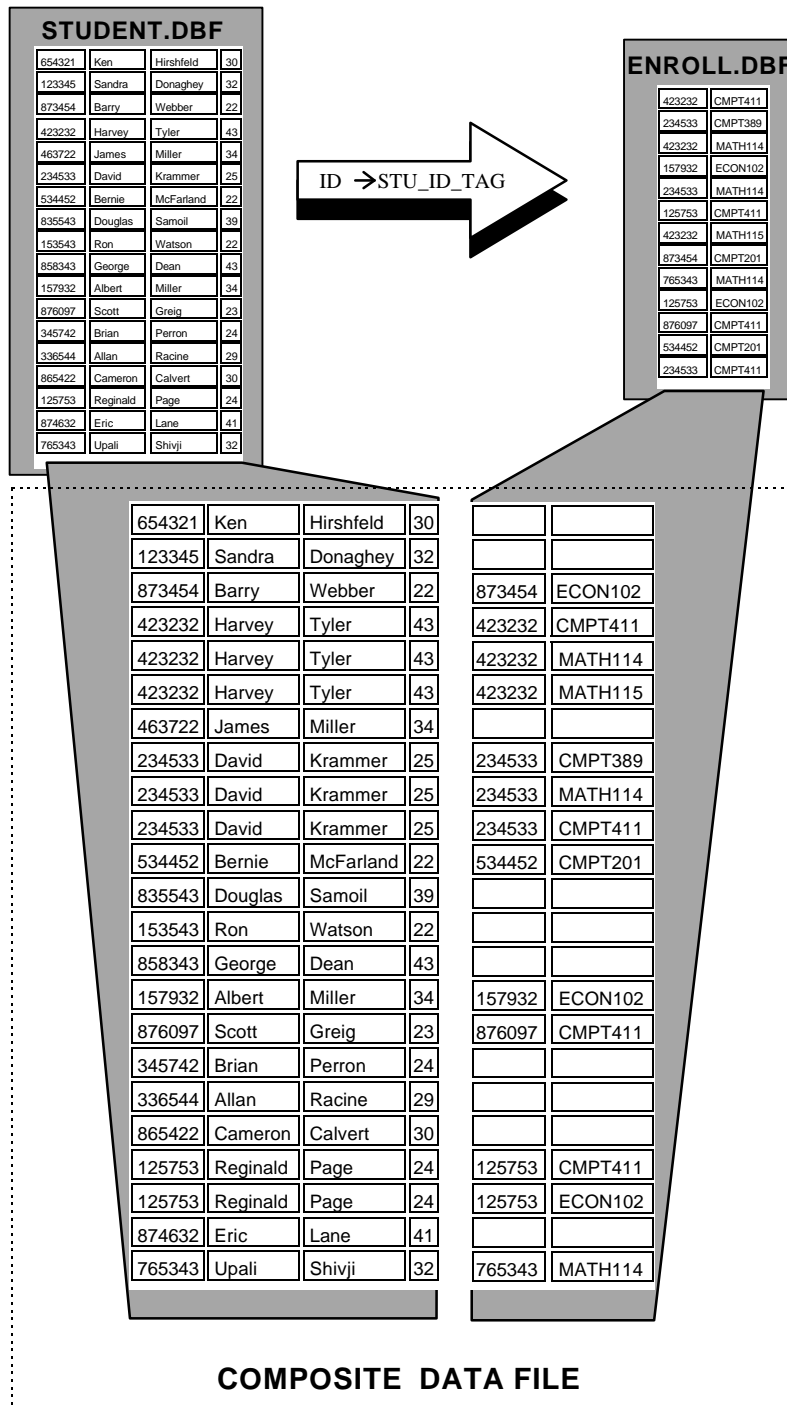


Figure 4.3 The Composite Data File



The composite data file does not physically exist on disk. It is merely a way of viewing the information in the data files of the relation.

For example, when the relation is applied to record number 3 of the master data file the following composite record is generated:

Composite Data File Record # 3

873454	Barrv	Webber	22	87345	CMPT201
--------	-------	--------	----	-------	---------

By default, if the relation does not find a match in the slave data file, the slave's fields in the composite record are left as blanks. For example record number 1 in the master data file does not have a corresponding slave record, so the generated composite record is:

Composite Data File Record # 1

654321	Ken	Hirshfield	30		
--------	-----	------------	----	--	--

## Complex Relations

Although having a relation between two data files is quite useful, it is often necessary to set relations between one master and several slaves, or to set relations between the slaves and other data files.

### Single Master, Multiple Slave Relations

Relations with one master and several slaves are very similar to the single master, single slave relations. In this case there is a master expression and slave tag for each slave, and the composite record consists of fields from all the data files. CodeBase permits any number of slaves for a single master. The only limitations are the resources of your system.

### Multi-Layered Relations

There are many situations that require a master with a relation to a slave, which in turn has a relation to other data files. CodeBase supports an unlimited number of these multi-layered relations and will automatically perform lookups in any associated slave data files.

#### The Top Master

A master data file is a master only in the context of a specific relation. The data file can also be slave in a different relation. If the data file is not a slave in any other relation, it is called a *top master*.

#### The Relation Set

A relation set consists of a top master and all other connected relations. There must be exactly one top master in a relation set.

Figure 4.4 describes a new relation between the ENROLL.DBF data file and the new data file COURSE.DBF. It shows the entire relation set and points out the various relations between the data files.


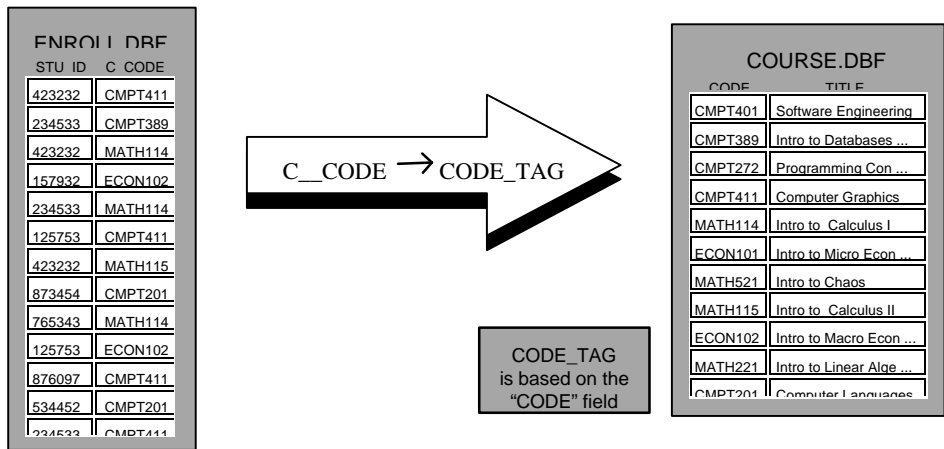


Figure 4.4 A Relation Set

## Relation Types

There are three basic types of relations. They are *exact match*, *scan* and *approximate match* relations. The type of relation determines what record or records are located during a lookup.

### Exact Match Relations

An *exact match* relation permits only a single match between the master and slave data files. Both one to one and many to one relations are exact match relations. In a one to one relation, there is only one record in the master that matches a single record in the slave. In a many to one relation, there are one or more records in the master that match a single slave record.

If there are multiple slave records that match a single master record, then a composite record is only generated for the first slave record, as illustrated below:

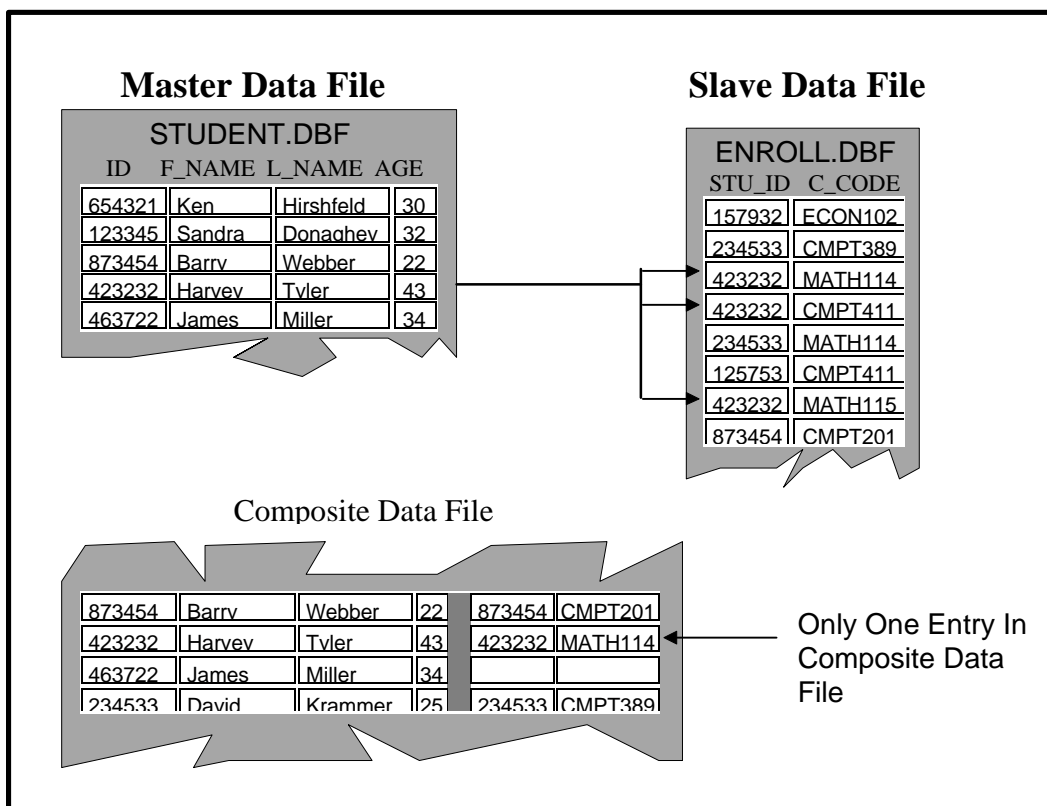


Figure 4.5 An Exact Match Relation

In the above example, there are three records in the slave data file that are matches for record number 4 of the master. Because this is an exact relation a composite record is made from only the first matching record of the slave data file.

**Scan Relations** In a *scan* relation, if there are multiple slave records for a master record, there is one composite record in the extended data file for each of the matching slave records. Both one to many and many to many relations are scan relations. In a one to many relation, each record in the master may have multiple matching slave records. A many to many relation is the same as one to many except that different master records may match the same slave record.

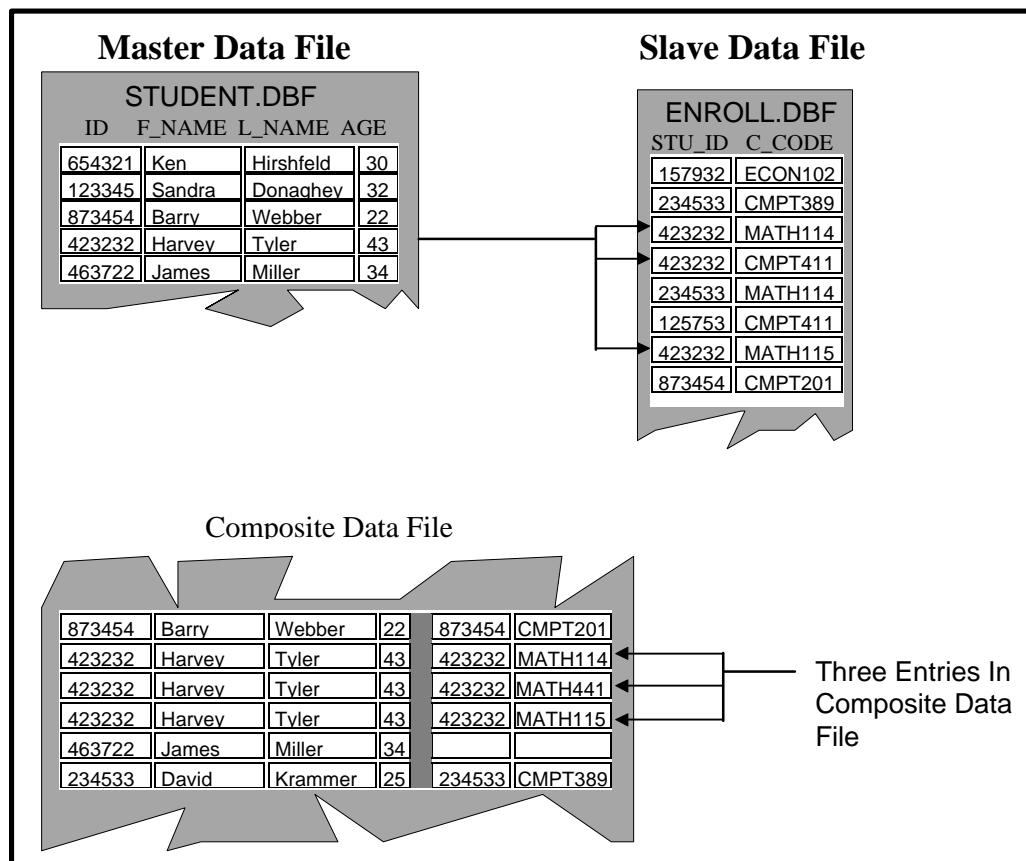


Figure 4.6 A Scan Relation

### Approximate Match Relations

The final type of relation is the approximate match relation. This is similar to the exact match relation in that it only permits one match for a master record. The only difference is the way it behaves when an exact match is not found in the slave data file. If the match fails, the slave record whose index key appears next in the slave tag is used instead. Approximate match relations are generally quite rare and are usually used only when a range of records are represented by a single high value.

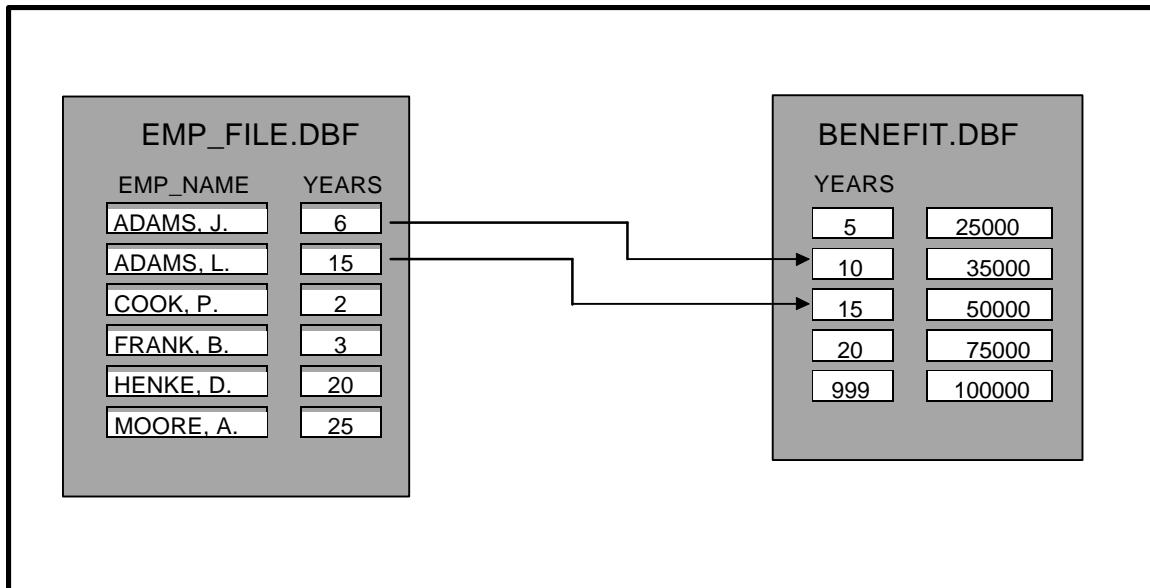


Figure 4.7 Approximate Match Relations

An approximate match relation is shown in Figure 4.7. In this case, the employees' retirement benefits are determined by the number of years that they put into the company. Instead of making an entry for each possible year served, the BENEFIT.DBF only lists the upper limit for each pay out level. The first pay out level is from zero to five years, the second is six to ten, et cetera until the maximum entry of twenty-one years and above pays out 100000.

## Creating Relations



### PROGRAM Relate1

Relations are created during the execution of your application. Before you can create a relation, you must have opened all the data files and any index files that are used in the relation set.

Program Relate1 sets up a scan type relation between the STUDENT.DBF and ENROLL.DBF data files that were illustrated in Figure 4.1. This program lists all of the records in composite data file.

```
program RELATE1 ;
uses CodeBase, WinCrt, Sysutils ;

var
  code_base : CODE4 ;
  student, enrollment : DATA4 ;
  master, slave : RELATE4 ;
  id_tag, name_tag : TAG4 ;

Procedure open_data_files ;
begin
  code_base := code4init ;

  {$IFDEF N4OTHER}
    code4autoOpen( code_base, 0 ) ;
  {$ENDIF}

  student := d4open( code_base, 'STUDENT' ) ;
  enrollment := d4open( code_base, 'ENROLL' ) ;

  {$IFDEF N4OTHER}
    name_tag := t4open( student, nil, 'NAME' ) ;
    id_tag := t4open( enrollment, nil, 'STU_ID' ) ;
```

```

    {$ELSE}
        name_tag := d4tag( student, 'NAME' ) ;
        id_tag   := d4tag( enrollment, 'STU_ID_TAG' ) ;
    {$ENDIF}

    error4exitTest( code_base ) ;
end ;

Procedure set_relation ;

begin
    master := relate4init( student ) ;
    if master = nil then Halt ;

    slave := relate4createSlave( master, enrollment, 'ID', id_tag ) ;

    relate4type( slave, relate4scan ) ;

    relate4top( master ) ;
end ;

Procedure print_record ;

var
    relation : RELATE4 ;
    data : DATA4 ;
    field : FIELD4 ;
    j : integer ;

begin
    relation := master ;
    while relation <> nil do
        begin
            data := relate4data( relation ) ;

            for j := 1 to d4numFields( data ) do
                write( ' ', f4memoStr( d4fieldJ( data, j ))) ;

                relate4next( @relation ) ;
            end ;
            writeln( ' ' ) ;
        end ;
    end ;

Procedure list_records ;

var
    rc : integer ;

begin
    rc := relate4top( master ) ;

    while rc <> r4eof do
        begin
            print_record ;
            rc := relate4skip( master, 1 ) ;
        end ;

        writeln( ' ' ) ;

        code4unlock( code_base ) ;
    end ;

begin
    student      := nil ;
    enrollment   := nil ;
    master       := nil ;
    slave        := nil ;

    open_data_files ;

    set_relation ;

    list_records ;

    code4unlock( code_base ) ;
    relate4free( master, 0 ) ;

    code4close( code_base ) ;
    code4initUndo( code_base ) ;

    readKey ;

```

```
DoneWinCrt ;

Halt ;
end.
```

## The RELATE4 Structure

The **RELATE4** structure contains the control information about a single data file in a relation. Every data file that is in the relation set must have its own **RELATE4** structure.

A data file's **RELATE4** structure contains information such as what data file (if any) is the master of this data file and what type of relation exists between the data file and its master.

In a single relation set, a data file can only have one **RELATE4** structure associated with it.

## Specifying The Top Master

The first step for creating a relation entails specifying which data file will be the top master of the relation set. This step is accomplished by a call to **relate4init**, which creates and returns a pointer to the **RELATE4** structure for that data file.

Code fragment  
from RELATE1

```
master := relate4init( student ) ;
```

At this point, you have a complete relation set consisting of one data file -- the top master. Why would you want a relation set with just a single data file in it? A single data file relation set may be used when a query involves only the one data file.

## Adding Slave Data Files

Adding slave data files to a relation is accomplished by a call to **relate4createSlave**. To do this, there are four pieces of information that you must specify.

The master's  
**RELATE4** pointer

First, the master data file's **RELATE4** structure must be provided. This can either be the **RELATE4** pointer returned from **relate4init** or from a previous call to **relate4createSlave**.

The slave's **DATA4**  
pointer

The second piece of information is a pointer to the slave's **DATA4** structure.



### WARNING

A data file should only be used once in a relation set. The result of using a data file more than once in a relation set is undefined and unpredictable.

The master  
expression

The master expression is the third piece of information needed when adding a slave. This dBASE expression is evaluated using the current record of the master data file to produce a lookup key.

A **TAG4** pointer  
from the slave data  
file

The final piece of information is a **TAG4** pointer to one of the slave's tags. During a lookup, a search for the lookup key (generated by the master expression) is made using this tag.

After the slave has been added to the relation set **relate4createSlave** returns a pointer to the slave's **RELATE4** structure. You should save this pointer in a variable if you plan on changing the default relation type or on adding slaves to the slave data file.

Code fragment  
from RELATE1

```
slave := relate4createSlave( master, enrollment, 'ID',
                             id_tag );
```

In the above code fragment, *master* is the **RELATE4** pointer returned from the call to **relate4init** and *enrollment* is the **DATA4** pointer of the slave data file. The master expression is "ID". As a result, the lookup key is the contents of the ID field from the master data file. The lookups are performed using a tag from the slave data file. This tag is identified by *idTag*.

---

### Creating Slaves Without Tags

CodeBase allows an alternative method of performing lookups that does not use tags from the slave data file. Instead of having the master expression evaluate to a lookup key, it evaluates to a record number. This record number specifies the physical record number of the slave record retrieved from the slave data file.

To use this method, simply pass nil for the **TAG4** parameter of **relate4createSlave**. This method is mainly useful on static unchanging slave data files that are never packed. This method has the advantage of being faster and more efficient than performing seeks on a tag. The following is an example of this method.

```
slave := relate4createSlave(master, enrollment, "RECNO()",
                             nil );
```

Since the master expression is "RECNO()", it returns the record number of master data file's current record. As a result, record *n* in the master must match record *n* in the slave.

If the master expression contains a reference to a non-existent record number ( $\leq 0$  or  $> \mathbf{d4recCount}$ ), CodeBase generates a -70 error ("Reading File", see Appendix A: Error Codes in the *Reference Guide*).

---

### Setting The Relation Type

The default type of relation is an exact match relation. You can change a relation's type by calling **relate4type**. This function takes a **RELATE4** pointer and an integer specifying the type of the relation. The valid integer values are specified by the following defined CodeBase constants:

- **relate4exact** (Default) This specifies an exact match relation.
- **relate4approx** This specifies an approximate match relation.
- **relate4scan** This specifies a scan relation.

Since the master data file may have multiple slaves each using a different type of relation, it is the slave's **RELATE4** pointer that is passed to the **relate4type** function. In Relate1 the **RELATE4** pointer *slave* is passed to the **relate4type** function.

Code fragment  
from RELATE1

```
relate4type( slave, relate4scan );
```



**Note**

**relate4type** always takes a pointer to the slave in the relationship, never the master.

---

**Setting The Error Action**

Sometimes a record in the slave data file cannot be located from a master data file record. This occurs when you are using an exact type relation and there is no match for the lookup key, or when the relation is using direct record lookup and the generated record number does not exist in the slave data file. Under either of these circumstances, there are several ways in which CodeBase can react. **relate4errorAction** specifies which method is used and it accepts a pointer to the slave's **RELATE4** structure and an integer that specifies the type error action. The valid error action codes are as follows:

- **relate4blank** (Default) This means that when a slave record cannot be located, it becomes a blank record.
- **relate4skipRec** This code means that the entire composite record is skipped as if it did not exist.
- **relate4terminate** This means that a CodeBase error is generated and the CodeBase relate function, usually **relate4skip**, returns an error code.

---

**Moving Through The Composite Data File**

CodeBase provides three functions for moving through the records in the composite data file. They are **relate4top**, **relate4bottom** and **relate4skip**. These functions are used mainly for obtaining composite records when a query is performed.

---

**Listing The Composite Data File**

When used together, these three functions allow you iterate through all of the composite data file. The following code segment demonstrates this process:

Code fragment  
from Relate1

```

Procedure list_records ;
var
  rc : integer ;
begin
  rc := relate4top( master ) ;

  while rc <> r4eof do
  begin
    print_record ;
    rc := relate4skip( master, 1 ) ;
  end ;

  writeln( ' ' ) ;

  code4unlock( code_base ) ;
end ;

```

Finding  
the first record in  
the Composite data  
file

**relate4top** sets the current record of the master data file to the first record in the composite data file. If there are slaves in the master's slave family, lookups are automatically performed. If there are no composite records in the composite data file, **r4eof** is returned.

Skipping through  
records

For each iteration of the **While** loop, **relate4skip** is used to move to the next record in the composite data file. Just as in the **d4skip** function, **relate4skip's** second parameter specifies the number of records to be skipped. When there are no more records in the composite data file, **r4eof** is returned.



## WARNING

Call **relate4top** or **relate4bottom** before **relate4skip** is used.

Skipping backwards through the composite data file is also permitted. Since skipping backwards requires extra internal overhead, you must either call **relate4bottom** or **relate4skipEnable** to initialize the backwards skipping abilities.

An example of this is given in the following code segment:

```
Procedure list_records ;

var
  rc : integer ;

begin
  rc := relate4bottom( master ) ;

  while rc <> r4bof do
  begin
    print_record ;
    rc := relate4skip( master, -1 ) ;
  end ;

  writeln( ' ' ) ;

  code4unlock( code_base ) ;
end ;
```

## Queries And Sorting

Once a relation has been set up, you can then perform queries on the composite data file. A *query* allows you to retrieve selected records from the composite data file by specifying a query expression. Essentially the query specifies a subset of the composite data file, which is called the *query set*, by filtering out any composite records that do not meet the search criterion.

In addition to performing queries, you can also specify the order in which the composite records are presented.

When a query is specified, CodeBase uses its Query Optimization, whenever possible, to minimize the number of necessary disk accesses. This greatly improves performance when retrieving composite records from the query set.



### PROGRAM Relate2

Program Relate2 demonstrates how queries can be performed on a single data file.

```
program RELATE2 ;

uses CodeBase, WinCrt, Sysutils ;

var
  code_base : CODE4 ;
  student   : DATA4 ;
  rc, num_fields, j : integer ;

Procedure open_data_files ;
```

```

begin
    student := d4open( code_base, 'STUDENT' ) ;
    error4exitTest( code_base ) ;
end ;

Procedure print_record( data_file : DATA4 ) ;

var
    j : integer ;

begin
    for j := 1 to d4numFields( data_file ) do
        write( ' ', f4memoStr( d4fieldJ( data_file, j ))) ;
        writeln( ' ' ) ;
    end ;

Procedure query( data_file : DATA4 ; expr : PChar ; order : PChar ) ;

var
    relation : RELATE4 ;
    rc : integer ;

begin
    relation := nil ;

    relation := relate4init( data_file ) ;
    if relation = nil then Halt ;

    relate4querySet( relation, expr ) ;
    relate4sortSet( relation, order ) ;

    rc := relate4top( relation ) ;
    while rc <> r4eof do
        begin
            print_record( data_file ) ;
            rc := relate4skip( relation, 1 ) ;
        end ;

        writeln( ' ' ) ;

    code4unlock( code_base ) ;
    relate4free( relation, 0 ) ;
end ;

begin
    student := nil ;

    code_base := code4init ;
    code4autoOpen( code_base, 0 ) ;

    open_data_files ;

    query( student, 'AGE > 30', ' ' ) ;

    query( student, 'UPPER(L_NAME) = 'MILLER'', 'L_NAME + F_NAME' ) ;

    code4close( code_base ) ;
    code4initUndo( code_base ) ;

    readKey ;
    DoneWinCrt ;

    Halt ;
end.

```

## The Query Expression

To generate a query, all you need to provide is a dBASE expression, called the *query expression*, which evaluates to a Logical value. This expression is used to determine whether a composite record should be included in the query. If the query expression evaluates to .TRUE. then the record is kept, otherwise it is discarded from the query.

Setting the query expression    The query expression is specified for a relation set by the **relate4querySet** function. This function takes a pointer to the top master's **RELATE4** structure and specifies the query expression for the entire relation set.

The query expression can reference fields from different files in the relation set. When a query is made, the default data base used is the master. This means that the field will automatically be associated with the master data base. To refer fields of different data files use the file's alias followed by an arrow (->) and the field name.

For example, suppose a relation consists of a master data file called FATHER.DBF and a slave data file called SON.DBF. Assume that each file has a field called NAME.

For these data files, the following queries are equivalent:

```
relate4querySet( relation, 'FATHER->NAME = "Ben Nyland"' );
relate4querySet( relation, 'NAME = "Ben Nyland"' );
```

These queries compare a field called NAME in FATHER.DBF with the string "Ben Nyland". The file name FATHER does not need to be specified in the query because the FATHER.DBF file is a master.

In order to check for the composite record that has Ben Nyland and his son Eric Nyland, the query can be specified by either of the following:

```
relate4querySet( relation, 'FATHER->NAME = "Ben Nyland"
                           .AND. SON->NAME = "Eric Nyland"');
relate4querySet( relation, 'NAME = "Ben Nyland" .AND. SON->NAME = "Eric Nyland"');
```

The query expression can be changed at any time, although **relate4top** or **relate4bottom** must then be called before **relate4skip** can be called.

---

## The Sort Expression

The *sort expression* of a relation is very similar to an index expression. They are both dBASE expressions, and are both used to determine the sort ordering.

The difference is that the sort expression determines the sort order of the query set and you are allowed to use fields from any data file in the relation set.

Setting the sort expression    The sort expression is specified for a relation set by **relate4sortSet**. This function takes a pointer to the top master's **RELATE4** structure and specifies the sort ordering for the entire query set.

The sort expression can contain field names from any data file in the relation set. It is required that any fields, except those from the top master, are qualified by the data file's alias.

```
relate4sortSet( relation, 'L_NAME + F_NAME + CLASSES->C_CODE');
```

Like the query expression, the sort expression can be changed at any time, although **relate4top** or **relate4bottom** must then be called before **relate4skip** can be called.

## Performing Queries on Relation Sets

Queries can be performed on relation sets of any size, including those consisting of a single data file. The following figures illustrate two queries. They are on the composite data file shown in Figure 4.3.

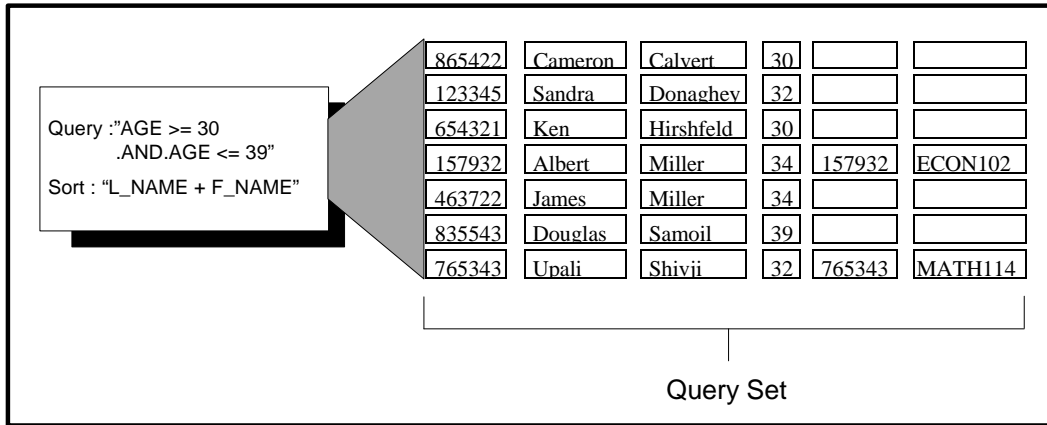


Figure 4.8 Query Example One

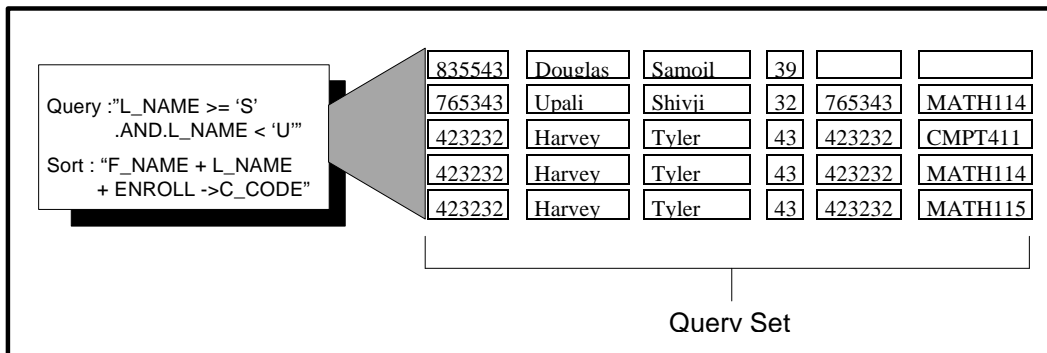


Figure 4.9 Query Example Two

## Accessing The Query Set

The query set is accessed by **relate4top**, **relate4bottom** and **relate4skip**. When a query has been specified, these three functions then access a query set instead of the entire composite data file.

## Generating The Query Set

The query set is initially formed by the first call to either **relate4top** or **relate4bottom**. These functions analyze the query expression for the most efficient way of performing the query. They then locate all of records that are in the query set. If a sort expression has been specified, the composite records are sorted. And finally, lookups are performed on the data files in the relation to locate the first or last composite record in the query set.



### WARNING

When complex relations and sort expressions are used on a relation set consisting of large data files, **relate4top** and **relate4bottom** can take considerable time to execute.

As long as the query and sort expressions are not changed, further calls to **relate4top** and **relate4bottom** do not cause the query set to be regenerated.

## Regenerating the Query Set

You can force CodeBase to completely regenerate the query set by calling **relate4changed**. This forces CodeBase to discard any buffered information and regenerate the query set based on the current state of the relation set's data files. After **relate4changed** is called, **relate4top** or **relate4bottom** must be called before any further calls to **relate4skip** may be made.

## Queries On Multi-Layered Relation Sets

The steps involved in performing queries on multi-layered relation sets are the same as on a relation set containing a single data file. You create your relations, set the types of the relations, specify query and sort expressions and retrieve the composite records from the query set.



### PROGRAM Relate3

Program Relate3 creates a multi-layered relation between the STUDENT, ENROLL and CLASSES data files and then performs queries on it.

```
program RELATE3 ;

uses CodeBase, WinCrt, Sysutils ;

var
  code_base : CODE4 ;
  student, enrollment, classes : DATA4 ;
  student_id, first_name, last_name, age, class_code, class_title : FIELD4 ;
  code_tag, id_tag : TAG4 ;
  class_rel, student_rel, enroll_rel : RELATE4 ;

Procedure open_data_files ;

begin
  code_base := code4init ;

  {$IFDEF N4OTHER}
    code4autoOpen( code_base, False ) ;
  {$ENDIF}

  student      := d4open( code_base, 'STUDENT' ) ;
  enrollment   := d4open( code_base, 'ENROLL' ) ;
  classes      := d4open( code_base, 'CLASSES' ) ;

  student_id   := d4field( student, 'ID' ) ;
  first_name   := d4field( student, 'F_NAME' ) ;
  last_name    := d4field( student, 'L_NAME' ) ;
  age          := d4field( student, 'AGE' ) ;
  class_code   := d4field( classes, 'CODE' ) ;
  class_title  := d4field( classes, 'TITLE' ) ;

  {$IFDEF N4OTHER}
    id_tag     := t4open( student, nil, 'ID' ) ;
    code_tag   := t4open( enrollment, nil, 'C_CODE' ) ;
  {$ELSE}
    id_tag     := d4tag( student, 'ID_TAG' ) ;
    code_tag   := d4tag( enrollment, 'C_CODE_TAG' ) ;
  {$ENDIF}

  error4exitTest( code_base ) ;
end ;

Procedure print_students ;

begin
  writeln( ' ' ) ;
  write( ' ', f4str( first_name ) ) ;
  write( ' ', f4str( last_name ) ) ;
  write( ' ', f4str( student_id ) ) ;
  write( ' ', f4str( age ) ) ;
  writeln( ' ' ) ;
end ;

Procedure set_relation ;

begin
```

```

class_rel := relate4init( classes ) ;

enroll_rel := relate4createSlave( class_rel, enrollment, 'CODE', code_tag ) ;

student_rel := relate4createSlave(enroll_rel, student, 'STU_ID_TAG', id_tag);
end ;

Procedure print_student_list( expr : PChar ; direction : Longint ) ;

var
  rc, end_value : integer ;

begin
  relate4querySet( class_rel, expr ) ;
  relate4sortSet( class_rel, 'STUDENT->L_NAME + STUDENT->F_NAME' ) ;

  relate4type( enroll_rel, relate4scan ) ;

  if direction > 0 then
  begin
    rc := relate4top( class_rel ) ;
    end_value := r4eof ;
  end
  else
  begin
    rc := relate4bottom( class_rel ) ;
    end_value := r4bof ;
  end ;

  writeln( ' ' ) ;
  write ( ' ', f4str( class_code ) ) ;
  writeln( ' ', f4str( class_title ) ) ;

  while rc <> end_value do
  begin
    print_students ;
    rc := relate4skip( class_rel, direction ) ;
  end ;
end ;

begin
  open_data_files ;

  set_relation ;

  print_student_list( 'CODE = 'MATH114 ' ', 1 ) ;

  print_student_list( 'CODE = 'CMPT411 ' ', -1 ) ;

  code4unlock( code_base ) ;
  relate4free( class_rel, 0 ) ;

  code4close( code_base ) ;
  code4initUndo( code_base ) ;

  readKey ;
  DoneWinCrt ;

  Halt ;
end.

```

## Lookups On Relations



### PROGRAM Relate4

The method illustrated above for retrieving composite records from the composite data file is well suited for performing queries and generating reports, but has unnecessary overhead if all you want to do is perform lookups to slave data files. As a result CodeBase provides an alternative method for performing lookups that is independent of the query set and sort orderings.

Program Relate4 sets up an exact match type relation between the STUDENT.DBF and ENROLL.DBF that were illustrated in Figure 4.1. This program performs some seeks and lookups.

## 80 CodeBase

```
program REL4 ;

uses CodeBase, WinCrt, Sysutils ;

var
  code_base : CODE4 ;
  student, enrolment : DATA4 ;
  id, f_name, l_name, age, c_code : FIELD4 ;
  id_tag, name_tag : TAG4 ;
  master, slave : RELATE4 ;

Procedure open_data_files ;

begin
  code_base := code4init ;

  {$IFDEF N4OTHER}
    code4autoOpen( code_base, False ) ;
  {$ENDIF}

  student := d4open( code_base, 'STUDENT' ) ;
  enrolment := d4open( code_base, 'ENROLL' ) ;

  id := d4field( student, 'ID' ) ;
  f_name := d4field( student, 'F_NAME' ) ;
  l_name := d4field( student, 'L_NAME' ) ;
  age := d4field( student, 'AGE' ) ;
  c_code := d4field( enrolment, 'C_CODE_TAG' ) ;

  {$IFDEF N4OTHER}
    name_tag := t4open( student, nil, 'NAME' ) ;
    id_tag := t4open( enrolment, nil, 'STU_ID' ) ;
  {$ELSE}
    name_tag := d4tag( student, 'NAME' ) ;
    id_tag := d4tag( enrolment, 'STU_ID_TAG' ) ;
  {$ENDIF}

  error4exitTest( code_base ) ;
end ;

Procedure set_relation ;

begin
  master := relate4init( student ) ;
  slave := relate4createSlave( master, enrolment, 'ID', id_tag ) ;
end ;

Procedure seek( data_file : DATA4 ; tag : TAG4 ; relation : RELATE4 ; key :
                                                         PChar ) ;

var
  old_tag : TAG4 ;

begin
  old_tag := d4tagSelected( data_file ) ;
  d4tagSelect( data_file, tag ) ;

  d4seek( data_file, key ) ;
  relate4doAll( relation ) ;

  d4tagSelect( data_file, old_tag ) ;
end ;

Procedure print_record ;

begin
  write ( ' ', f4str( f_name ) ) ;
  write ( ' ', f4str( l_name ) ) ;
  write ( ' ', f4str( id ) ) ;
  write ( ' ', f4str( age ) ) ;
  writeln( ' ', f4str( c_code ) ) ;
end ;

begin
  student := nil ;
  enrolment := nil ;
  master := nil ;
  slave := nil ;

  open_data_files ;
```



```

set_relation ;

seek( student, name_tag, master, 'Tyler          Harvey          ' ) ;

print_record ;

seek( student, name_tag, master, 'Miller          Albert          ' ) ;

print_record ;

code4unlock( code_base ) ;
relate4free( master, 0 ) ;

code4close( code_base ) ;
code4initUndo( code_base ) ;

readKey ;
DoneWinCrt ;

Halt ;
end.

```

## Performing A Lookup

Performing lookups is quite simple. First you locate the record in the master data file that you wish to perform the lookup from, using the normal **d4top**, **d4bottom**, **d4go**, **d4seek**, **d4seekDouble**, **d4seekN**, **d4seekNext**, **d4seekNextDouble**, **d4seekNextN** and **d4skip** functions. The lookup is then accomplished by a call to **relate4doAll**. This function is passed a **RELATE4** pointer of a master data file and it performs lookups on the slave data files of that master. In addition, it travels down the relation set performing lookups on the slave's slaves.

The following function performs a seek on a data file and then performs the lookups on it's slaves.

Code fragment  
from Relate4

```

Procedure seek( data_file : DATA4 ; tag : TAG4 ;
               relation : RELATE4 ; key : PChar ) ;

var
  old_tag : TAG4 ;

begin
  old_tag := d4tagSelected( data_file ) ;
  d4tagSelect( data_file, tag ) ;

  d4seek( data_file, key ) ;
  relate4doAll( relation ) ;

  d4tagSelect( data_file, old_tag ) ;
end ;

```

**relate4doOne**

A similar function to **relate4doAll** is **relate4doOne**. This function accepts the **RELATE4** pointer to a slave and performs a lookup on that slave.



### WARNING

Any query and sort expressions are ignored by functions **relate4doAll** and **relate4doOne**. Consequently, these functions provide somewhat independent functionality. Using them in conjunction with relate functions such as **relate4top**, **relate4bottom**, and **relate4skip** is not particularly useful.

## Iterating Through The Relations



### WARNING

As an aid in writing generic reusable code, CodeBase provides the function **relate4next**, which allows you to iterate through the relations in a relation set.

This function takes a **RELATE4** pointer ( passed by reference ) and changes it to a pointer to the next **RELATE4** structure in the relation set.

The **relate4next** function changes the value of the pointer that is passed in as a parameter. Consequently, you should make a copy of this pointer if you need it later.

The following code segment displays all of the records from all of the relations in the relation set.

Code Fragment  
from Relate1

```

Procedure print_record ;
var
  relation : RELATE4 ;
  data : DATA4 ;
  field : FIELD4 ;
  j : integer ;

begin
  relation := master ;
  while relation <> nil do
    begin
      data := relate4data( relation ) ;

      for j := 1 to d4numFields( data ) do
        write( ' ', f4memoStr( d4fieldJ( data, j )) ) ;

        relate4next( @relation ) ;
      end ;
      writeln( ' ' ) ;
    end ;
  end ;

```

---

## 5 Query Optimization

---

### What is Query Optimization

Query Optimization is an application invisible method of returning query results at high speed. The CodeBase Relation/Query modules use Query Optimization to analyze the query condition and return the matching set of records with a minimum number of disk accesses.

This is accomplished by comparing the query expression to the tag sort orderings of the top master tag. If part of the query matches the tag expression, the tag itself is used to filter out records that do not match the expression.

This results in lightning quick performance, even on the largest data files, since very few records are actually physically read. For example, if there were an index built on PRODUCT and the query was:

PRODUCT = 'GIZMO'

The Query Optimization would automatically use the index file to quickly seek to the appropriate record and, using the PRODUCT tag, and skip to retrieve all the records where PRODUCT is equal to 'GIZMO'. Once this is no longer the case, Query Optimization ignores the remaining records.

It makes little difference whether the data file has 500 records or 500,000 records, since the same number of disk reads are performed in each case.

Complex  
expressions

Query Optimization can also work on more complex expressions involving the .AND. and .OR. operators. In these cases, Query Optimization breaks down the whole expression into sub-expressions that can be optimized. For example, the two sub-expressions in the query:

L\_NAME = "SMITH" .AND. F\_NAME = "JOE"

could both be optimized if there were two tags, one based on the expression L\_NAME, the other on F\_NAME. CodeBase could still partially optimize the query if there was a tag on either of the two sub-expressions. Even a partially optimized query can lead to very fast performance.

### When is Query Optimization used

To start with, the Query Optimization capabilities are built into the CodeBase relate/query module, the report module and into CodeReporter. Your applications can use the Query Optimization provided a few simple conditions are met.

## Requirements of Query Optimization

To ensure that the Query Optimization is used to return your queries, the following steps must be taken:

1. Insure that the Master data file's index file(s) are open. In general, the more tags open on the master data files, the greater the chance that query optimization can take effect.
2. Specify a query expression that contains in whole or in part, a top master tag key compared to a constant.
3. The query expression is set using **relate4querySet**. Query optimization is only effective when the query expression involves the top master data file.

The following tables illustrate some of the situations where query optimization is used automatically, and where it cannot be used. The data file assumed by these tables have the following fields and tags:

```
FieldInfo : array[1..11] of FIELD4INFO = (
  (name:'L_NAME'      ; atype:r4str  ; len:10 ; dec:0),
  (name:'F_NAME'      ; atype:r4str  ; len:10 ; dec:0),
  (name:'COMPANY'     ; atype:r4str  ; len:2  ; dec:0),
  (name:'AGE'         ; atype:r4num  ; len:10 ; dec:0),
  (name:'DT'          ; atype:r4date ; len:8  ; dec:0),
  (name:'AMOUNT'      ; atype:r4num  ; len:7  ; dec:2),
  (name:'ADDRESS'     ; atype:r4str  ; len:15 ; dec:0),
  (name:'PRODUCT'     ; atype:r4num  ; len:10 ; dec:0),
  (name:'ID'          ; atype:r4str  ; len:5  ; dec:0),
  (name:'PHONE'       ; atype:r4str  ; len:8  ; dec:0),
  (name:nil           ; atype:0      ; len: 0 ; dec:0) );

TagInfo : array[1..10] of As TAG4INFO = (
  (name:'L_NAME_TAG' ; expression:'L_NAME'      ;
   filter:'' ; unique:0 ; descending:0),

  (name:'COMPANY_TAG' ; expression:'UPPER(COMPANY)' ;
   filter:'' ; unique:0 ; descending:0 ),

  (name:'AGE_TAG'      ; expression:'AGE' ;
   filter:'' ; unique:0 ; descending:0 ),

  (name:'DT_TAG'       ; expression:'DT' ;
   filter:'' ; unique:0 ; descending:0 ),

  (name:'AMNT_TAG'     ; expression:'AMOUNT' ;
   filter:'.NOT.DELETED()'; unique:0 ; descending:0 ),

  (name:'ADDR_TAG'     ; expression:'ADDRESS' ;
   filter:'DELETED().AND..NOT.DELETED()';
   unique:0 ; descending:0 ),

  (name:'COMMENT_TAG' ; expression:'COMMENT' ;
   filter:'' ; unique:0 ; descending:r4descending ),

  (name:'PROD_TAG'     ; expression:'PRODUCT' ;
   filter:'' ; unique:r4uniqueContinue ; descending:0 ),

  (name:'PRO_ID_TAG'   ; expression:'PRODUCT+ID' ;
   filter:'' ; unique:0 ; descending:0 ),

  (name:'PHONE_TAG'    ; expression:'PHONE' ;
   filter:'' ; unique:r4unique ; descending:0 ),

  (name:nil ; expression:nil ; filter:nil ; unique:0 ;
   descending:0) );
```

Query optimization is automatically used in the following situations.

Query Expression	Explanation
L_NAME = "SMITH"	L_NAME is a tag expression which is compared to a constant.
L_NAME >= "S" .AND. L_NAME <= "T"	Each sub expression compares a tag expression against a constant, and so the entire query can use Query Optimization.
L_NAME = "SMITH" .AND. F_NAME = "JOE"	Only the first sub expression can use Query Optimization because there is no tag based on F_NAME.
UPPER(COMPANY) = "IBM"	Again, this is a tag expression compared to a constant. Even though the expression contains a function (UPPER), Query Optimization is still used.
UPPER(COMPANY) = "IBM" .AND. L_NAME="SMITH"	Both parts of the expression use Query Optimization, since both sides compare a tag expression to a constant value.
AGE > 50	Even though the tag is Numeric, the expression can still be optimized.
AGE > 25*2	25*2 is still a constant even though it is a complicated constant.
DTOS(DT) >= "19930101"	This would return all records where the date value in the DT field is greater than or equal to January 1, 1993. If the tag expression was DT instead of DTOS(DT), then a query expression such as DT>=CTOD("01/01/93") would also use Query Optimization. ( Be careful when using CTOD because the format of its parameter depends on the data picture set in <b>code4dateFormat</b> .)
COMMENT = "SALE"	In this case the tag was set using <b>r4descending</b> . Query Optimization can be used because CodeBase is indifferent to the ordering of the tag.
PHONE = "555 6031"	<p>Query Optimization can utilize unique tags as long the tag has not been specified with <b>r4uniqueContinue</b>. <b>r4unique</b> prevents duplicate records from being added to the data base. Therefore, <b>r4unique</b> does not act as a filter for the tag, unlike <b>r4uniqueContinue</b>.</p> <p>Note that when the index is opened, the unique code is automatically set to the default value of <b>CODE4.errDefaultUnique</b>, which is <b>r4uniqueContinue</b>. Use <b>t4uniqueSet</b> to set unique code to <b>r4unique</b> before making a query. See the "Unique Keys" section of the "Indexing" chapter of this guide for more information.</p>

The following expressions cannot use the Query Optimization. Queries on these expressions are not optimized and take a longer time to execute.

Query Expression	Explanation
F_NAME = "JOE"	There is no tag based on F_NAME.
COMPANY = "IBM"	There is no tag based on COMPANY. The tag is UPPER(COMPANY). Again, the tag expression must match the expression being compared to the constant exactly.
L_NAME = F_NAME	This expression cannot be optimized because F_NAME is not a constant.
AMOUNT = 1000.00	Query Optimization will not be used because the tag contains a filter.
ADDRESS = "104 ELM ST"	Query Optimization will not be used because the tag contains a filter, even though the filter is useless. Tags with filters are useless, regardless of the filter result.
PRODUCT ="GIZMO"	Query Optimization will not be used because the tag specifies <b>r4uniqueContinue</b> . <b>r4uniqueContinue</b> acts like a filter, allowing duplicate records in the database but not in the index.
PRODUCT + ID ="GIZMO 13445"	Query Optimization will not be used because the tag consists of two fields.

### How To Use Query Optimization

Query Optimization is automatically enabled whenever a relation operation occurs. CodeBase uses query optimization on the top master data file transparently when retrieving records from the composite data file.

Query optimization, as described above, operates only on the top master data file and only when a tag sort ordering corresponds to a portion of the query expression. All you need to do is ensure that the top master data file has an appropriate index file open.

### How to tell if Query Optimization can be used

It is possible to determine whether Query Optimization can be used by calling the function **relate4optimizeable**. This function returns true (non-zero) when the relation is able to use the Query Optimization and it returns false (zero) when it can not. If there is insufficient memory, the Query Optimization will not be invoked, even if **relate4optimizeable** returns true (non-zero).

Call **relate4optimizeable** before functions that may use query optimization. If false (zero) is returned, it may be possible to create a new index file for those queries that can not use Query Optimization. In the above scenario, the expression "F\_NAME = 'JOE' ", is not able to use query optimization. If a new index were built on "F\_NAME", the execution time of the query would be improved.

---

**Memory  
Requirements  
of Query  
Optimization**

The activation of the Query Optimization also depends on the compiler. Therefore the following restrictions apply:

When a 16 bit compiler is used, the Query Optimization will be able to handle 500,000 records. If the data base has more than half a million records, the query will be made without using query optimization.

It is recommended that 32 bit compiler be used when using data bases that have more than 500,000 records. Query Optimization will be able to handle 32 billion records when a 32 bit compiler is used.





## 6 Date Functions

The date functions allow you to convert between a variety of date formats and to perform date arithmetic. The date functions also allow you to retrieve the various components of a date, such as the day of the week, in numeric form.

### Date Pictures

The format of a date string is represented by a date picture string. A date picture is a string containing several special formatting characters and other characters. The special formatting characters are:

- **C** Century. A 'C' represents the first digit of the century. If two 'C's appear together, then both digits of the century are represented. Additional 'C' s are not used as formatting characters.
- **Y** Year. A 'Y' represents the first digit of the year. If two 'Y's appear together, they represent both digits of the year. Additional 'Y' s are not used as formatting characters.
- **M** Month. One or two 'M's represent the first or both digits of the month. If there are more than two consecutive 'M's, a character representation of the month is returned.
- **D** Day. One or two 'D' s represent the first or both digits of the day of the month. Additional 'D' s are not used as formatting characters.
- **Other Characters** Any character which is not mentioned above is placed in the date string when it is formatted.

For example, if the date August 4 1994 is converted to a date string using the date picture "MM/DD/CCYY", the resulting date string is "08/04/1994". If the same date is converted with the date picture "MMMMMMMM DD, YY" the resulting date string is "August 04, 94".

### Date Formats



#### PROGRAM Date

Dates can be stored in many formats. The two used by CodeBase are the standard format and the Julian day format.

Program Date uses several of the date functions to calculate the number of days until Christmas and until your next birthday.

```
program DATE ;

uses CodeBase, WinCrt, Sysutils;

var
  birthdate : array[0..79] of char ;
  standard  : array[0..8]  of char ;

Function valid_date( date : PChar ) : integer ;

var
  rc : Longint ;

begin
  rc := date4long( date ) ;

  if rc < 1 then
    valid_date := 0
  else
    valid_date := 1 ;
  end ;
end ;
```

```

Procedure how_long_until( month : integer ; day : integer ; title : PChar ) ;

var
  today_standard : array[0..8] of char ;
  today          : array[0..24] of char ;
  date           : array[0..8] of char ;
  dow : PChar ;
  i, year, days : integer ;
  julian_today, julian_date : Longint ;
  year_str      : array[0..4] of char ;
  month_str     : array[0..2] of char ;
  day_str       : array[0..2] of char ;

begin
  for i := 0 to SizeOf( today_standard ) do today_standard[i] := #0 ;
  for i := 0 to SizeOf( today ) do today[i] := #0 ;
  for i := 0 to SizeOf( date ) do date[i] := #0 ;

  date4today( today_standard ) ;
  date4format( today_standard, today, 'MMM DD/CCYY' ) ;

  writeln( '' ) ;
  writeln( 'Today's date is ', today ) ;

  julian_today := date4long( today_standard ) ;

  year := date4year( today_standard ) ;

  Str( year:5, year_str ) ;
  Str( month:3, month_str ) ;
  Str( day:3, day_str ) ;
  StrLCopy( date, year_str, SizeOf( date ) - 1 ) ;
  StrLCat( date, month_str, SizeOf( date ) - 1 ) ;
  StrLCat( date, day_str, SizeOf( date ) - 1 ) ;

  julian_date := date4long( date ) ;

  if julian_date < julian_today then
  begin
    year := year + 1 ;

    Str( year:5, year_str ) ;
    Str( month:3, month_str ) ;
    Str( day:3, day_str ) ;
    StrLCopy( date, year_str, SizeOf( date ) - 1 ) ;
    StrLCat( date, month_str, SizeOf( date ) - 1 ) ;
    StrLCat( date, day_str, SizeOf( date ) - 1 ) ;

    julian_date := date4long( date ) ;
  end ;

  days := julian_date - julian_today ;

  writeln( 'There are ', days, ' days until ', title ) ;
  dow := date4cdow( date ) ;

  writeln( '(which is a ', dow, ' this year)' ) ;
end ;

begin
  how_long_until( 12, 25, 'Christmas' ) ;

  repeat
    writeln( 'Please enter your birthdate' ) ;
    write ( 'in "DEC 20/1993" format: ' ) ;
    readln( birthdate ) ;
    date4init( standard, birthdate, 'MMM DD/CCYY' ) ;
  until valid_date( standard ) = 1 ;

  how_long_until( date4month( standard ), date4day( standard ), 'your next
    birthday' ) ;

  readKey ;
  DoneWinCrt ;
end.

```

## Standard Format

Dates are stored in the data file in "CCYYMMDD" format. This is known as the standard format. The **f4str** function returns a date from a Date field in this standard format. **f4assign** performs the opposite operation by storing a standard format date string to the data file.



### WARNING

If **f4assign** is passed a date string in anything other than standard format, indexing and seeking will not be performed correctly.

## Julian Day Format

Another important date format is the Julian day format, which is stored as **longint**. A Julian day is defined as the number of days since Jan. 1, 4713 BC. Having the date accessible in this format lets you perform date arithmetic. Two dates can be subtracted to find the number of days separating them, or an integer number of days can be added to a date.

## Converting Between Formats

Since dates are only stored in the data file in standard format, conversion functions have been provided for converting between standard and julian day format, as well as converting between standard and any other format.

The interrelationships between the data file, date formats and date functions are illustrated by Figure 6.1.

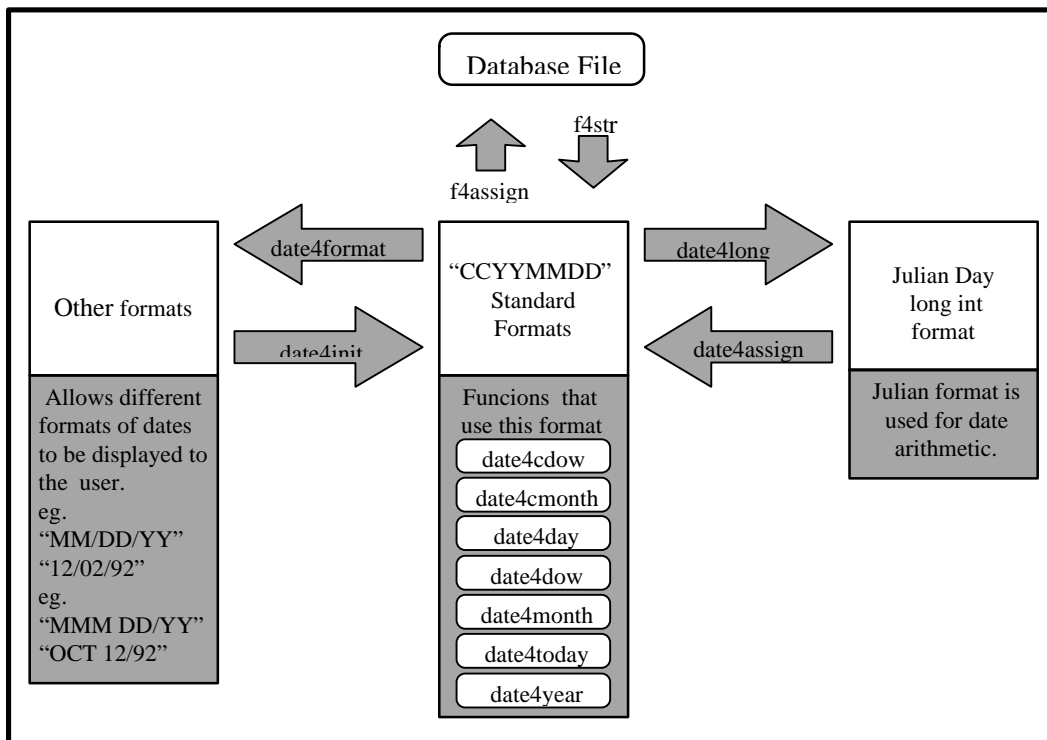


Figure 6.1 Date Conversions

As mentioned above, dates are stored and retrieved from the data file using the functions **f4assign** and **f4str**.

Date strings in standard format can be converted to Julian day format by using the **date4long** function. The inverse of this function is **date4assign**, which converts a Julian date back into a standard date string.

Code fragment  
from Date

```
julian_today := date4long( today_standard ) ;
```

To convert from a standard date string to another format, the **date4format** function is used. This function uses a provided date picture to format the date string.

Code fragment  
from Date

```
date4format( today_standard, today, 'MMM DD/CCYY' ) ;
```

To convert from a date string in any other format to standard format, the **date4init** function is used. This time the date picture string is used to specify what format the original string is in. If any part of the date is missing, the missing portion is filled in using the date January 1, 1980.

Code fragment  
from Date

```
date4init( standard, birthdate, 'MMM DD/CCYY' ) ;
```

---

## Testing For Invalid Dates

You can test for invalid dates using the **date4long** function. If the date string, in standard format, is not a valid date, the **date4long** function will return **-1**. If it is passed a zero length string, the function returns zero. The following example function tests for valid dates:

Code fragment  
from Date

```
Function valid_date( date : PChar ) : integer ;
var
  rc : Longint ;
begin
  rc := date4long( date ) ;

  if rc < 1 then
    valid_date := 0
  else
    valid_date := 1 ;
end ;
```

---

## Other Date Functions

There are several other date functions that perform various date related tasks:

---

### Days and Months as Characters

There are two functions that return the day or month portion of the date string as a character strings. They are:

**date4cdow** The day of the week in character form is returned.

**date4cmonth** The day of the month in character form is returned.

Code fragment  
from Date

```
dow := date4cdow( date ) ;
```

---

## Days, Months and Years as Integers

There are four functions that return a portion of the date string as integer values:

- **date4day** The day of the month is returned as an integer value from 1 to 31.
- **date4dow** The day of the week is returned as an integer value from 1 to 7.
- **date4month** The month of the year is returned as an integer value from 1 to 12.
- **date4year**. Returns the century/year portion of the date as an integer value.

Code fragment  
from Date

```
how_long_until( date4month( standard ), date4day( standard ),
               'your next birthday' ) ;
```

## Miscellaneous Functions

CodeBase also provides the following miscellaneous date functions:

- **date4isLeap** Returns true (non-zero) if the date falls within a leap year and false (zero) otherwise.
- **date4today** Initializes a date string with today's date.



---

## 7 Memory Optimizations

---

Memory optimization is accomplished by buffering portions of data, index and memo files in memory. This improves performance because physically accessing the disk takes longer than accessing memory. In addition, reading and writing several records at once is considerably quicker than reading and writing individual records separately.

Memory optimization is more effective in some operating environments than others. This is because some operating systems already do some memory optimizations at the operating system level and some hard disks do memory optimizations at the hardware level.

Regardless, when memory is available, proper use of CodeBase memory optimizations will always improve performance. This is because CodeBase memory optimizations are designed specifically to optimize database operations. In addition, in network applications, they can reduce requests to the network server.

---

### Using Memory Optimizations

Using memory optimization is quite easy. All you have to do is turn it on. CodeBase has defaults for determining which files should be optimized and whether the files should be optimized for reading and/or writing. If you wish, you can override the defaults for any particular file(s).

---

### Specifying Files To Optimize

Any file that you can open with a CodeBase function can be optimized. This includes data files, index files and memo files. Each file opened by CodeBase has two flags associated with it. The first specifies whether the file is optimized. When a file is optimized, the second flag specifies whether the file is optimized for both reading and writing or just reading.

There are two ways of setting a file's optimization flags. You can use the **CODE4** default settings when the file is opened or you can explicitly set them.

## Changing The Default Settings

When the file is opened, its optimization flags are set from the **CODE4.optimize** and **CODE4.optimizeWrite** flag settings. Valid settings for **CODE4.optimize** are:

- **OPT4EXCLUSIVE** (Default) Read optimize files when the files are opened exclusively or when the read-only attribute is set. Otherwise do not read-optimize the file.
- **OPT4OFF** Do not read optimize the file.
- **OPT4ALL** Same as **OPT4EXCLUSIVE** except that shared files are also optimized.

Valid settings for **CODE4.optimizeWrite** are:

- **OPT4EXCLUSIVE** (Default) Write optimize files when they are opened exclusively. Otherwise do not write optimize.
- **OPT4OFF** Do not write optimize.
- **OPT4ALL** Same as **OPT4EXCLUSIVE** except that shared files that are locked are also write optimized.



### Note

If read optimization is disabled, then write optimization is also disabled.

The following code segment opens three data files and changes their optimization flag settings. When memory optimization is activated, the first data file is read optimized, the second is read/write optimized, and the third has no optimization at all.

```
Function OptTest;
var
  cb: CODE4;
  db1, db2, db3: DATA4;
begin
  cb := code4init;

  code4optimize( cb, OPT4ALL ) ;
  code4optimizeWrite( cb, OPT4OFF ) ;
  db1 := d4open( cb, 'DATA1' ) ;

  code4optimizeWrite( cb, OPT4ALL ) ;
  db2 := d4open( cb, 'DATA2' ) ;

  code4optimize( cb, OPT4OFF ) ;
  code4optimizeWrite( cb, OPT4OFF ) ;
  db3 := d4open( cb, 'DATA3' ) ;

end;
```



---

## Overriding The Default Settings

A file's optimization flags can be changed at any time by calling one of the following functions: **d4optimize** or **d4optimizeWrite**. The **d4optimize** and **d4optimizeWrite** change the optimization flags of a data file and all of its open index and memo files. Following is the above example converted to use the **d4optimize** and **d4optimizeWrite** functions.

```
Function OptTest;
var
  cb: CODE4;
  db1, db2, db3: DATA4;
begin
  cb := code4init;

  db1 := d4open( cb, 'DATA1' ) ;
  d4optimize( db1, OPT4ALL );
  d4optimizeWrite( db1, OPT4OFF );

  db2 := d4open( cb, 'DATA2' ) ;
  d4optimize( db2, OPT4ALL );
  d4optimizeWrite( db2, OPT4ALL );

  db3 := d4open( cb, 'DATA3' ) ;
  d4optimize( db3, OPT4OFF );
  d4optimizeWrite( db3, OPT4OFF );

end;
```

---

## Activating The Optimizations

When the **code4optStart** function is called, the files whose optimization flags are set are memory optimized. Memory optimization is disabled by **code4optSuspend**.

```
code4optStart( cb ) ;
. . .
code4optSuspend( cb ) ;
```

---

## Refreshing The Buffers

You can also force a refresh of the optimization buffers. The **d4refresh** causes any buffered portion of the data, index or memo file to be discarded and reread into the buffer from disk. Finally, the **d4refreshRecord** rereads the current record from the disk.

---

## Memory Requirements

When using memory optimization, you can limit the amount of memory which CodeBase uses for this purpose by setting **code4memStartMax**. The natural question is what is an appropriate maximum? In general, the best maximum is the amount of memory which is likely to be available. For more information please refer to the CodeBase Members and Functions chapter of the *Reference Guide*.

Since most applications run under a number of hardware environments where varying amounts of memory are available, the application should assume that all extra available memory was used by CodeBase memory optimization. Therefore, the application should allocate its own memory before calling **code4optStart**, or after calling **code4optSuspend**. If you must allocate memory when the optimization is active, use **u4allocFree**. If this function fails to allocate memory, it will free memory from the CodeBase memory buffers and try again. These techniques make CodeBase memory optimization a lower priority.

If you expect very little memory to be available for CodeBase memory optimization, you should probably just not use it.

---

## When To Use Memory Optimization

---

Using memory optimization is not especially difficult. The most difficult part about memory optimization is knowing when to use it.

### Single User

The single user case is the most straight forward. Essentially, a single user application can safely memory read and write optimize all files. If an application explicitly flushes to disk by calling CodeBase flushing functions, write optimizations are ineffective. In all other single user cases write optimization is useful for improving performance.

When writing single user applications with memory optimization, it is a good idea to set **CODE4.accessMode** to **OPEN4DENY\_RW**. This way CodeBase performs memory optimizations by default.

---

### Multi-User

Whether or not to use memory optimization in a multi-user or multi-tasking environment is a more difficult decision. This is because memory optimization interferes with the multi-user sharing of information. On the other hand, the speed improvements resulting from the use of memory optimization can be even more dramatic because accessing a network server can be slower than accessing a local drive. In addition, memory optimization causes the server to be used less which can improve response time for other users.

Using memory read optimization in a network environment means that previously read database information is buffered in memory. The next time the information needs to be read, it can be quickly fetched from local memory. The disadvantage of this is that if the information has been changed by another application, the most recent piece of disk information may not be retrieved. At worst, one half of a read record could be an up-to-date version and the other half an older version. If read optimization is being used on memo files being updated by another user, it is theoretically possible to be returned an old or garbled memo entry. If the file is locked, using read optimization is completely safe because it cannot be updated by another user. Consequently, it is necessary to be careful to only use read optimization under appropriate circumstances.

Using memory write-optimization on shared files is potentially even more hazardous than using read-optimization. This is because information is written to disk only when the memory buffers are full. Consequently, from the perspective of any user reading the changed information, the information can appear as corrupt. If other applications are using index files which appear corrupt, they can generate errors. It is not quite as bad as it might sound because CodeBase is programmed with extensive error checking and reacts appropriately to most apparent corruption.



**WARNING**

Allowing one application to write optimize an index file, while another application may use the same index file, can lead to unpredictable results in the application reading the index file.

**Note**

CodeBase only allows write optimization when the entire data file is locked or when it is opened exclusively. This restriction is necessary in order to guarantee that index and memo files are not corrupt after they have been flushed.

For more examples of memory optimization, refer to the "Optimization" section of the "Multi-user Applications" chapter.

100 CodeBase

---

## 8 Multi-User Applications

---

A multi-user application can take many forms. For example, several people could be entering data, over a local area network, into the same data file at the same time. Another possibility is one person entering data while several others look at the data file.

When using CodeBase, you have many options in how you design the multi-user aspects of your application. For example, you can lock data areas before you read them or you can read them without locking. You can choose memory optimizations to improve performance or you can write/read directly to/from disk in order to ensure information is current. The exact options you choose depend on the requirements of the application and hardware resources available.

---

### Locking

At the heart of multi-user applications is locking. Locking is a way in which multi-user database applications communicate with each other. When an application locks some data, it is telling other applications "you cannot modify this data". When data is locked, other applications can still read the data. However, no other application can lock or modify the data.

Operations that  
require locking

Locking is automatically performed by CodeBase before data is written to disk. It is necessary to lock data before it is written to keep two applications from updating the same data at the same time. This avoids the corruption due to several applications updating the same index or memo file at the same time.

When a field is to be modified, the record should be locked to prevent multiple users from changing the same record at the same time. This principle is enforced when **CODE4.lockEnforce** is true (non-zero) and the field is modified with a field function or **d4blank**, **d4changed**, **d4delete** or **d4recall**.

Sometimes it is appropriate for an application to lock data before reading it. This is done to keep other applications from updating the data while the lock is present.

---

**Supported  
Locking  
Protocols**

CodeBase supports a variety of locking protocols that allow applications to be multi-user compatible with applications built using other products. When you use a CodeBase library that uses a specific type of index file compatibility, you also automatically get the multi-user locking compatibility for the DBMS that uses those types of files. The supported locking protocols are as follows:

- **FoxPro** (This is the default locking protocol). This locking protocol is employed when using a FoxPro compatible CodeBase library. This provides multi-user compatibility with FoxPro 2.0/2.5.
- **dBASE IV** The locking protocol is compatible with dBASE IV when a dBASE IV compatible CodeBase library is used.
- **Clipper** Clipper 5.2 locking compatibility is provided when using one of the Clipper compatible CodeBase libraries.

The product versions listed above were the versions which were multi-user compatible with CodeBase when this document was written. CodeBase may be updated to support additional product versions. Check the '\$READ.ME' file on your CodeBase diskette to determine exactly which versions are currently supported.

---

**Types Of  
Locking**

CodeBase performs several types of locking, although the only locks that you need to be concerned with are record and file locks. The types of locks are listed as follows:

- **Record Locking** When a record is locked, that data file record cannot be updated by other applications. This is the lowest level of locking (ie. you cannot lock fields).
- **Data File Locking** When a data file is locked, no records in the data file may be updated by other applications. In addition, a data file lock means that no other application may append records while the data file lock is in place.
- **Index and Memo Locking** CodeBase often locks and unlocks index and memo files when they are updated. However, you do not need to be concerned about this since it is automatic.

---

**Creating  
Multi-User  
Applications**

Creating well behaved multi-user applications, that is applications that do not lock portions of files for long lengths of time, is relatively easy using CodeBase. If you use the default **CODE4** flag settings, there are only a few things you must consider when writing multi-user applications.

---

**Recommended  
CODE4 Flag  
Settings**

The following **CODE4** flag settings are often appropriate when you are writing multi-user applications. Unless otherwise specified, the discussions in the following sections assume that these settings are being used. For details on the effects of changing these settings, please refer to last sections of this chapter.

- **CODE4.accessMode = OPEN4DENY\_NONE** (default) Files are opened in non-exclusive mode. This means that other applications can share the files with your application and have read and write access.
  - **CODE4.readOnly = 0** (default) Opens files in read/write mode. This allows you to read and write records to and from the data file.
  - **CODE4.readLock = 0** (default) There is no automatic record locking when a record is read.
  - **CODE4.lockAttempts = WAIT4EVER** (default) If a lock fails, CodeBase will keep retrying the lock until it succeeds.
  - **CODE4.lockEnforce = 1** An error is generated if an attempt is made to modify an unlocked record with a field function or **d4blank**, **d4changed**, **d4delete** or **d4recall**. This member variable must be set explicitly set to true (non-zero), since the default setting is false (zero).
- 

**Automatic  
Record  
Locking**

Since locking and unlocking are time consuming operations (in the same order of magnitude as a write to disk), CodeBase functions that write a record to disk lock that record without unlocking it afterwards. This prevents redundant unlocking calls and allows you the option of leaving the record locked.

The only functions that modify records and perform automatic locking are: **d4append**, **d4appendBlank**, **d4flush**, **d4flushRecord** and **d4write**.

Normally, you do not want to leave the record locked after these operations. To unlock the record, a call can be made to **d4unlock**:

```
d4append( db ) ;
d4unlock( db ) ;
```

---

## Automatic Data File Locking

In addition to functions that automatically lock a data file record, there are CodeBase functions that automatically lock the entire data file. These are **d4memoCompress**, **d4pack**, **d4reindex**, **d4zap** and **i4reindex**. It is strongly recommended that not only the data file be locked, but that the file be opened exclusively before performing these operations. Please refer to the section on **CODE4.accessMode** in the CodeBase Reference Guide.

These functions leave the data file locked after they finish executing. As a result, if the data file was opened non-exclusively, these functions should be immediately followed by a call to **d4unlock**.

## Automatic Unlocking Of Records

As a rule, CodeBase functions that move from an old record to a new record automatically remove any locks on the open data file according to **code4unlockAuto**. These functions are **d4bottom**, **d4go**, **d4goEof**, **d4positionSet**, **d4seek**, **d4seekDouble**, **d4seekN**, **d4seekNext**, **d4seekNextDouble**, **d4seekNextN**, **d4skip**, and **d4top**. Refer to **code4unlockAuto** in the *Reference Guide* for more information on how the automatic unlocking works.

The only exception to this rule is when the new record is already locked. In that case, no unlocking is performed.

## Common Multi-User Tasks

If you follow the advice about calling **d4unlock** after calling functions that automatically lock records and data files and are aware of functions that perform automatic unlocking, you should have little difficulty when writing multi-user applications.

To help you write multi-user applications, examples of common tasks performed by multi-user applications are provided below.



### PROGRAM Multi

Application Multi, which is perhaps simplistic, illustrates some basic operations which are present in almost any multi-user application: adding a record, modifying a record, finding a record, and reporting. It handles its multi-user aspects in a manner which is appropriate for many applications.

To run this application on a network, copy the files NAMES.\* from the samples directory to a network drive and modify the path in **d4open** to point to the network path. A multi-user scenerio can be simulated by running two instances of the application, in which case SHARE.EXE must be loaded first.

Multi assumes that a data file "NAMES.DBF" is present along with a production index file containing a tag named "NAME". Data file "NAMES.DBF" also contains a field named "NAME".

```
program MULTI ;

uses CodeBase, WinCrt, Sysutils ;

var
  cb : CODE4 ;
  field_name : FIELD4 ;
  data_names : DATA4 ;
  tag_name : TAG4 ;
  command : char ;
  b : array[0..99] of char ;
```



```

Procedure add_record ;
var
    oldAttempts: Integer;
begin
    writeln( 'Enter New Record' ) ;
    readln( b ) ;

    oldAttempts := code4lockAttempts( cb, 1 );
    if d4lock( data_names , d4recNo( data_names ) ) = r4success then
    begin
        d4appendStart( data_names, 0 ) ;
        f4assign( field_name, b ) ;
        d4append( data_names ) ;

        d4unlock( data_names ) ;
    end;
    code4lockAttempts( cb, oldAttempts );
end ;

Procedure find_record ;

begin
    writeln( 'Enter Name to Find' ) ;
    readln( b ) ;

    d4tagSelect( data_names, tag_name ) ;
    d4seek( data_names, b ) ;
end ;

Procedure modify_record ;
var
    oldAttempts: Integer;
begin
    writeln( 'Enter Replacement Record' ) ;
    readln( b ) ;
    oldAttempts := code4lockAttempts( cb, 1 );
    if d4lock( data_names , d4recNo( data_names ) ) = r4success then
    begin
        f4assign( field_name, b ) ;

        d4flush( data_names ) ;
        d4unlock( data_names ) ;
    end;
    code4lockAttempts( cb, oldAttempts );
end ;

Procedure list_data ;
var
    oldOpt: Integer;
begin
    code4optStart( cb ) ;
    oldOpt := d4optimize( data_names, OPT4ALL ) ;

    d4tagSelect( data_names, tag_name ) ;

    d4top( data_names ) ;
    while d4eof( data_names ) = 0 do
    begin
        writeln( d4recno( data_names ):4, ' ', f4str( field_name ) ) ;
        d4skip( data_names, 1 ) ;
    end ;
    d4optimize( data_names, oldOpt ) ;
    code4optSuspend( cb ) ;
end ;

begin
    cb := code4init ;

    code4accessMode( cb, OPEN4DENY_NONE ) ;
    code4readOnly( cb, 0 ) ;
    code4readLock( cb, 0 ) ;
    code4lockAttempts( cb, -1 ) ;
    code4lockEnforce( cb, 1 );

    data_names := d4open( cb, 'NAMES' ) ;
    error4exitTest( cb ) ;

    field_name := d4field( data_names, 'NAME' ) ;
    tag_name := d4tag( data_names, 'NAME' ) ;

```

```

d4top( data_names ) ;

while True do
begin
  code4errorCode( cb, 0 ) ;

  writeln( 'Record #: ', d4recno( data_names ), '   Name: ', f4str( field_name ) ) ;

  writeln( 'Enter Command ("a","f","l","m" or "x")' ) ;

  command := ReadKey ;
  case command of
    'a' : add_record ;
    'f' : find_record ;
    'l' : list_data ;
    'm' : modify_record ;
    'x' :
      begin
        code4close( cb ) ;
        Halt ;
      end ;
  end ;
end ;

code4initUndo( cb ) ;

readKey ;
DoneWinCrt ;
end.

```

---

**Opening Files** It is especially important to check for any file open errors in multi-user applications because there is a chance that the file might be opened exclusively by another application.

Code fragment  
from Multi

```

data_names := d4open( cb, 'NAMES' ) ;
error4exitTest( cb ) ;

```

---

## Control Loops

Most data file editing software has some kind of loop where the end user can enter various editing and possibly reporting commands. The Multi application is no exception. It allows you to 'add', 'find', 'list', 'modify' or 'exit'.

The start of the loop sets the CodeBase error code to zero. An error can occur if the user tries to modify a record before any have been added.

```

d4top( data_names ) ;

while True do
begin
  code4errorCode( cb, 0 ) ;

  writeln( 'Record #: ', d4recno( data_names ), '   Name: ',
    f4str( field_name ) ) ;

  writeln( 'Enter Command ("a","f","l","m" or "x")' ) ;

  command := ReadKey ;
  case command of
    'a' : add_record ;
    'f' : find_record ;
    'l' : list_data ;
    'm' : modify_record ;
    'x' :
      begin
        code4close( cb ) ;
        Halt ;
      end ;
  end ;
end ;

rc = d4top( db )

```

## Adding Records

In a single user situation, new records may simply be appended with a call to **d4append**. The multi-user situation is exactly the same, except that after appending the record **d4unlock** should be called (as in the *addRecord* function above). This call is necessary because the newly appended record remains locked after **d4append** completes.

```
Procedure AddRecord( dbf: DATA4; field: FIELD4 );
var
  buf: array[0..20] as Char;
begin
  StrPCopy( buf, InputBox( 'Record', 'Enter new
    record', 'New!' ) );

  d4appendStart( dbf, 0 );
  f4assign( field, buf );
  d4append( dbf );

  d4unlock( dbf );
end;
```

## Finding Records

The *findRecord* procedure is significant, from a multi-user perspective, more for what it does not do rather than what it does. Specifically, there is no multi-user logic necessary in this function. When **CODE4.readLock** is false (zero), the data file functions do not perform any automatic locking as the data file is being read. Accordingly, you can search using index files and read data file records without having any extra multi-user logic present.

Code fragment  
from Multi

```
Procedure find_record ;
begin
  writeln( 'Enter Name to Find' );
  readln( b );

  d4tagSelect( data_names, tag_name );
  d4seek( data_names, b );
end ;
```

## Modifying Records

The field functions are used to assign values to the fields, and when appropriate, the changes are flushed to disk. Explicit flushing can be accomplished by calling **d4flush**.

After flushing a modified record in a multi-user situation, an application should call **d4unlock**. When **d4flush** is called, the changed record gets written to disk. In order to accomplish this, **d4flush** locks the record and leaves it locked. Consequently, the call to **d4unlock** is suggested once the modified record is no longer required, in order to give other users an opportunity to modify the record.

The call to **d4flush** is not strictly necessary. This is because **d4unlock** is smart enough to recognize when the record buffer has changed. In this case, **d4unlock** will lock the record, flush the record, and then unlock the record.

Code fragment  
from Multi

```

Procedure modify_record ;
var
  oldAttempts: Integer;
begin
  writeln( 'Enter Replacement Record' ) ;
  readln( b ) ;
  oldAttempts := code4lockAttempts( cb, 1 );
  if d4lock( data_names , d4recNo( data_names ) ) = r4success then
  begin
    f4assign( field_name, b ) ;

    d4flush( data_names ) ;
    d4unlock( data_names ) ;
  end;
  code4lockAttempts( cb, oldAttempts );
end ;

```

Locking the  
Record before  
modifying

In the MULTI.PAS example, the record must be explicitly locked before it is modified by a field function, since the **CODE4.lockEnforce** member is set to true (non-zero). This serves to ensure that only one application can edit a record at a time. The example sets the **CODE4.lockAttempts** to 1, to cause **d4lock** to immediately return **r4locked** if the record is already locked by another user, in which case the *modifyRecord* function is aborted. Adhering to this locking procedure will prevent the following undesirable scenario from occurring.

Consider the case where two users of the application decide to modify the same record at the same time. When this happens, one application may write out its changes to disk before the second. The second application then writes to the disk, overwriting the first application's changes. Neither application knows that this has happened. User two makes changes based on an out of date version of the record, and user one loses all changes.

---

### Multi-User Optimizations

This next section deals with using the memory optimization functions in a multi-user application. The optimizations can be used for both appending large quantities of records and for efficiently generating lists of records.

---

### Listing Records

Function *listData* from the program Multi is an example of reporting. Notice that memory optimization is turned on at the start of the routine and turned off at the end of the routine. This speeds up the report. In addition, when a network is being used it minimizes requests to the server.

In a multi-user application, the use of memory optimization is recommended when reporting but is strongly discouraged when editing. Refer to the memory optimization chapter for additional information.

Notice that there are no unlocking function calls. As explained under Finding Records section (above), when a record is read using the data file functions no automatic locking takes place. Consequently there is no need for any unlocking.

Code fragment  
from Multi

```

Procedure list_data ;
var
  oldOpt: Integer
begin
  code4optStart( cb ) ;
  oldOpt := d4optimize( data_names, OPT4ALL ) ;

  d4tagSelect( data_names, tag_name ) ;

  d4top( data_names ) ;
  while d4eof( data_names ) = 0 do
  begin
    writeln( d4recno( data_names ):4, ' ',
             f4str( field_name ) ) ;
    d4skip( data_names, 1 ) ;
  end ;
  d4optimize( data_names, oldOpt ) ;
  code4optSuspend( cb ) ;
end ;

```

## Repeated Appending



### PROGRAM Append

The next example multi-user application demonstrates how to append records from one data file to another.

The program Append demonstrates how to append records from one data file to another. It assumes that data files "TO\_DBF.DBF" and "FROM\_DBF.DBF" are both present and that they both have a field named "INFO".

```

program APPEND ;

uses CodeBase, WinCrt;

var
  cb : CODE4 ;
  data_from, data_to : DATA4 ;
  info_from, info_to : FIELD4 ;
  rc1, rc2, rc : integer ;

begin
  cb := code4init ;

  code4optimize( cb, OPT4ALL ) ;
  code4optimizeWrite( cb, OPT4ALL ) ;
  code4memStartMax( cb, 262140 ) ;

  data_from := d4open( cb, 'FROM_DB.DBF' ) ;
  data_to := d4open( cb, 'TO_DB.DBF' ) ;
  error4exitTest( cb ) ;

  code4optStart( cb ) ;

  info_from := d4field( data_from, 'F_NAME' ) ;
  info_to := d4field( data_to, 'F_NAME' ) ;

  code4lockAttempts( cb, 1 ) ;
  rc1 := d4lockFile( data_from ) ;
  rc2 := d4lockFile( data_to ) ;
  if (rc1 <> 0) or (rc2 <> 0) then
  begin
    writeln( 'Locking Failed' ) ;
    Halt ;
  end ;

  rc := d4top( data_from ) ;
  while rc = 0 do
  begin
    d4appendStart( data_to, 0 ) ;
    f4assignField( info_to, info_from ) ;
    d4append( data_to ) ;
    rc := d4skip( data_from, 1 ) ;
  end ;

  d4unlock( data_from ) ;
  d4unlock( data_to ) ;

  code4close( cb ) ;
  code4initUndo( cb ) ;

```

```

readKey ;
DoneWinCrt ;
end.

```

In this application, both memory read optimization and memory write optimization are used. In order to do this, the **CODE4.optimize** and **CODE4.optimizeWrite** defaults are changed before the files are opened. Please refer to chapter 7 for details.

Code fragment  
from Append

```

code4optimize( cb, OPT4ALL ) ;
code4optimizeWrite( cb, OPT4ALL ) ;
code4memStartMax( cb, 262140 ) ;

data_from := d4open( cb, 'FROM_DB.DBF' ) ;
data_to   := d4open( cb, 'TO_DB.DBF' ) ;
error4exitTest( cb ) ;

code4optStart( cb ) ;

```

Note that **code4optStart** is called *after* the files are opened. This is because a call to **code4optStart** may use up the available memory. In this case, the calls to **d4open** could be slowed down as optimization was repeatedly suspended and re-invoked inside **d4open**. Internally, **d4open** calls **u4allocFree** when it allocates memory. When memory is not available, **u4allocFree** temporarily suspends memory optimization in an attempt to allocate the requested memory.

This example illustrates the only situation in which write optimization should be considered in a multi-user application. When writing to a file with write optimizations enabled, the file should be locked.

With the file locked and write optimization enabled, the repeated appending goes considerably faster. The trade off is that if any other application attempts to read the records being appended, the other application could, at worst, generate CodeBase errors or could read garbage information. To avoid this, you can open the files exclusively or do not use write optimization.

In this specific case, the read optimization is a very good idea. This is because there is no chance that information being returned could be out of date because the file is locked. The speed improvements in this example, as a result of the read optimization, are considerable.

This application skips sequentially through data file "FROM\_DBF" using memory optimization. Consequently, it is important to use function **d4skip** rather than **d4go** because function **d4skip** detects the sequential reading and does special performance optimizations.

Since the data files are locked going into the loop, they stay locked throughout the entire loop with no automatic unlocking occurring.

Code fragment  
from Append

```
rc := d4top( data_from ) ;
while rc = 0 do
begin
  d4appendStart( data_to, 0 ) ;
  f4assignField( info_to, info_from ) ;
  d4append( data_to ) ;
  rc := d4skip( data_from, 1 ) ;
end ;
```

Once the records have been appended, the files are unlocked and closed. When data file "TO\_DBF.DBF" is unlocked its changes are automatically flushed to disk.

Code fragment  
from Append

```
d4unlock( data_from ) ;
d4unlock( data_to ) ;

code4close( cb ) ;
```

The calls to **d4unlock** are not strictly necessary in this example because **code4close** does flushing and unlocking automatically. Alternatively, **code4unlock** could have been called to unlock both files at once.

---

## Avoiding Deadlock

When locking multiple data files at once, you need to be careful to avoid deadlock. Deadlock happens when one application waits for a second application to unlock something and the second application waits for the first application to unlock something else. Since they are both waiting for each other, they both wait forever. This program avoids deadlock by changing **CODE4.lockAttempts** from the default of unlimited retries to one try. If the lock attempt fails, the program exits.

Code fragment  
from Append

```
code4lockAttempts( cb, 1 ) ;
rc1 := d4lockFile( data_from ) ;
rc2 := d4lockFile( data_to ) ;
if (rc1 <> 0) or (rc2 <> 0) then
begin
  writeln( 'Locking Failed' ) ;
  Halt ;
end ;
```

Another way to avoid deadlock is to be careful to always lock data files in the same order. If one application locks data file "FROM\_DBF" before "TO\_DBF" using unlimited retries, then all applications using unlimited retries should do the same.

If applications always lock data files in the same order, deadlock is not possible. This is because once the first application succeeds in locking the first data file, other applications will wait for the first application to finish its locking, do what it needs to do and then unlock all of the data files. The best way to ensure you are always locking data files in the same order is to lock and unlock all data files together.

---

## Group Locks

Sometimes it is desirable to lock many items as a group. In this case, either all items are locked or no items are locked. This functionality will help minimize the chance of deadlock. CodeBase supplies functions that will put an item on a queue so that the whole queue can be locked as a group. The functions **d4lockAdd**, **d4lockAddAppend**, **d4lockAddFile**, **d4lockAddAll** and **relate4lockAdd** put their respective items on a queue to be locked. These functions do NOT lock any items. The items are locked by a call to **code4lock**. If any item in the queue fails to be locked, the successful locks are removed, **code4lock** returns **r4locked** and the queue remains intact for a future lock attempt by a subsequent call to **code4lock**. When **code4lock** succeeds, all the locks are in place and the queue is emptied. Refer to **code4lock** in the *Reference Guide* for more details.

---

## Exclusive Access

The simplest type of multi-user applications are those that open all of their files in exclusive access mode. If any other application tries to open a file which has been opened exclusively by another user, the application gets a file open error.

The big disadvantage of this method is that no other applications can use files that you are using. This is not a recommended method of creating multi-user applications.

The main use of opening files exclusively is when you are performing packing, zapping, indexing or reindexing. This prevents other applications from generating errors or garbage output if they use the file while one of these activities is occurring.

---

## Opening Files Exclusively

Whether a file is opened exclusively is determined by the current status of the **CODE4.accessMode** flag. This flag specifies what access OTHER users have to the current file. **CODE4.accessMode** has three possible values:

- **OPEN4DENY\_NONE** Open the database files in shared mode. Other users have read and write access. This is the default value.
- **OPEN4DENY\_WRITE** Open the database files in shared mode. The application may read and write to the files, but other users may only open the file in read only mode and may not change the files.
- **OPEN4DENY\_RW** Open the database files exclusively. Other users may not open the files.

This flag can be changed any time after the **CODE4** structure has been initialized with **code4init**. Changes to the flag only affect those files that are opened thereafter. Therefore if the access mode must be altered for an open file, the file must be closed and then reopened.



```
code4accessMode( cb, OPEN4DENY_RW ) ;
db := d4open( cb, 'DATAFILE' ) ;      {Open exclusively}

d4close( db ) ;

code4accessMode( cb, OPEN4DENY_NONE ) ;
db := d4open( cb, 'DATAFILE' ) ;      {Open in shared mode}
```

If a file cannot be opened exclusively due to the fact that some other application already has it open, the file open fails. When this happens, and the **CODE4.errOpen** flag is set to true, an error message is generated.

---

## Read Only Mode

If your application reads information from a file and never modifies it, you can open the file in read only mode. This mode must be used when you only have read access to a file. This is a common situation on a network where a database is shared among many users, but only a few users have the authority to make changes to it.

---

## Opening Files In Read Only Mode

Whether a file is opened in read only mode is determined by the current status of the **CODE4.readOnly** flag. If the flag is set to its default value of false (zero), files are opened in read / write mode. If this flag is set to true (non zero), files are opened in read only mode.

Code fragment  
from Multi

```
rc = code4readOnly( cb, 0 )
```

---

## Lock Attempts

When a lock is performed on a file, the possibility exists that some other application has already locked that file or a portion of the file. When this occurs, what CodeBase does next is determined by the status of the **CODE4.lockAttempts** flag.

The default value of **CODE4.lockAttempts** is **WAIT4EVER**. This indicates that CodeBase will keep trying to establish a lock until it succeeds. The advantage of this method is that it simplifies programming because the application can assume that all lock attempts succeed. Unfortunately, this can increase the chances of deadlock bugs being present in an application. In addition, if applications attempt locking for long periods of time, users can be left waiting.

Code fragment  
from Multi

```
code4lockAttempts( cb, 1 ) ; {Try lock once}
```

If **CODE4.lockAttempts** is set to any value greater than one, CodeBase keeps trying the lock at approximately one second intervals until either a lock is established or the number of attempts equals **CODE4.lockAttempts**. If the lock fails, a value of **r4locked** is returned from the function attempting the lock.

If you wish to perform a special operation if the lock fails, such as displaying a message to the user, you should set **CODE4.lockAttempts** to one and test for return values. A value of zero makes CodeBase functions return **r4locked** after the first lock attempt fails.

---

## Automatic Record Locking

When **CODE4.readLock** is true, several CodeBase functions automatically lock records before reading them. This ensures that no other application can modify a record that your application has in its record buffer.

Unfortunately, locking data file records when they are only being read can create unnecessary locking contention. Simply put, this would increase the chance of a record being locked when an end user just wants to look at it. Consequently, the default of no automatic read locking is appropriate for most multi-user applications.

Code fragment  
from Multi

```
code4readLock( cb, 0 );
```

The following list of functions perform record locking when the **CODE4.readLock** flag is set to true: **d4bottom**, **d4go**, **d4position**, **d4seek**, **d4seekDouble**, **d4seekN**, **d4seekNext**, **d4seekNextDouble**, **d4seekNextN**, **d4skip**, and **d4top**.

---

## Enforced Locking

Locking a record before it is modified is very important in the multi-user configuration, since it ensures that only one application can edit a record at a time. The **CODE4** member variable **CODE4.lockEnforce** can be used to ensure that an application has explicitly locked a record before it is modified with a field function or the following data file functions: **d4blank**, **d4changed**, **d4delete** or **d4recall**. When **CODE4.lockEnforce** is true (non-zero), an **e4lock** error is generated when an attempt is made to modify an unlocked record.

An alternative method of ensuring that only one application can modify a record at a time is to deny all other applications write access to a data file. Write access can be denied to other applications by passing **OPEN4DENY\_WRITE** or **OPEN4DENY\_RW** to **code4accessMode** before opening the file. Thus, it is unnecessary to explicitly lock the record before editing it, even when **CODE4.lockEnforce** is true (non-zero), since no other application can write to the data file. This method suffers from the same disadvantage as discussed in *Exclusive Access* section of this chapter. The restricted access to the files prevents the creation of practical multi-user applications.

---

## 9 Performance Tips

---

Programmers are always looking for ways to improve the performance of their applications. This chapter provides the programmer with tips on how to achieve the highest performance when building applications with CodeBase.

---

### Memory Optimization

The prudent use memory optimization can enhance the performance of an application. For details on when to use memory optimization refer the Memory Optimization chapter of this guide. Write optimization is most useful when making multiple updates on files, such as appending many records at once. In general, when memory optimization is used, **d4skip** is faster than **d4go**.

#### Memory Requirements

The **CODE4** member variable **CODE4.memStartMax** specifies the maximum amount of memory allocated for memory optimization. Theoretically, the more memory that CodeBase can use for memory optimization the greater the improvement in performance. In practise, if **CODE4.memStartMax** is too large, CodeBase may not be able to allocate all of the requested memory causing the memory optimization to work less efficiently. In some operating systems all requested memory is allocated; but as virtual memory, which is counter-productive to memory optimization. On the other hand, if **CODE4.memStartMax** too small, the potential benefits of memory optimization will not warrant the overhead.

#### Observing the performance improvement

The improvement in performance can be observed by using **code4optAll**. **code4optAll** ensures that memory optimization is fully implemented by locking and fully read and write optimizing all opened data, index and memo files. A comparison of the performance can then be made between the application that employs memory optimization with one that does not.

#### Functions that can reduce performance

Memory optimization acts by buffering portions of files in memory and thus decreases the number of disk accesses and improves performance. Therefore, frequent calls to CodeBase functions that read data from disk will negate the effects of read optimization. For this reason, functions such as **d4refresh** and **d4refreshRecord** should not be called repeatedly while memory optimization is invoked. Frequent flushing also reduces the benefits of write optimization.

---

### Locking

Locking can take as long a disk access, so repeated record locks are inefficient. It is more efficient to lock an entire file when multiple updates are necessary. The best way to lock a file that requires many modifications is to use **d4lockAddAll** with **code4lock**, or call **d4lockAll**.

A **CODE4** setting that can influence performance is **CODE4.readLock**. When **CODE4.readLock** is set to true (non-zero), each record is locked before it is read, which can reduce performance. It is useful to set **CODE4.readLock** to true when the records are to be modified.

Another important **CODE4** variable that can affect performance in a network environment is **CODE4.lockDelay**, which determines how long to wait between lock attempts. Consider this variable carefully because if the value is too small, it can cause an increase in network traffic and possibly a decrease network performance.

---

### Appending

Whenever possible, it is more efficient to append many records at once rather than one at a time. When appending many records, it is sometimes faster to close the index files first, append the records and then reindex. The files must be locked before the records may be appended and the most efficient method is to call **d4lockAddAll** and **code4lock**, or **d4lockAll**, as discussed above. Batch appending can also take advantage of write optimization, which will improve performance.

---

### Time Consuming Functions

Certain CodeBase functions are inherently time consuming and this aspect should be considered when incorporating them into an application. Avoid calling time consuming functions repeatedly, since this can reduce performance. Some functions are meant to be used for debugging purposes and should not be used in the final user application.

The following functions reindex index files, which is a time consuming procedure, so they should be used with care. **d4reindex** and **i4reindex** can be called to explicitly reindex files. Both **d4pack** and **d4zap** automatically reindex the associated index files when they are called.

**d4pack** and **d4zap** are time consuming functions even when there are no tags to reindex. These functions physically remove records from a data file and then reorganize the remaining records, which is a time consuming process.

**d4check** is a function that determines whether an index file has been corrupted, which is useful for debugging. This function can take as long as reindexing and therefore application performance can be hindered.

---

### Queries and Relations

CodeBase uses Query Optimization to greatly increase the performance of queries in relation sets. To ensure that the Query Optimization can be used by CodeBase, the query expression must have a corresponding tag expression. See the Query Optimization chapter in this guide for more details on how to ensure that the query optimization will be invoked.

Another important issue to consider when manipulating relations is how to specify the sort order of the query set. The most efficient manner will depend on the size of the database. There are three ways in which a sorted order may be specified for the query set.

1. The sorted order can be determined by the selected tag in the master data file.
2. If there is no selected tag for the master data file, then the natural order of query set is used.
3. The function **relate4sortSet** can be used to specify the sorted order.

If there is a tag that specifies the same sort order as the **relate4sortSet** expression, then use the tag to specify the sort order when the query set is almost the same size as the data file. If the query set is almost the same size as the data file, then **relate4sortSet** is less efficient when compared with the tag sorted order or the natural order. Use **relate4sortSet** when the query set is small compared to the size of the data file. This will result in better performance when compared with the selected tag ordering.

If there is no tag that specifies a desired sort order, then use **relate4sortSet** to sort the query set. Using **relate4sortSet** will be faster than creating a new tag to specify the sort order, regardless of the size of the query set.

---

## General Tips

The following section discusses miscellaneous issues that may influence the performance of an application.

- Be sure to use the field functions to manipulate fields for the best results. Avoid using the expression module to perform calculations on fields, otherwise the application will execute at a slower rate.
- It is recommended that all necessary tags be constructed when the index file is created with either **d4create** or **i4create**. This method is more efficient than adding new tags by calling **i4tagAdd** later in the application.
- File access is quicker when the file is opened exclusively, since no future record or file locks will be necessary.
- The **CODE4** member variables that determine how many memory blocks are allocated should be defined at the beginning of the program and should not be changed during the program. This allows CodeBase to share memory pools efficiently. Set the following **CODE4** settings at the beginning of the application.

**CODE4.memStartData**  
**CODE4.memStartIndex**  
**CODE4.memStartBlock**  
**CODE4.memStartLock**  
**CODE4.memStartMax**  
**CODE4.memStartTag**



# 10 Transaction Processing

At times it is necessary to have a group of actions executed as a unit. In this case, either all actions are completed or none of the actions are completed.

To illustrate this, consider the everyday example of transferring funds from one bank account to another. There are two steps to this process: first one account has the money debited and then the second account has the sum credited. A failure could occur after the debit but before the credit, in which case the money would be removed from the first account but the second account remains unchanged, resulting in lost funds. In this type of situation, one would like an all or nothing response, where either both the debit and credit are completed or nothing happens and the accounts remain unchanged.

## Transactions

A transaction is one way of treating a group of actions as a whole. The scope of a transaction is delimited by a specified "start" and "end" between which the actions are treated as a unit. A transaction has the all or nothing property, so that either all the actions within the scope of a transaction are completed or none are completed. When a transaction has successfully completed all the operations within its scope, it is "committed". A transaction can be aborted and all the changes that were made up until the failure can be reversed. This constitutes a "rollback". After the transaction has terminated, the changes are permanent and can not be reversed.



### PROGRAM Transfer

The following program shows how a simple bank transfer can be accomplished using a CodeBase transaction.

```
program TRANSFER;

uses CodeBase, WinCrt, Sysutils ;

var
  cb: CODE4 ;
  datafile: DATA4 ;

  acctNo, balance: FIELD4 ;
  acctTag, balTag: TAG4 ;

Procedure OpenLogFile ;           {Opening or creating a log file is }
                                   {only required for STAND-ALONE mode.}
var
  rc: Integer;
begin
  { NIL is passed as the log file name }
  {so the default 'C4.LOG' is used as }
  {the log file name. }
  code4errOpen( cb, 0 ) ;
  rc := code4logOpen( cb, nil, 'user1' ) ;
  if rc = r4noOpen then
  begin
    code4errorCode( cb, 0 ) ;
    rc := code4logCreate( cb, nil, 'user1' ) ;
  end;
end;

Procedure OpenDataFile ;
begin
  datafile := d4open( cb, 'BANK.DBF' ) ;

  acctNo := d4field( datafile, 'ACCT_NO' ) ;
  balance := d4field( datafile, 'BALANCE' ) ;

  acctTag := d4tag( datafile, 'ACCT_TAG' ) ;
  balTag := d4tag( datafile, 'BAL_TAG' ) ;
```

```

end;

Function Credit( toAcct: Longint; amt: double ): Integer;
var
    rc: Integer;
    newBal: Double;
begin
    rc := d4tagSelect( datafile, acctTag );
    rc := d4seekDouble( datafile, toAcct );

    if rc = r4success then
    begin
        newBal := f4double( balance ) + amt ;
        f4assignDouble( balance, newBal );
        rc := r4success;
    end;
    Credit := rc;
end;

Function Debit( fromAcct: Longint; amt: Double ): Integer;
begin
    Debit := Credit( fromAcct, -amt );
end;

Procedure TransferFunds( fromAcct: Longint; toAcct: Longint; amt: double );
var
    rc, rc1, rc2: Integer;
begin
    code4tranStart( cb );

    rc1 := Debit( fromAcct, amt );
    rc2 := Credit( toAcct, amt );

    if (rc1 = r4success) and (rc2 = r4success) then
        rc := code4tranCommit( cb )
    else
        rc := code4tranRollback( cb );
end;

Procedure PrintRecords ;
var
    rc: Integer;
begin
    rc := d4top( datafile );
    while rc = r4success do
    begin
        writeln('-----') ;
        writeln('Account Number: ', f4long( acctNo )) ;
        writeln('Balance      : ', Format('%m', [f4double( balance )])) ;
        rc := d4skip( datafile, 1 );
    end;
    writeln('=====') ;
end;

begin
    cb := code4init ;

    code4errOpen( cb, 0 );
    code4safety( cb, 0 );
    code4lockAttempts( cb, 5 );
    OpenLogFile ;

    OpenDataFile ;

    PrintRecords ;

    { The account number 56789 doesn't exist in the database,
    { the transfer is aborted and database is not changed      }

    TransferFunds( 12345, 56789, 200.00 );
    PrintRecords ;

    { Both accounts exist so the TransferFunds is completed
    { and the database is updated                                }

    TransferFunds( 12345, 55555, 150.50 );
    PrintRecords ;

    code4initUndo( cb );
end.

```



In the Transfer program, a transfer takes place when there is a debit from one account followed by a credit to another, as shown by the following function calls.

Code fragment for  
Transfer

```
rc1 := Debit( fromAcct, amt ) ;
rc2 := Credit( toAcct, amt ) ;
```

These two operations must both be successful in order for the transfer to succeed. Therefore, these two actions must be contained within a transaction where either both the debit and the credit succeed or neither will be completed and the database will not be changed. In CodeBase, a transaction is delimited by the functions **code4tranStart** and **code4tranCommit**, which initiate and terminate the transaction respectively.

```
code4tranStart( cb ) ;

{...These two statements are treated as one unit }

code4tranRollback( cb ) ;
```

Events may occur that cause failure during a transaction. For example, if there is a power failure or the hardware crashes during a transaction, one would want to be able to reverse any changes that had occurred before the failure. Programmers must anticipate and test for possible failure points within a transaction. For instance, the example program Transfer checks to see if both specified accounts exist in the database and a failure occurs if either account did not exist. When the program encounters a failure of this sort, it calls **code4tranRollback** to abort the transaction and restore the database to its original state.

If the transfer succeeds, the transaction is completed by **code4tranCommit**, which commits the changes to the database. After the transaction has been committed, **code4tranRollback** can NOT be used restore the database to the state that existed before the transaction was started. An error is generated if **code4tranRollback** is called after **code4tranCommit**.

Code fragment  
from Transfer

```
if (rc1 = r4success) and (rc2 = r4success) then
    rc := code4tranCommit( cb )
else
    rc := code4tranRollback( cb ) ;
```

## How a Transaction Works

CodeBase has implemented transactions in the following manner. When a transaction is initiated with **code4tranStart**, all the changes to the database are recorded in a log file. The log file stores both the old value and the new value, as well as who made the change. At every step during a transaction, each change is recorded in the log file and then the database is modified. After the transaction has been committed using **code4tranCommit**, the changes to the database are permanent. During a rollback, **code4tranRollback** uses the original values from the log file to reverse all the changes made to the database.

## Logging in the Stand-Alone Case

All modifications that are made while an application is running, which are outside the scope of a transaction, can be automatically recorded in the log file. Automatic logging is useful when a back up copy of all changes is required. There are times when logging is not necessary; for example, when copying data files. In the stand-alone configuration, the logging can be turned off to save time and disk space. The following discussion applies only to logging in the stand-alone configuration.

**code4log** The **CODE4** member variable **CODE4.log** is used to specify whether changes to the data files are automatically recorded in a log file. This setting can be changed by using the function **code4log**. Changing this setting only affects the logging status of the files that are opened subsequently. The files that were opened before the setting was changed retain the logging status that was set when the file was opened. **CODE4.log** has three possible values:

- **LOG4ALWAYS** The automatic logging takes place for all the data files for as long as the application is running. The logging can NOT be turned off for the current data file by calling **d4log**.
- **LOG4ON** The automatic logging takes place for all the data files for as long as the application is running. In this case, **d4log** can be used to turn the logging off and on for the current data file. CodeBase uses this value as the default.
- **LOG4TRANS** Only the changes made during a transaction are automatically recorded in the log file. Logging may be turned on and off for the current data file by calling **d4log**.

**d4log** When a file is opened with **CODE4.log** set to either **LOG4ON** or **LOG4TRANS**, **d4log** can be used to turn the logging off and on for the current data file. **d4log** only has an effect on logging while the data file is open and it has no effect on the logging that is done during a transaction.

Pass false (zero) to **d4log**, to turn the logging off and pass it true (non-zero) to turn the logging back on. **d4log** also returns the previous logging setting. If -1 is passed to **d4log** the current logging status is returned. The possible return values are as follows:

- **r4logOn** This return value means that logging is turned on for the data file.

- **r4logOff** This return value means that either the logging is turned off for the data file or that logging in general is not enabled.
- **<0** An error has occurred.

turning logging off

```
{ Logging of changes takes place
  for all data files opened subsequently }

rc := code4log( cb, LOG4ON ) ;

datafile := d4open( cb, 'DATA.DBF' ) ;

rc := d4log( datafile, 0 ) ;
      { logging is turned off for this particular data
        file and the previous logging status is returned }
rc := d4log( datafile, -1 ) ; { the current logging status is returned }
```

## Log files

In the stand-alone configuration, a log file must exist or be explicitly created before any automatic logging or transaction logging can take place. If a log file does not exist, **code4tranStart** will generate an error.

The log file may be manually opened or created by using **code4logOpen** and **code4logCreate** respectively. **code4logOpen** and **code4logCreate** both take the **CODE4** pointer, the name of the log file and user identification as parameters. If a zero length string is passed as the file name, then the default file name "C4.LOG" is used. **code4logOpen** and **code4logCreate** must be called before **d4open** or **d4create**.

Generally, one would like to open a log file if it exists, otherwise create a new log file. This concept is the same as discussed in the "Opening or Creating" section of the Database Access chapter in this guide.

Code fragment  
from Transfer

```
Procedure OpenLogFile ;
var
  rc: Integer;
begin
  code4errOpen( cb, 0 ) ;
  rc := code4logOpen( cb, nil, 'user1' ) ;
  if rc = r4noOpen then
  begin
    code4errorCode( cb, 0 ) ;
    rc := code4logCreate( cb, nil, 'user1' ) ;
  end;
end;
```

If **code4logOpen** or **code4logCreate** have not been explicitly called, **d4open**, **d4create** and **code4tranStart** automatically try to open a log file by calling **code4logOpen**. **code4logOpenOff** can be used to instruct **d4open**, **d4create** and **code4tranStart** not to automatically open the log file. When **code4logOpenOff** is used, no transactions can start unless the log file is opened explicitly by the application.

the log file is not  
automatically  
opened

```
rc := code4logOpenOff( cb ) ;
datafile := d4open( cb, 'DATA.DBF' ) ;
```

## Logging in the Client-Server case

In the client-server configuration, the automatic logging and transaction logging are handled by the server. Calling the functions **code4logCreate**, **code4logOpen** and **code4logOpenOff** in a client application will have no effect, since the log file is controlled by the server. These functions will always return **r4success** when called by a client application and will have no effect on the log file.

## Locking

Every time a data file is being modified one must consider which locking procedure is appropriate. Locking is necessary step in preserving the integrity of the database. CodeBase has both automatic and explicit locking features, as discussed in the Multi-user Applications chapter in this guide.

### Automatic locking during a transaction

While a transaction is in progress, any automatic locking that is required during a transaction is performed, but CodeBase acts as though the **code4unlockAuto** is set to **LOCK4OFF**, so no automatic unlocking will occur.

Automatic locking and unlocking can occur when a transaction is committed or rolled back. In both cases, record buffer flushing may be required. The locking and unlocking procedure follows that of **d4flush**. If a new lock is needed, everything is unlocked according to **code4unlockAuto** and then the lock is placed. Otherwise, if no new locks are required then nothing is unlocked. It may be necessary to explicitly unlock files after the transaction is committed or rolled back depending on the **code4unlockAuto** setting.

### Automatic locking vs. Explicit locking

Whether using the automatic or explicit locking one must decide on a value for the **CODE4** member variable **CODE4.lockAttempts**. This member has a default value of **WAIT4EVER**, which may result in deadlock in a multi-user environment. To prevent deadlock, set **CODE4.lockAttempts** to a reasonable finite number.

The example program Transfer sets **CODE4.lockAttempts** to a finite value and takes advantage of the automatic locking properties of CodeBase. When **d4seek** is called during the transaction, the current record may need to be flushed before the new record can be loaded into the recorded buffer. Normally, if a new lock is required, everything is unlocked according to **code4unlockAuto** and then the lock is placed, but the seek takes place during a transaction so no unlocking is performed.

### Code fragment from Transfer

```
code4lockAttempts( cb, 5 ) ;
```

When an application is running in a multi-user environment, it is strongly recommended that records be locked before they are modified to prevent more than one user access to the record at the same time. In this case, it is prudent to set **CODE4.lockEnforce** to true ( non-zero) to ensure that a record is explicitly locked before it is modified. If it is known before hand that many changes will be made, use **d4lockAll**, or use **d4lockAddAll** followed by **code4lock** to explicitly lock the files.

If the file is explicitly locked, it must be explicitly unlocked by calling **code4unlock** or **d4unlock** after the transaction is committed or rolled back. If an attempt is made to unlock files during a transaction, an error is generated.



# 11 Index

—.—

.CDX. *See* Index Files  
 .CGP. *See* Group Files  
 .DBF. *See* Data Files  
 .MDX. *See* Index Files  
 .NTX. *See* Index Files

—A—

Appending Records. *See* Records  
 Automatic Opening  
   index files. *See* Index Files, production indexes

—B—

Buffering. *See* Memory Optimizations

—C—

CCYYMMDD. *See* Date Operations  
 Character Fields  
   assigning values to, 27  
   creating, 24  
   retrieving contents, 21  
 Clipper  
   index files, 34  
   locking protocol, 100  
 CODE4, 13, 17  
   Flag setting functions  
     code4accessMode, 17  
     code4autoOpen, 41, 54  
     code4errDefaultUnique, 46, 47, 83  
     code4errOpen, 26, 111  
     code4lockAttempts, 101, 106, 109, 111, 122  
     code4lockDelay, 114  
     code4lockEnforce, 27, 99, 101, 106, 112, 123  
     code4log, 120  
     code4memStartMax, 113, 115  
     code4optimize, 94, 95, 108  
     code4optimizeWrite, 94, 108  
     code4readLock, 101, 105, 112, 113  
     code4readOnly, 101, 111  
     code4safety, 25  
   initializing, 17  
   multi-user settings, 101  
   uninitializing, 18  
 CodeBase  
   constants, 12  
   functions, 11  
   structures, 12, 13  
 CodeBase Functions  
   code4close, 18, 109  
   code4init, 17, 25, 110

code4initUndo, 18  
 code4lock, 110, 113, 114, 123  
 code4logCreate, 121, 122  
 code4logOpen, 121, 122  
 code4logOpenOff, 121, 122  
 code4optStart, 95, 108  
 code4optSuspend, 95  
 code4tranCommit, 119, 120  
 code4tranRollback, 119, 120  
 code4tranStart, 119, 120, 121  
 code4unlockAuto, 102, 122  
 Composite Data Files. *See* Relations  
   moving between composite records, 71  
 Composite Record. *See* Relations  
 Constants  
   CodeBase constants, 12  
 Converting Date Formats. *See* Date Operations  
 Current Record. *See* Data Files

—D—

Data File Functions  
 d4alias, 32  
 d4aliasSet, 32  
 d4append, 26, 101, 105  
 d4appendBlank, 26, 101  
 d4appendStart, 26  
 d4bottom, 18, 43, 79, 102, 112  
 d4close, 16, 18  
 d4create, 24, 25, 39, 53, 115, 121  
 d4createCB, 58  
 d4delete, 27, 30, 99, 101  
 d4deleted, 30  
 d4flush, 101, 105  
 d4go, 18  
 d4lockAdd, 110  
 d4lockAddAll, 110, 113, 114, 123  
 d4lockAddAppend, 110  
 d4lockAddFile, 110  
 d4lockAll, 113, 114, 123  
 d4log, 120  
 d4memoCompress, 31, 102  
 d4numFields, 16, 20  
 d4open, 15, 17, 25, 26, 41, 108, 121  
 d4optimize, 95  
 d4optimizeWrite, 95  
 d4pack, 102  
 d4position, 112  
 d4recall, 27, 30, 99, 101  
 d4recCount, 32, 70  
 d4recWidth, 33  
 d4refresh, 95, 113  
 d4refreshRecord, 95, 113  
 d4reindex, 57, 102, 114  
 d4seek, 18, 49, 50, 51, 52, 79, 102, 112, 122

- d4seekDouble, 18, 49, 50, 79, 102, 112
- d4seekN, 18, 49, 50, 79, 102, 112
- d4seekNext, 18, 49, 50, 51, 52, 79, 102, 112
- d4seekNextDouble, 18, 49, 50, 52
- d4seekNextN, 18, 49, 50, 52, 79, 102, 112
- d4skip, 16, 18, 19, 43, 79, 108, 113
- d4tagSelect, 43
- d4top, 18, 19, 43, 79, 102, 112
- d4unlock, 101, 102, 105, 109, 123
- d4write, 101
- d4zap, 102, 114

#### Data Files, 7

- alias, 32
- alias, changing, 32
- buffering. *See* Memory Optimizations
- closing, 18
- composite. *See* Relations
- compressing memo files, 31
- creating, 22, 24, 25
- current record, 14
- deletion flag. *See* Deletion Flag
- exclusive access. *See* Exclusive Access
- field structure, 14, 24, 33
- index files, creating with, 39
- locking. *See* Locking
- master. *See* Relations
- moving between records, 18
- natural order. *See* Natural Order
- opening, 17, 25
- packing, 31
- queries. *See* Queries
- read only mode. *See* Read Only Mode
- record buffer, 14
- record count, 32
- record number. *See* Record Number
- relations. *See* Relations
- slaves. *See* Relations
- top master. *See* Relations

DATA4, 13, 26, 32, 69, 70

- obtaining the pointer to, 13, 17, 25

Databases. *See* Data Files

#### Date Fields

- assigning values, 28, 89
- creating, 24
- retrieving contents, 22, 89

Date Formats. *See* Date Operations

#### Date Functions

- date4assign, 90
- date4cdow, 90
- date4cmonth, 90
- date4day, 91
- date4dow, 91
- date4init, 90
- date4isleap, 91
- date4long, 90
- date4month, 91
- date4today, 91
- date4year, 91

#### Date Operation

- leap year, 91

#### Date Operations, 87

- assigning values to date fields, 28, 89
- CCYYMMDD, 89
- converting between date formats, 89
- date formats, 87
- date pictures, 87
- day of the month, 90, 91
- day of the week, 90, 91
- julian day format, 89
- retrieving contents from date fields, 22, 89
- retrieving the current day, 91
- standard format, 89
- testing for valid dates, 90
- years, 91

Date Pictures. *See* Date Operations

#### dBASE Expressions. *See* Reference Guide

- filter expression, 45
- index expression, 36
- master expression, 59, 69
- query expression, 73
- sort expression, 74

#### dBASE Functions, 36

- DELETED(), 45
- STR(), 52
- UPPER(), 36, 83

#### dBASE IV

- descending order tags, 38, 39
- index files, 34
- index filters, 43
- locking protocol, 100

#### dBASE operators

- +', ', 52

Deadlock. *See* Multi-User

Deleting records. *See* Records

Deletion Flag, 8, 28

Descending Order. *See* Tags

## —E—

e4unique, 46

#### Exact Matches

- seeks. *See* Seeking

#### Exclusive Access

- opening files exclusively, 110, 112

## —F—

#### Field Attributes

- decimal, 7, 24
- length, 7, 24
- name, 7, 24
- type, 7, 24

#### Field Functions

- f4assign, 27, 89
- f4assignDouble, 28
- f4assignInt, 28
- f4assingLong, 28
- f4decimals, 33



- f4double, 22
- f4int, 22
- f4len, 33
- f4long, 22
- f4memoAssign, 27
- f4memoStr, 21
- f4name, 33
- f4str, 21, 89
- f4type, 33
- Field Types
  - character. *See* Character Fields
  - Date. *See* Date Fields
  - determining, 33
  - Floating Point. *See* Numeric Fields
  - Logical. *See* Logical Fields
  - Memo. *See* Memo Fields
  - Numeric. *See* Numeric Fields
- FIELD4, 14, 20, 32, 33
- FIELD4INFO
  - dec member, 24
  - len member, 24
  - name member, 24
  - obtaining a copy of, 33
  - type member, 24
- Fields
  - accessing, 19
  - assigning contents, 27
  - character fields. *See* Character Fields
  - creating, 24
  - date fields. *See* Date Fields
  - decimal, 33
  - field attributes. *See* Field Attributes
  - field number, 20
  - field type, 21
  - floating point fields. *See* Numeric Fields
  - generic access, 20, 21
  - length, 33
  - name, 33
  - numeric fields. *See* Numeric Fields
  - numeric values, retrieving, 22
  - qualifying, 32, 61, 74
  - referencing, 19
  - referencing by field number, 20
  - referencing by name, 20
  - retrieving contents, 21
- Files
  - buffering. *See* Memory Optimizations
  - locking. *See* Locking
  - opening exclusively, 101, 110, 112
  - overwriting, 25
- Filters
  - queries. *See* Queries
  - tag filters. *See* Tags
- FoxPro
  - descending order tags, 38, 39
  - index files, 34
  - index filters, 43
  - locking protocol, 100

## —G—

- Group Files, 53
  - bypassing, 53
  - creating, 53

## —I—

- Index Files, 34
  - .CDX, 10, 34, 36, 40, 53
  - .MDX, 10, 34, 38, 40, 53
  - .NTX, 10, 35, 38, 40, 53, 54, 55, 57
  - buffering. *See* Memory Optimizations
  - creating, 36, 39, 40
  - descending order. *See* Tags
  - exclusive access. *See* Exclusive Access
  - filter expression, 45
  - index expression, 36
  - index key, 36
  - locking. *See* Locking
  - non-production indexes, 36, 38
  - opening, 41, 55
  - production indexes, 10, 36, 38, 39, 40
  - reindexing, 57
  - reindexing a single file, 57
  - reindexing all files, 57
  - search key. *See* Seeking
  - seeking. *See* Seeking
  - tags. *See* Tags
  - unique keys, 38, 39, 46
- INDEX4, 40, 57
- Index4 Functions
  - i4create, 40, 57
  - i4open, 41, 55, 57
  - i4reindex, 102
  - i4tagInfo, 58
- Indexes. *See* Index Files

## —J—

- Julian Day Format. *See* Date Operations

## —K—

- Keys
  - index key. *See* Index Files
  - lookup key. *See* Relations
  - search key. *See* Seeking

## —L—

- Locking
  - automatic data file locking, 102
  - automatic record locking, 101, 112
  - automatic record unlocking, 102, 122
  - Clipper locking protocol, 100
  - data file locking, 100
  - dBASE IV locking protocol, 100
  - efficient locking, 113

- enforced, 99, 101, 106, 112
- explicit, 27, 112, 122, 123
- FoxPro locking protocol, 100
- group locks, 110
- index and memo file locking, 100
- lock attempts, 111, 114
- protocols, 100
- record locking, 100
- unlocking, 102, 122

Locking Protocols. *See* Locking

Log Files, 120, 121

LOG4ALWAYS, 120

LOG4ON, 120

LOG4TRANS, 120

Logging

- automatic, 120
- changing the status, 28, 120
- Client-Server configuration, 120
- current status, 110, 111
- transaction, 117, 120

Logical Fields

- assigning values to, 28
- creating, 24
- retrieving contents, 21

Lookups. *See* Relations

## —M—

Master Data File. *See* Relations

Memo Fields

- assigning values to, 27
- creating, 24

Memo File

- buffering. *See* Memory Optimizations
- compressing, 31
- exclusive access. *See* Exclusive Access
- locking. *See* Locking

Memory

- buffering. *See* Memory Optimizations
- dynamic allocation, 33, 58

Memory Optimizations, 93

- activating, 95
- buffering, 93, 113
- memory requirements, 95, 113
- multi-user mode, 96, 106
- read optimizations, 94, 96
- refreshing the buffers, 95
- single user mode, 96
- specifying files, 93, 95
- suspending, 95
- using, 93
- write optimizations, 94, 96

Multi-User

- appending records, 105, 107
- applications, 99
- applications, creating, 100
- common tasks, 102
- deadlock, 109
- exclusive access. *See* Exclusive Access

- finding records, 105
- listing records, 106
- lock attempts. *See* Locking
- locking protocols. *See* Locking
- memory optimizations. *See* Memory Optimizations
- modifying records, 105
- opening files, 104
- optimizations, 106
- read only mode. *See* Read Only Mode

## —N—

Natural Order

- selecting, 43

Numeric Fields

- assigning values, 27, 28
- creating, 24
- retrieving contents, 22

## —O—

Optimizations. *See* Memory Optimizations

## —P—

Packing. *See* Data Files

Partial Matches

- relations. *See* Relations, approximate match
- seeks. *See* Seeking

Performance Issues

- efficient appending, 114
- efficient locking, 113
- efficiently sorting a query set, 75
- functions that reduce performance, 113
- general tips, 115
- memory optimization, 113
- time consuming functions, 114

Production Indexes. *See* Index Files, production indexes

## —Q—

Queries, 59, 72

- creating, 73
- generating, 73, 75
- query expression, 73
- query set, 72, 74, 75
- sort expression, 74
- sorting, 74
- using Query Optimization, 81

Query Expression. *See* Queries

Query Optimization, 81

- how to use, 84
- queries, 82
- relate4optimizeable, 84
- relations, 81
- requirements, 82

Query Set. *See* Queries

## —R—

- r4after, 51
- r4bin, 24
- r4date, 24
- r4descending, 38, 39
- r4eof, 52
- r4float, 24
- r4gen, 24
- r4locked, 111
- r4log, 24
- r4logOff, 120, 121
- r4logOn, 120
- r4memo, 24
- r4num, 24
- r4str, 24
- r4success, 12, 19, 51
- r4unique, 46
- r4uniqueContinue, 46
- Read Only Mode
  - opening files in, 101, 110, 111
- Read Optimizations. *See* Memory Optimizations
- Record Buffer. *See* Data Files
- Record Buffering. *See* Memory Optimizations
- Record Count, 16
- Record Number, 8
- Records, 8, 71
  - adding, 26
  - appending, 26
  - buffering. *See* Memory Optimizations
  - deleting, 28, 30
  - deletion status, 30
  - listing records, 67
  - physically removing, 31
  - recalling, 30
  - removing, 30
  - undeleting, 30
  - width, 33
- Reindexing. *See* Index Files
- RELATE4, 69
- relate4approx, 70
- relate4blank, 71
- relate4exact, 70
- relate4scan, 70
- relate4skipRec, 71
- relate4terminate, 71
- Relation Functions
  - relate4bottom, 71
  - relate4changed, 76
  - relate4createSlave, 69
  - relate4doAll, 79
  - relate4doOne, 79
  - relate4errorAction, 71
  - relate4init, 69
  - relate4lockAdd, 110
  - relate4skip, 71
  - relate4skipEnable, 72
  - relate4sortSet, 74
  - relate4top, 71

- relate4type, 70
- Relation Set. *See* Relations
- Relation Types. *See* Relations
- Relations, 59
  - approximate match relation, 66
  - complex, 63
  - composite data file, 61
  - composite record, 60
  - creating, 67
  - error action, 71
  - exact match relation, 65
  - lookup key, 59
  - lookups, 77, 79
  - many to many relation, 66
  - many to one relation, 65
  - master, 59, 69
  - master expression, 59, 69
  - multi-layered, 63
  - one to many relation, 66
  - one to one relation, 65
  - relation set, 63, 69
  - scan relation, 66
  - skipping, 72
  - skipping, backwards, 72
  - slave tag, 59, 69
  - slaves, adding, 69
  - slaves, creating without tags, 70
  - top master, 63, 69
  - top master, specifying, 69
  - types, 65
  - using Query Optimization, 81
- Removing
  - removing records. *See* Records
- Rereading, 95

## —S—

- Seeking
  - character tags, 49
  - compound keys, 52
  - date tags, 49
  - exact matches, 51
  - incremental, 50
  - no matches, 52
  - numeric tags, 49
  - partial matches, 51
  - performing seeks, 49
  - search key, 47
  - soft seek. *See* Seeking, partial matches
- Single User
  - memory optimizations, 96
- Slave Data Files. *See* Relations
- Soft Seek. *See* Seeking, partial matches
- Sorting
  - queries. *See* Queries
  - records. *See* Indexes
- Standard Format. *See* Date Operations
- Strings
  - copying to fields, 27

retrieving from fields, 21  
 Structures  
   CodeBase structures, 12

## —T—

Tables, 7  
 Tag Functions  
   t4open, 57  
   t4uniqueSet, 83  
 TAG4, 37, 41, 42, 43, 44, 47, 55, 56, 57, 67, 69, 70,  
   76, 78, 79, 102, 117  
 TAG4INFO, 38  
   arrays, 39  
   descending member, 38, 39  
   expression member, 38, 39  
   filter member, 38, 39  
   name member, 38  
   unique member, 38, 39  
 Tags  
   creating, 38  
   currently selected, 42  
   default tag, 42  
   deleting, 40  
   descending order, 38, 39  
   effects of, 43  
   filters, 38, 39, 43  
   filters, creating, 45  
   index expression, 36, 38, 39  
   index expression, compound, 36

index key, 36, 51  
 name, 38  
 natural order. *See* Natural Order  
 referencing, 41  
 referencing by name, 42  
 search key. *See* Seeking  
 seeking. *See* Seeking  
 selecting, 43  
 slave tag. *See* Relations  
 Top Master. *See* Relations  
 Transactions, 117  
   client-server configuration, 120  
   committing, 117, 119  
   implementation, 120  
   log files, 121  
   logging status, 120  
   rollback, 117, 119  
   stand-alone configuration, 122  
   starting, 117, 119  
 Tuples, 7

## —U—

u4free, 58  
 Unlocking. *See* Locking

## —W—

Write Optimizations. *See* Memory Optimizations

